

# Tiled & GPU-accelerated Image Processing

## Robert Haase

Funded by



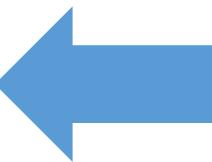
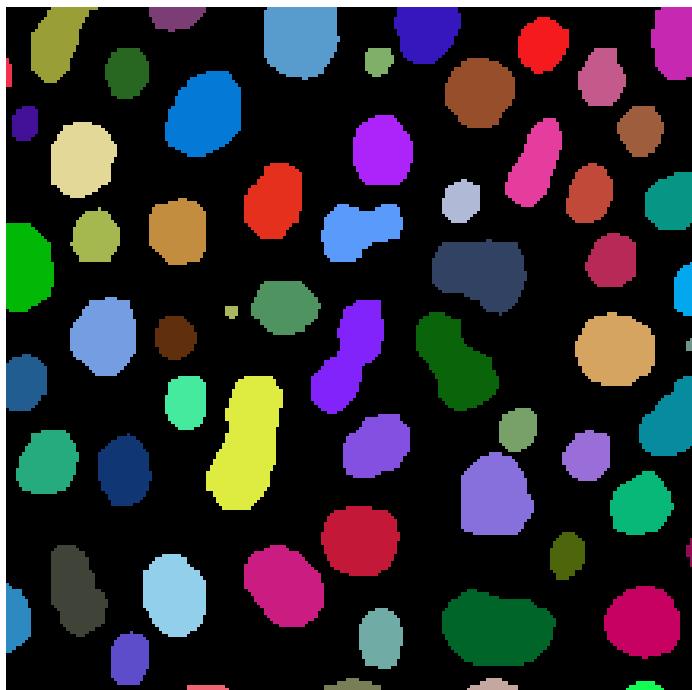
Bundesministerium  
für Bildung  
und Forschung



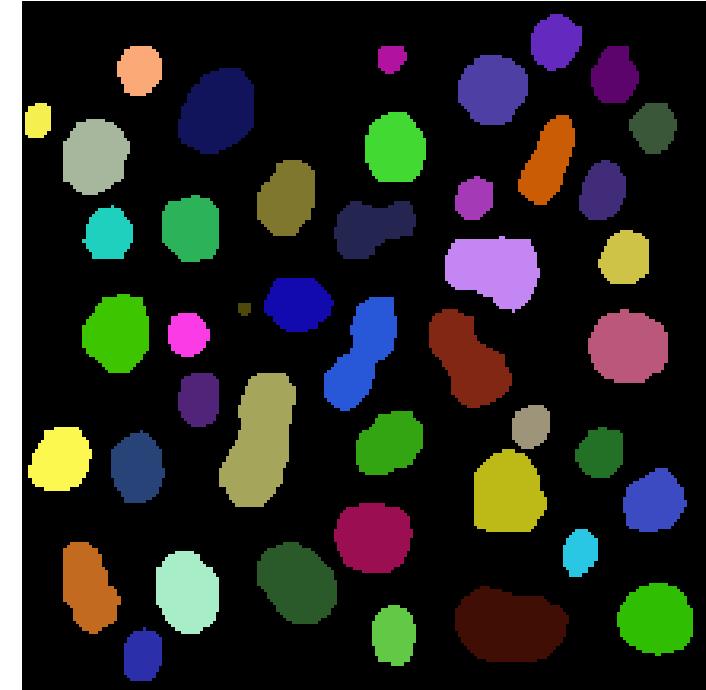
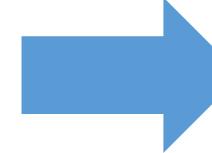
Diese Maßnahme wird gefördert durch die Bundesregierung  
aufgrund eines Beschlusses des Deutschen Bundestages.  
Diese Maßnahme wird mitfinanziert durch Steuermittel auf  
der Grundlage des von den Abgeordneten des Sächsischen  
Landtags beschlossenen Haushaltes.

# Quiz:

- How can we estimate the number of cells in an image?

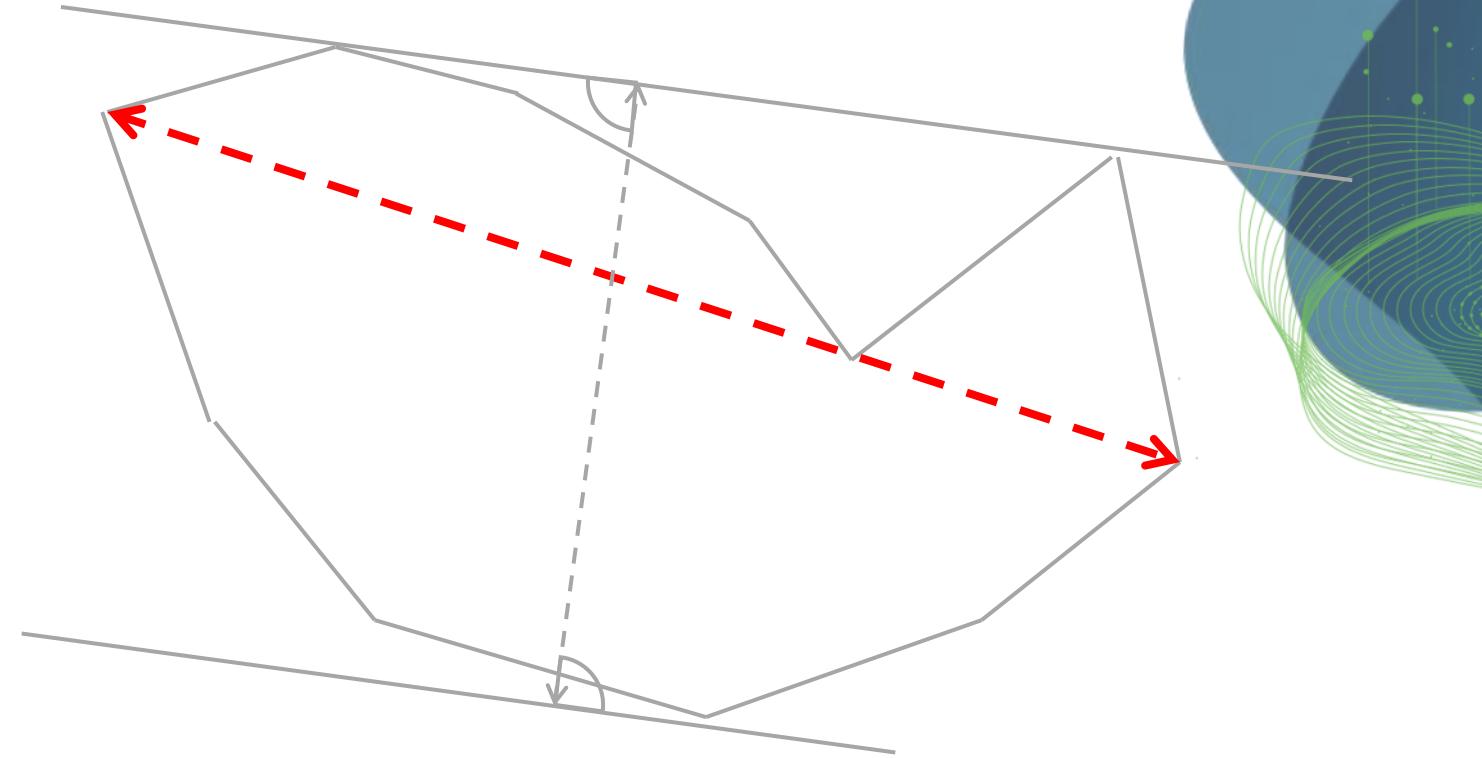


?



# Quiz: Recap

How is this feature called?



Feret's diameter



Minimum Caliper



Minor Axis length



Major Axis length



# GPU-accelerated Image Processing

## Robert Haase

Funded by



Bundesministerium  
für Bildung  
und Forschung

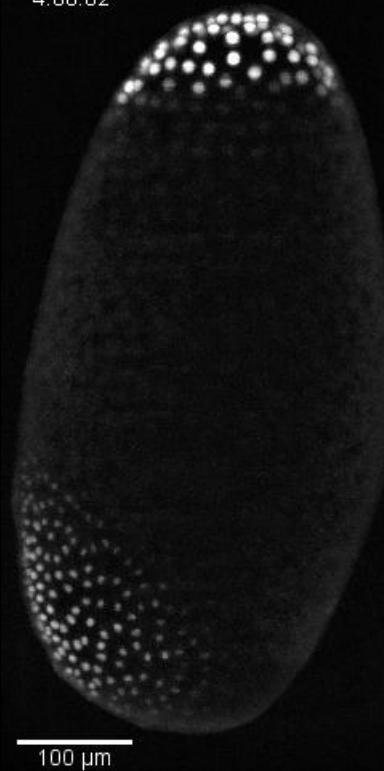


Diese Maßnahme wird gefördert durch die Bundesregierung  
aufgrund eines Beschlusses des Deutschen Bundestages.  
Diese Maßnahme wird mitfinanziert durch Steuermittel auf  
der Grundlage des von den Abgeordneten des Sächsischen  
Landtags beschlossenen Haushaltes.

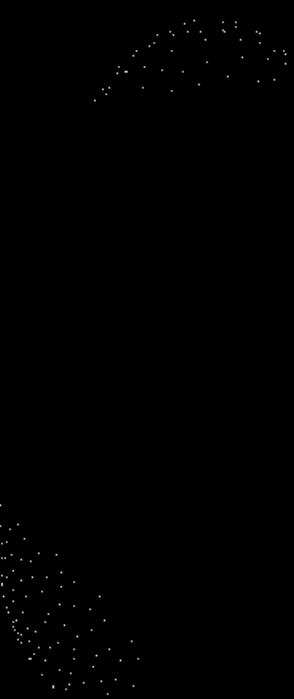
# GPU accelerated image processing in life sciences

- ... to study embryo development

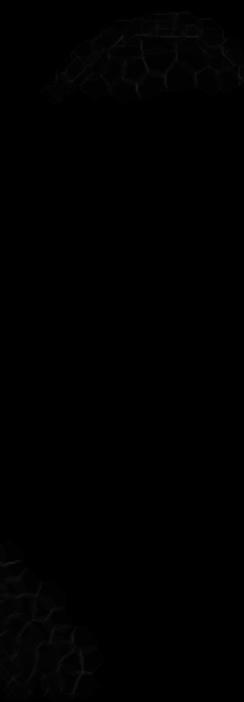
Tribolium castaneum nuclei-  
GFP,  
Background subtracted  
4:00:02



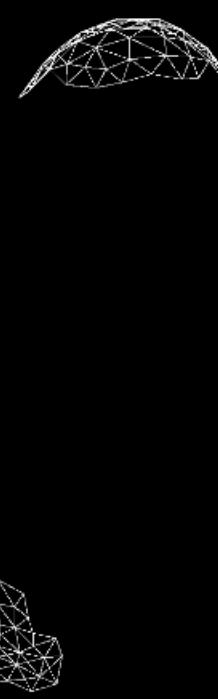
Spot detection (3D)



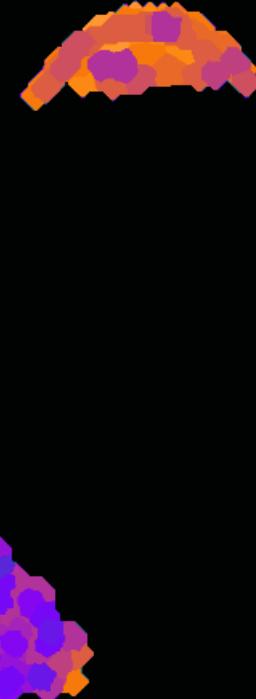
Theoretical membranes  
(pseudo Voronoi map)



Neighbor mesh



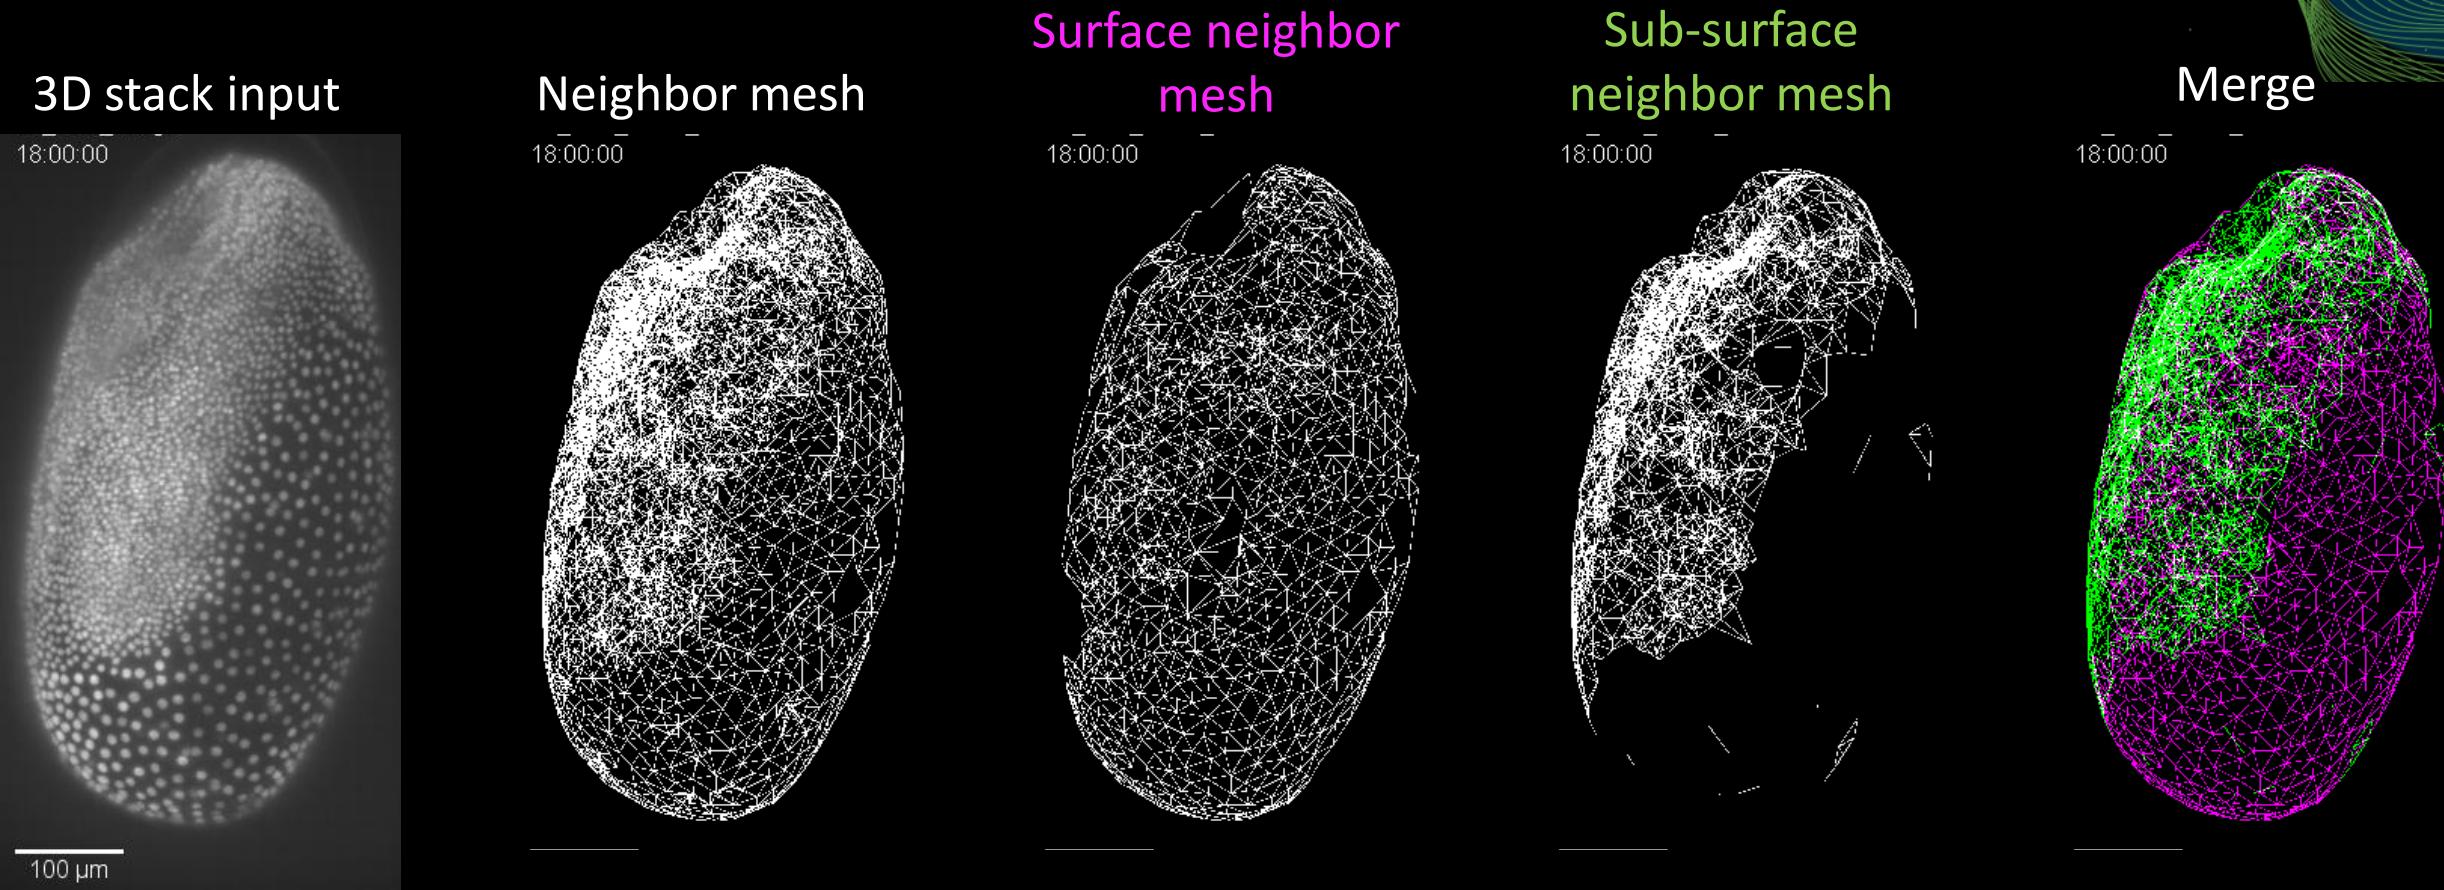
Average centroid distance of  
neighbors



0  $\mu\text{m}$  35

# GPU accelerated image processing in life sciences

- Raytracing enables differentiating surface and sub-surface mesh nodes



# Image processing in life-sciences

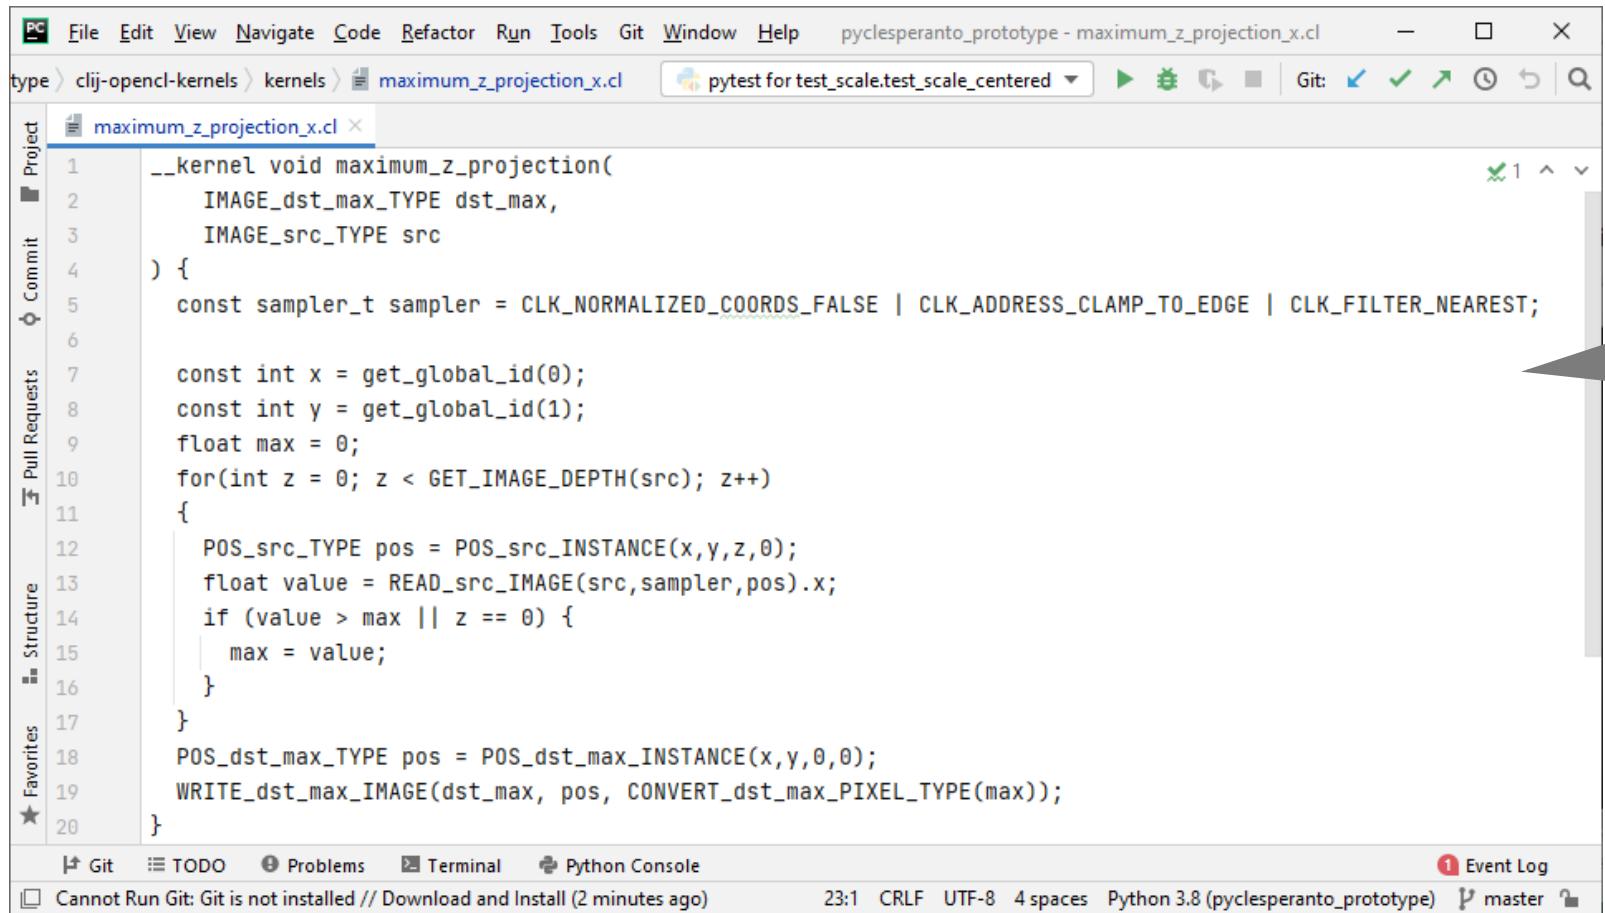
- State-of-the-art software for more than 20 years: ImageJ / Fiji



2x

# OpenCL-based GPU-acceleration

- GPU-acceleration? Learn the Open Computing Language (OpenCL)!



The screenshot shows an IDE interface with the following details:

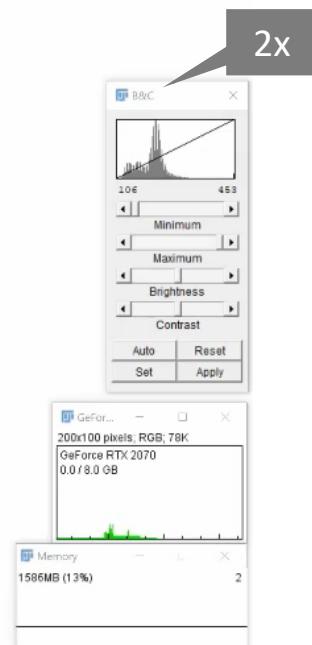
- Title Bar:** pyclesperanto\_prototype - maximum\_z\_projection\_x.cl
- Toolbar:** File, Edit, View, Navigate, Code, Refactor, Run, Tools, Git, Window, Help.
- Project Explorer:** Shows a project named "clij-opencl-kernels" with a file "maximum\_z\_projection\_x.cl" selected.
- Code Editor:** Displays the following OpenCL kernel code:

```
__kernel void maximum_z_projection(
    IMAGE_dst_max_TYPE dst_max,
    IMAGE_src_TYPE src
) {
    const sampler_t sampler = CLK_NORMALIZED_COORDS_FALSE | CLK_ADDRESS_CLAMP_TO_EDGE | CLK_FILTER_NEAREST;

    const int x = get_global_id(0);
    const int y = get_global_id(1);
    float max = 0;
    for(int z = 0; z < GET_IMAGE_DEPTH(src); z++)
    {
        POS_src_TYPE pos = POS_src_INSTANCE(x,y,z,0);
        float value = READ_src_IMAGE(src,sampler,pos).x;
        if (value > max || z == 0) {
            max = value;
        }
    }
    POS_dst_max_TYPE pos = POS_dst_max_INSTANCE(x,y,0,0);
    WRITE_dst_max_IMAGE(dst_max, pos, CONVERT_dst_max_PIXEL_TYPE(max));
}
```
- Bottom Status Bar:** Git: Cannot Run Git: Git is not installed // Download and Install (2 minutes ago), 23:1 CRLF UTF-8 4 spaces Python 3.8 (pyclesperanto\_prototype), master.

Maximum  
intensity  
projection  
along Z

# User-friendly GPU-acceleration



# GPU-accelerated image processing

- Performance depends on operation, image size, parameters, hardware, ....

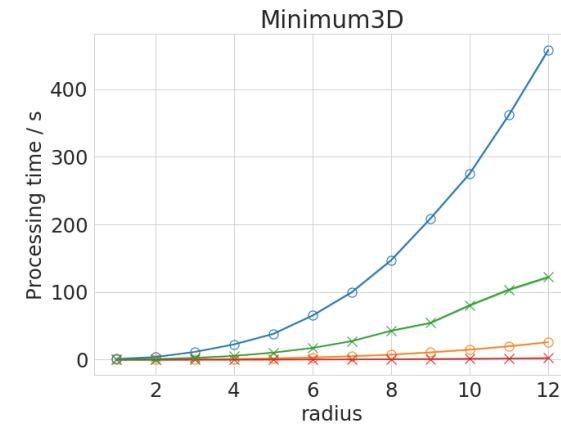
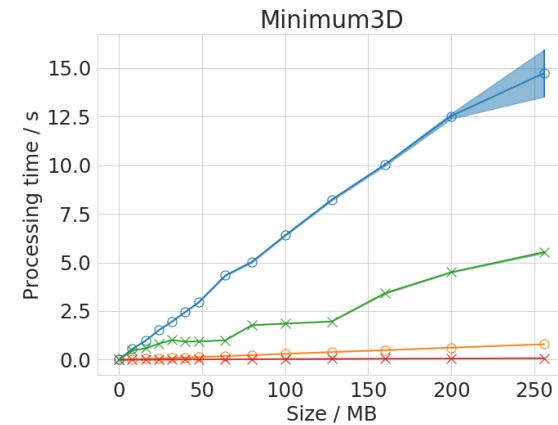
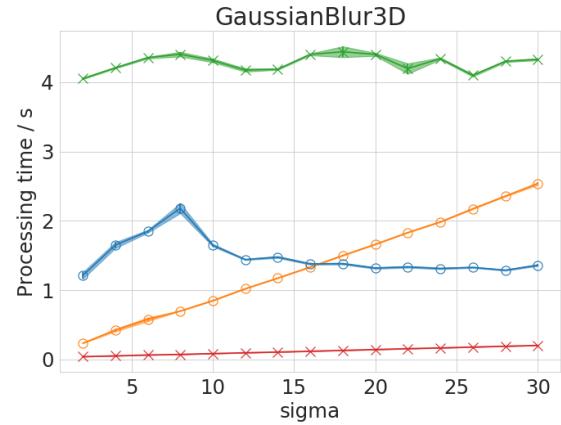
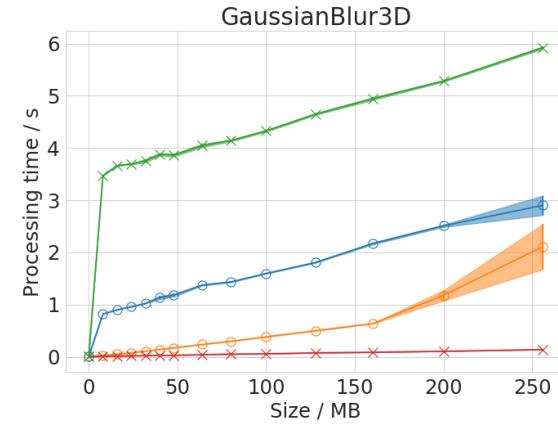


vs.



Workstation CPU  
2x Intel Xeon Silver  
4110  
Laptop CPU  
Intel Core i7-8650U

Workstation GPU  
Nvidia Quadro  
P6000  
Laptop GPU  
Intel UHD 620 GPU



# GPU-acceleration only makes sense if images are large

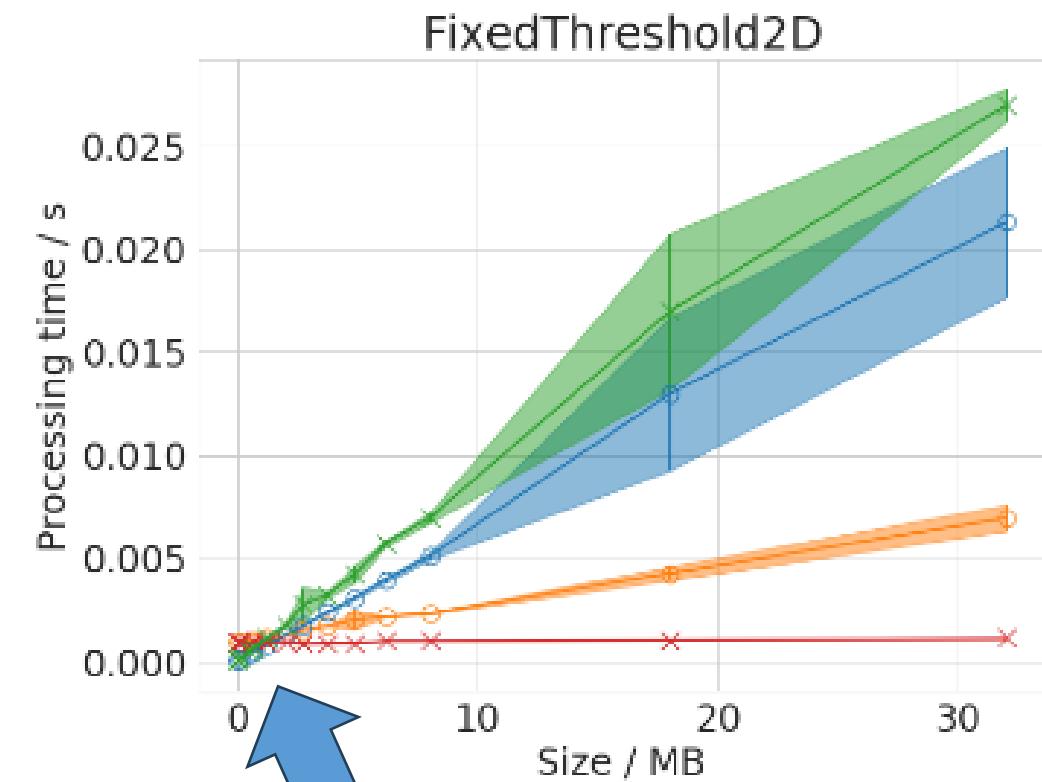
- Because the CPU cache is faster than GPU-memory.
- CPU cache is also very small though.
- Example:

Intel Core i7-8650u\*

Cache	
Cache L1:	64 KB (per core)
Cache L2:	256 KB (per core)
Cache L3:	8 MB (shared)

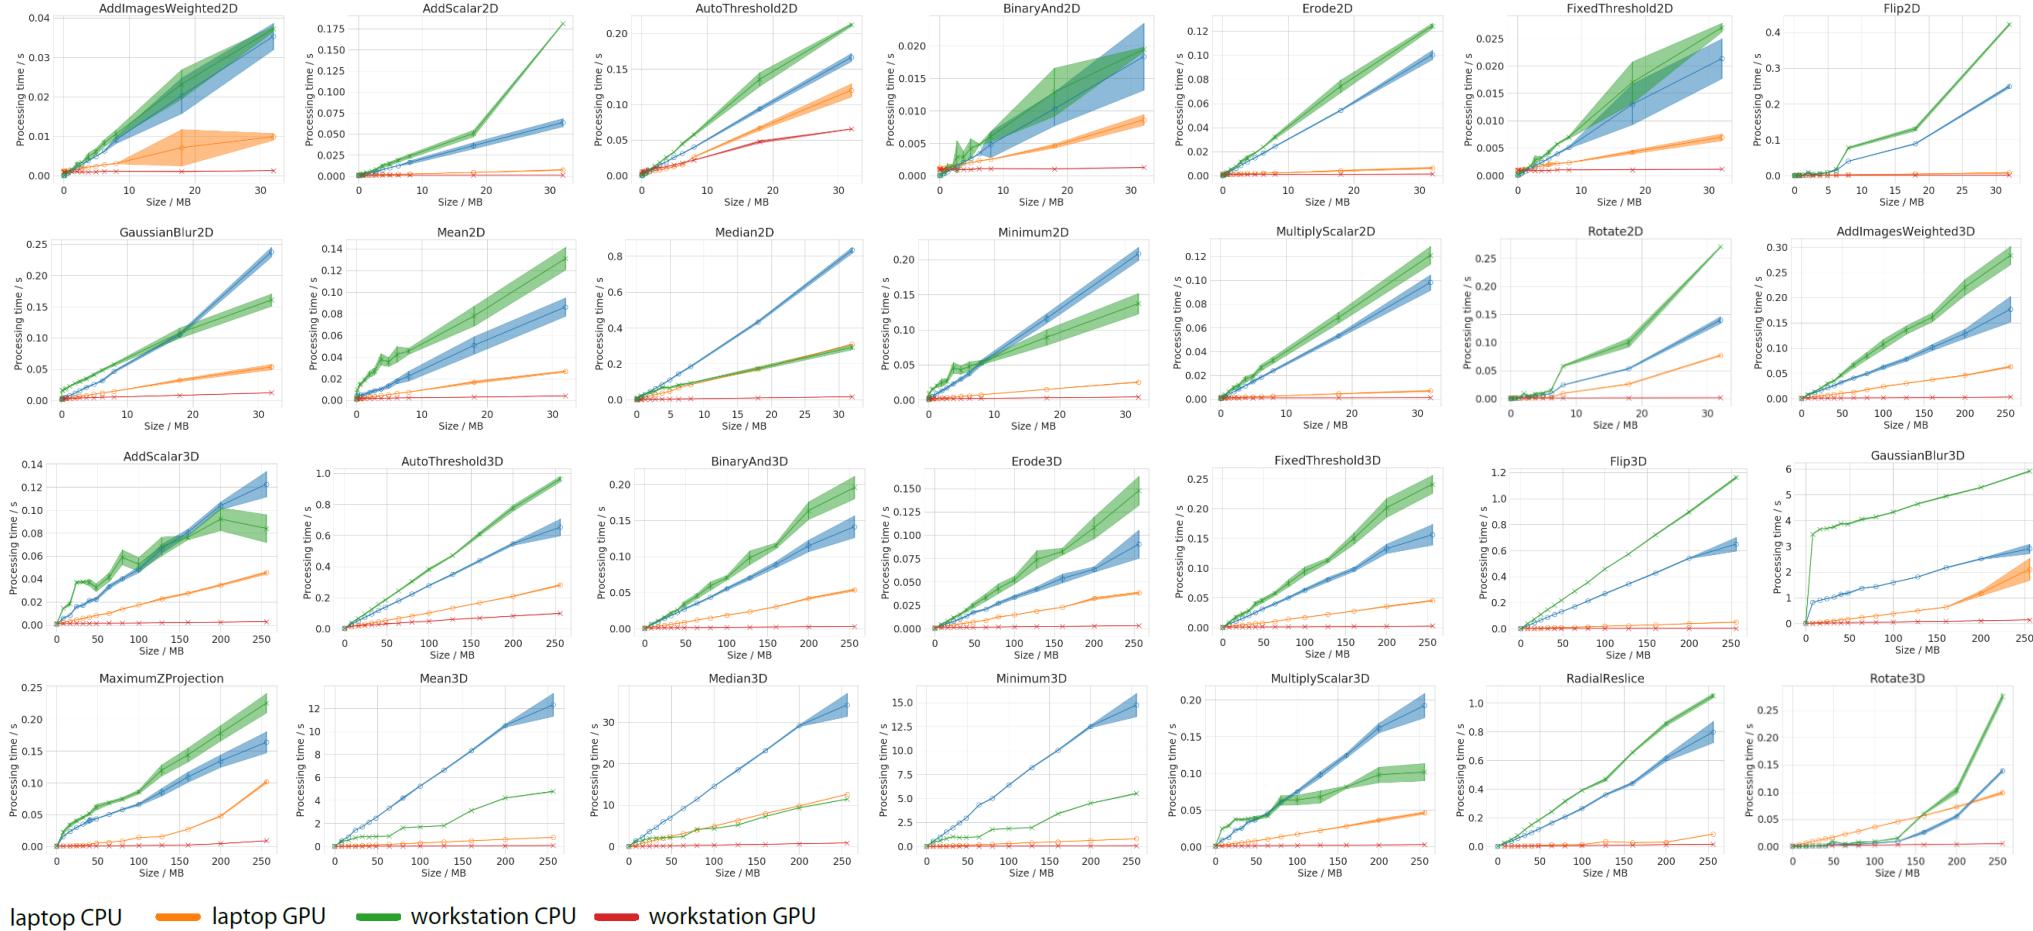
Workstation CPU  
2x Intel Xeon Silver  
4110  
Laptop CPU  
Intel Core i7-8650U

Workstation GPU  
Nvidia Quadro  
P6000  
Laptop GPU  
Intel UHD 620 GPU



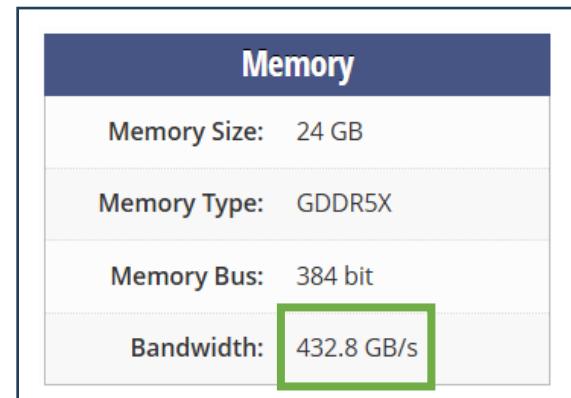
# GPU-accelerated image processing

- Performance depends on operation, image size, parameters, hardware, ....



# GPU-accelerated image processing

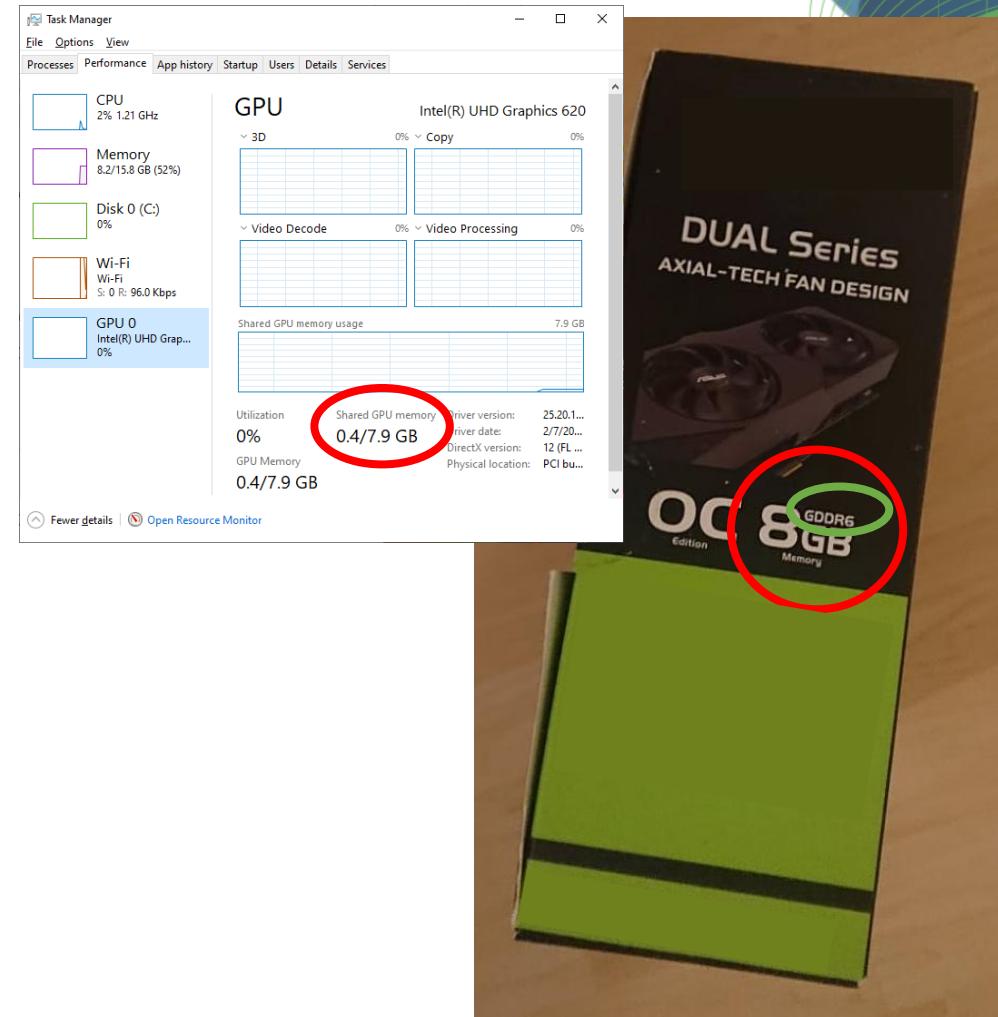
Speedup compared to Laptop CPU			
	Intel UHD 620 GPU (Laptop)	Nvidia Quadro P6000 (workstation)	
8 MB (2D)	AddImagesWeighted2D	3	8
	AddScalar2D	7	14
	AutoThreshold2D	2	2
	BinaryAnd2D	2	4
	Erode2D	11	20
	FixedThreshold2D	2	5
	Flip2D	16	37
	GaussianBlur2D	3	9
	Mean2D	3	10
	Median2D	2	35
	Minimum2D	7	22
	MultiplyScalar2D	10	21
	Rotate2D	3	22
	AddImagesWeighted3D	3	26
64 MB (3D)	AddScalar3D	3	23
	AutoThreshold3D	3	5
	BinaryAnd3D	3	24
	Erode3D	2	13
	FixedThreshold3D	4	30
	Flip3D	15	119
	GaussianBlur3D	6	35
	MaximumZProjection	7	46
	Mean3D	18	150
	Median3D	3	43
	Minimum3D	23	188
	MultiplyScalar3D	4	28
	RadialReslice	14	42
	Rotate3D	0.1	2



# Checklist: When does GPU-accelerated image processing make sense?

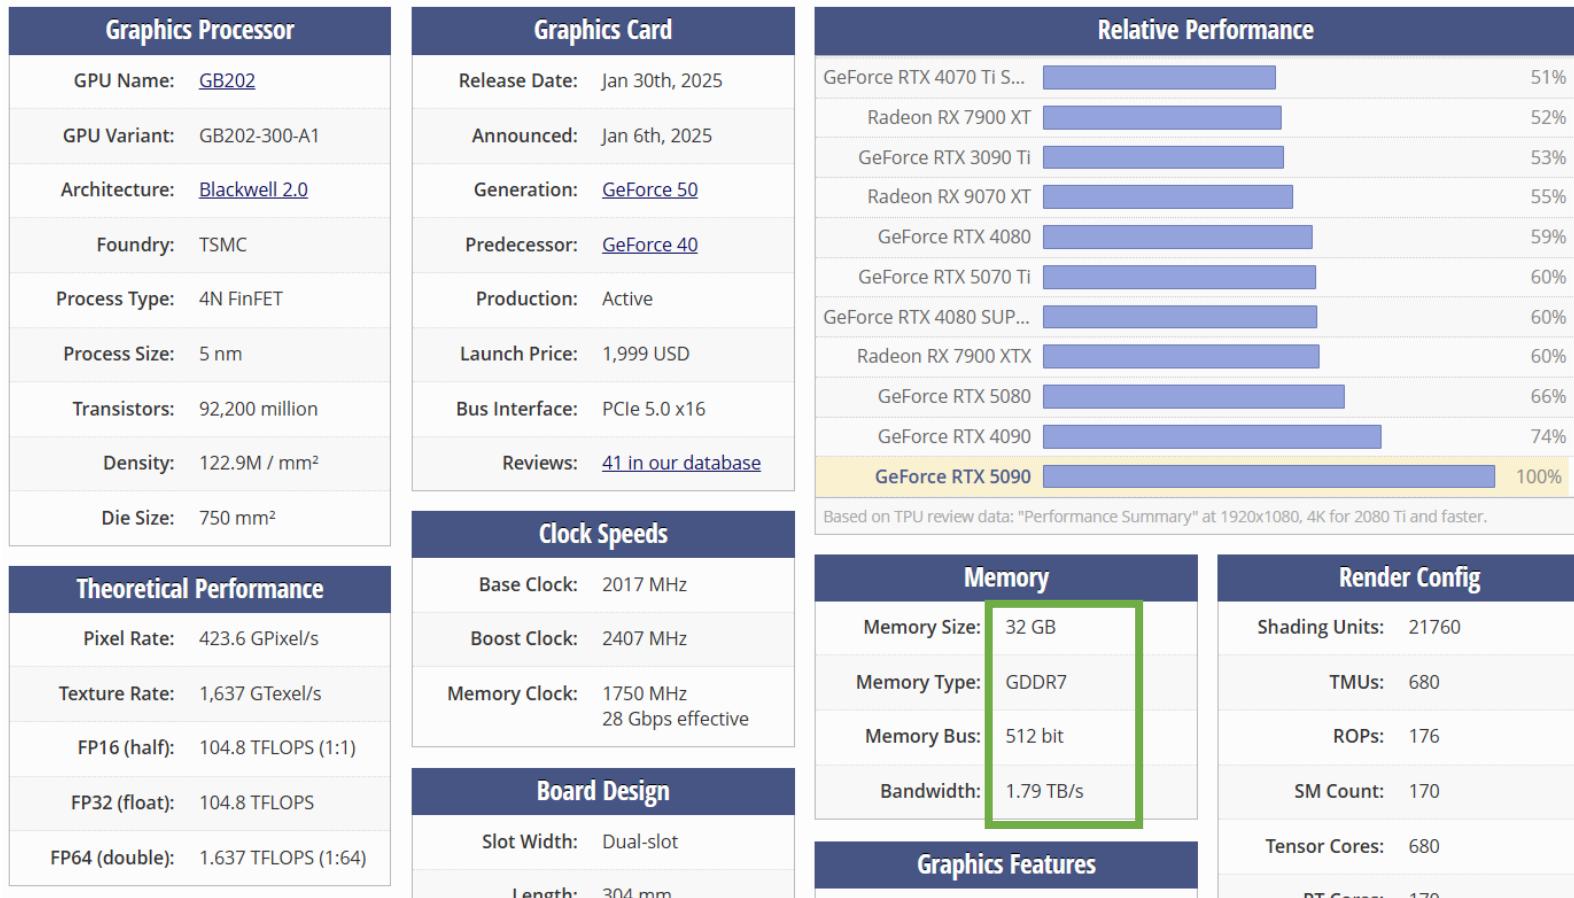
- In order to accelerate your image analysis workflow
- The pre-existing workflow should be slow; ideally:  $(\text{processing time} / \text{loading time}) > 10$ ,
- Memory requirements: rule of thumb: image size in GB  $\times$  number of processing steps + 1) should **fit in your graphics card memory**,

- If you really want to get the most out of it, you should own a graphics card with **GDDR6** memory (memory bandwidth  $> 400 \text{ Gb/s}$ ).
- Comparison: common DDR4 memory has a bandwidth of about 40 GB/s



# Nvidia RTX 5090 specs

- Image processing is “memory bound”; important specs:



Quiz: What does “memory bound” mean?

# Nvidia RTX 4070 versus 4070 mobile

- Vendors name GPUs similarly, despite their configuration is quite different.

## RTX 4070

Graphics Processor	
GPU Name:	AD104
GPU Variant:	AD104-250-A1
Architecture:	<a href="#">Ada Lovelace</a>
Foundry:	TSMC
Process Size:	5 nm
Transistors:	35,800 million
Density:	121.8M / mm <sup>2</sup>
Die Size:	294 mm <sup>2</sup>

## RTX 4070 mobile

Graphics Processor	
GPU Name:	AD106
GPU Variant:	GN21-X6
Architecture:	<a href="#">Ada Lovelace</a>
Foundry:	TSMC
Process Size:	5 nm
Transistors:	22,900 million
Density:	121.8M / mm <sup>2</sup>
Die Size:	188 mm <sup>2</sup>

# Nvidia RTX 4070 versus 4070 mobile

- Vendors name GPUs similarly, despite their configuration is quite different

Quiz: Why are mobile GPUs throttled?

## RTX 4070

Clock Speeds	
Base Clock:	1920 MHz
Boost Clock:	2475 MHz
Memory Clock:	1313 MHz 21 Gbps effective
Memory	
Memory Size:	12 GB
Memory Type:	GDDR6X
Memory Bus:	192 bit
Bandwidth:	504.2 GB/s

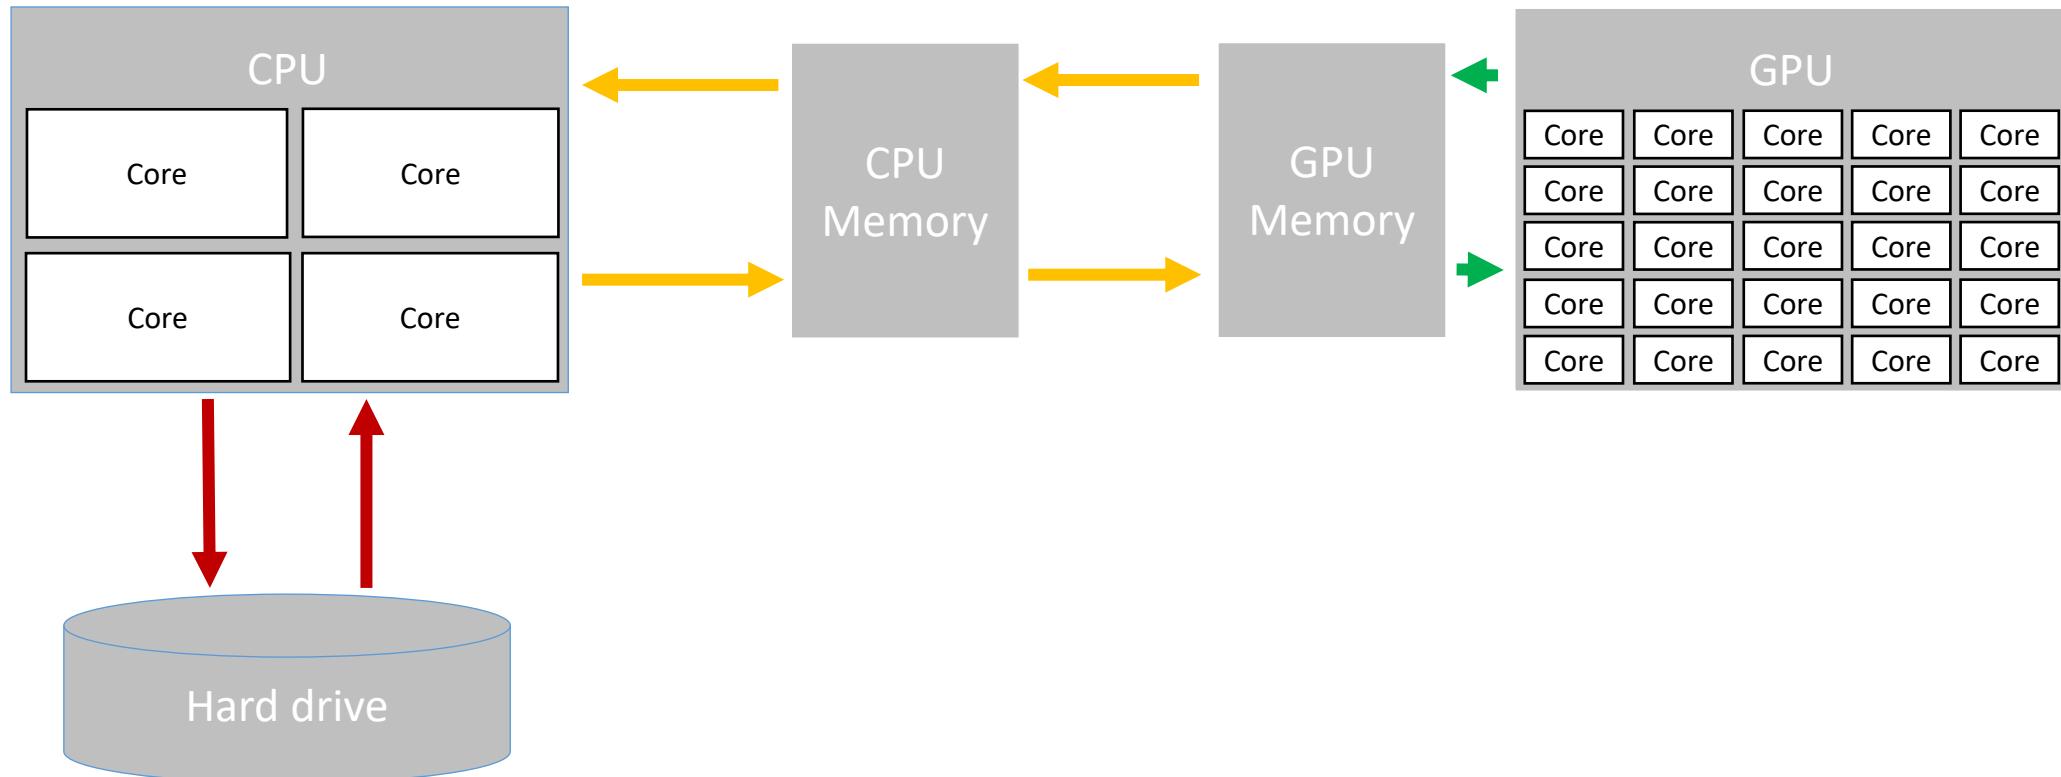
Board Design	
Slot Width:	Dual-slot
Length:	240 mm 9.4 inches
Width:	110 mm 4.3 inches
Height:	40 mm 1.6 inches
TDP:	200 W
Suggested PSU:	550 W

Clock Speeds	
Base Clock:	1395 MHz
Boost Clock:	1695 MHz
Memory Clock:	2000 MHz 16 Gbps effective
Memory	
Memory Size:	8 GB
Memory Type:	GDDR6
Memory Bus:	128 bit
Bandwidth:	256.0 GB/s

Board Design	
Slot Width:	IGP
TDP:	115 W
Outputs:	Portable Device Dependent
Power Connectors:	None

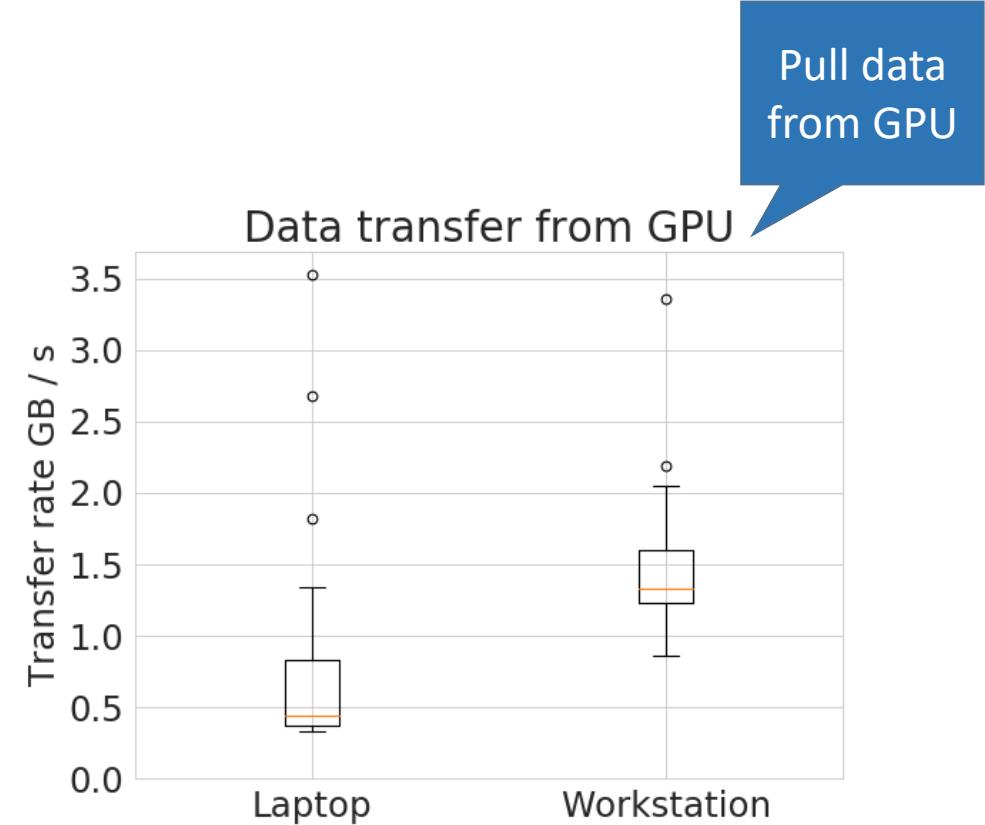
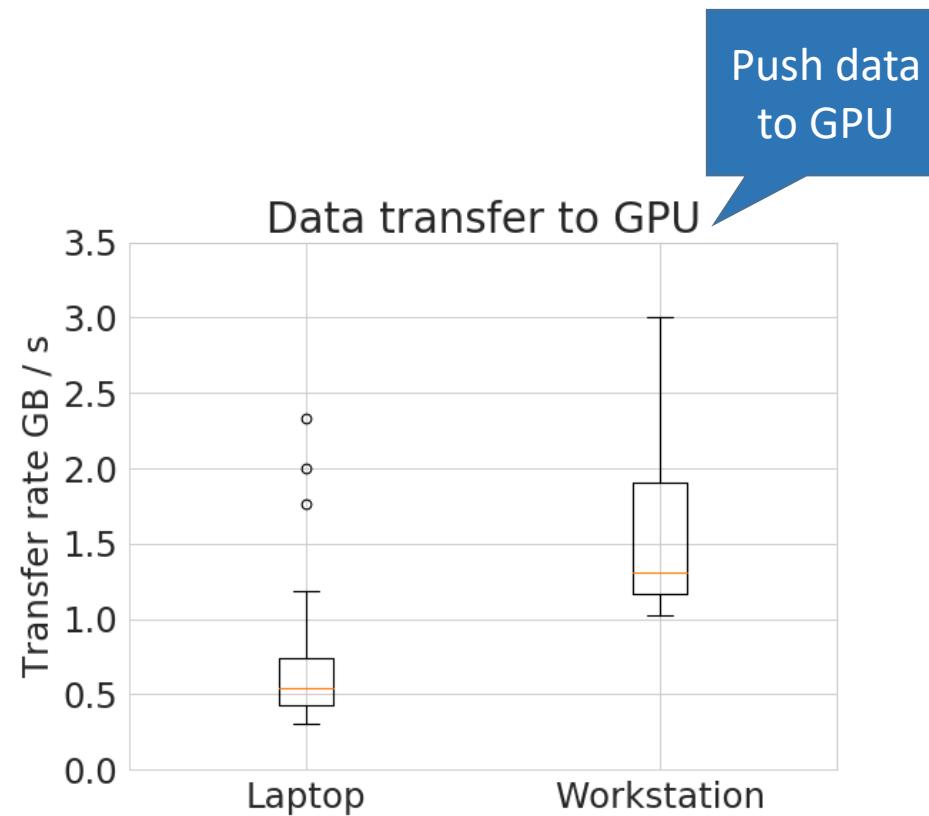
# GPUs allow real-time image processing

- GPUs are specialised in processing, very fast thanks to many cores and fast memory access



# Data transfer takes time

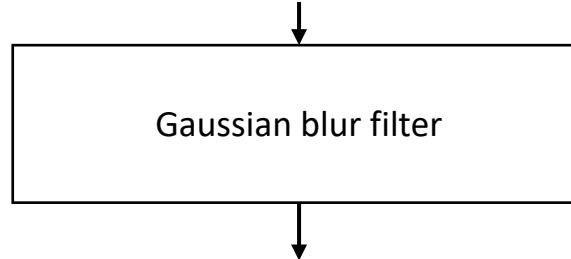
- Data transfer is the bottle neck



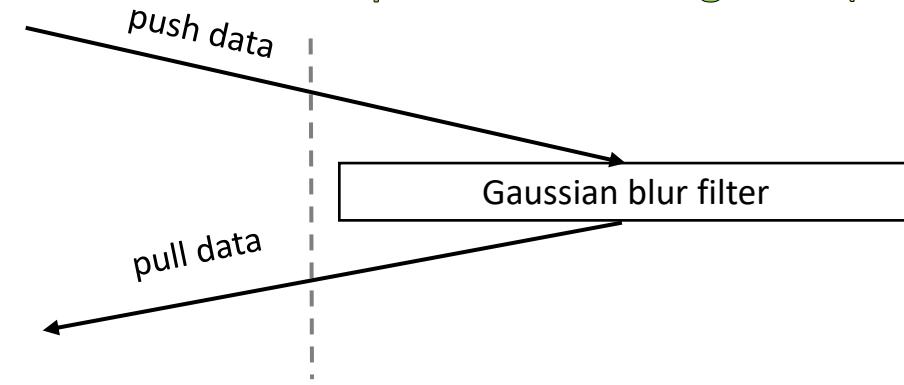
# Build workflows consisting of many operations

GPU acceleration may suffer from data transfer between CPU and GPU

## Central Processing Unit (CPU)



## Graphics Processing Unit (GPU)



# Optimal performance through smart memory management

- Example workflow processing a *Drosophila melanogaster* embryo, histone-GFP

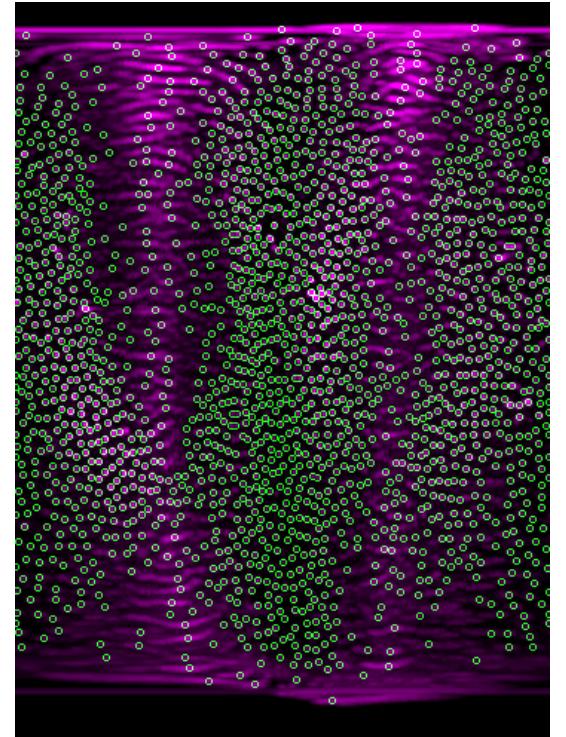
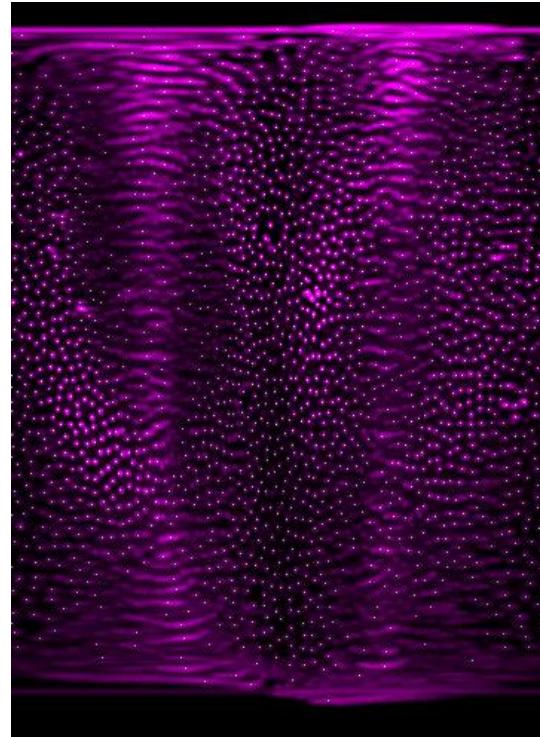
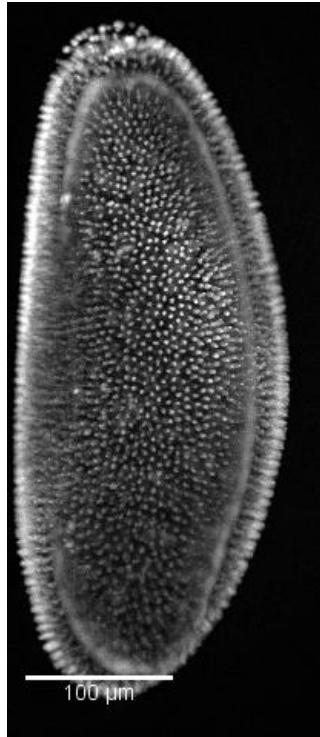
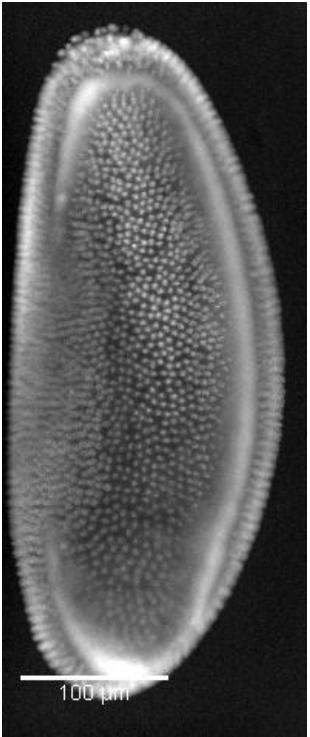
Load data

Preprocessing

Transformation

Segmentation

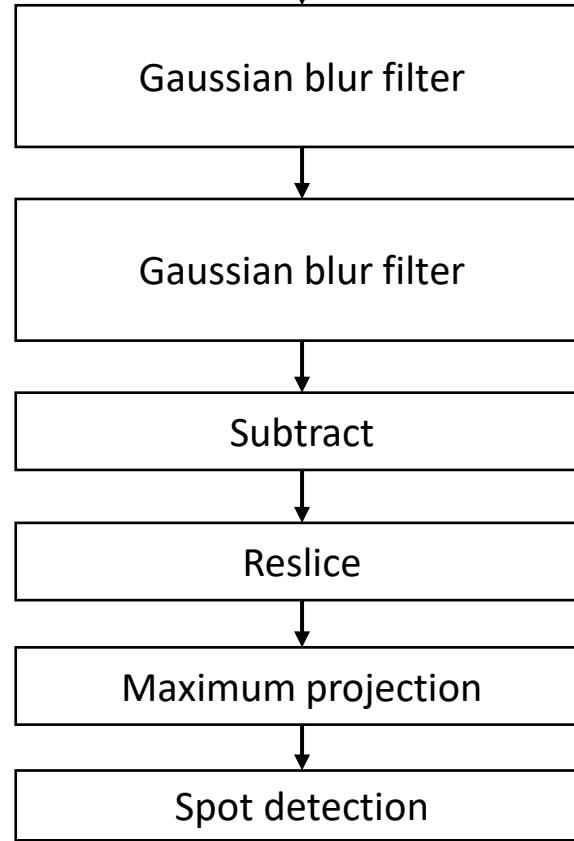
Save data



# Build workflows consisting of many operations

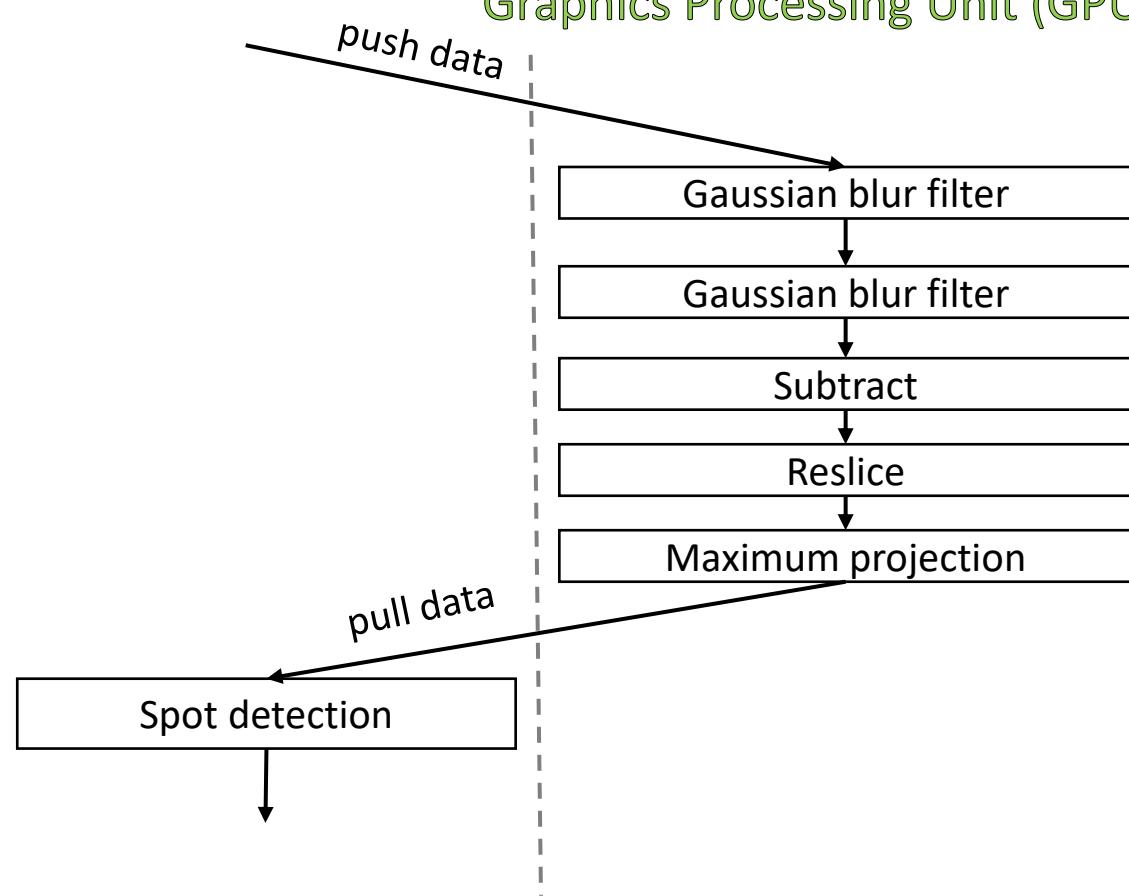
GPU acceleration may suffer from data transfer between CPU and GPU

## Central Processing Unit (CPU)



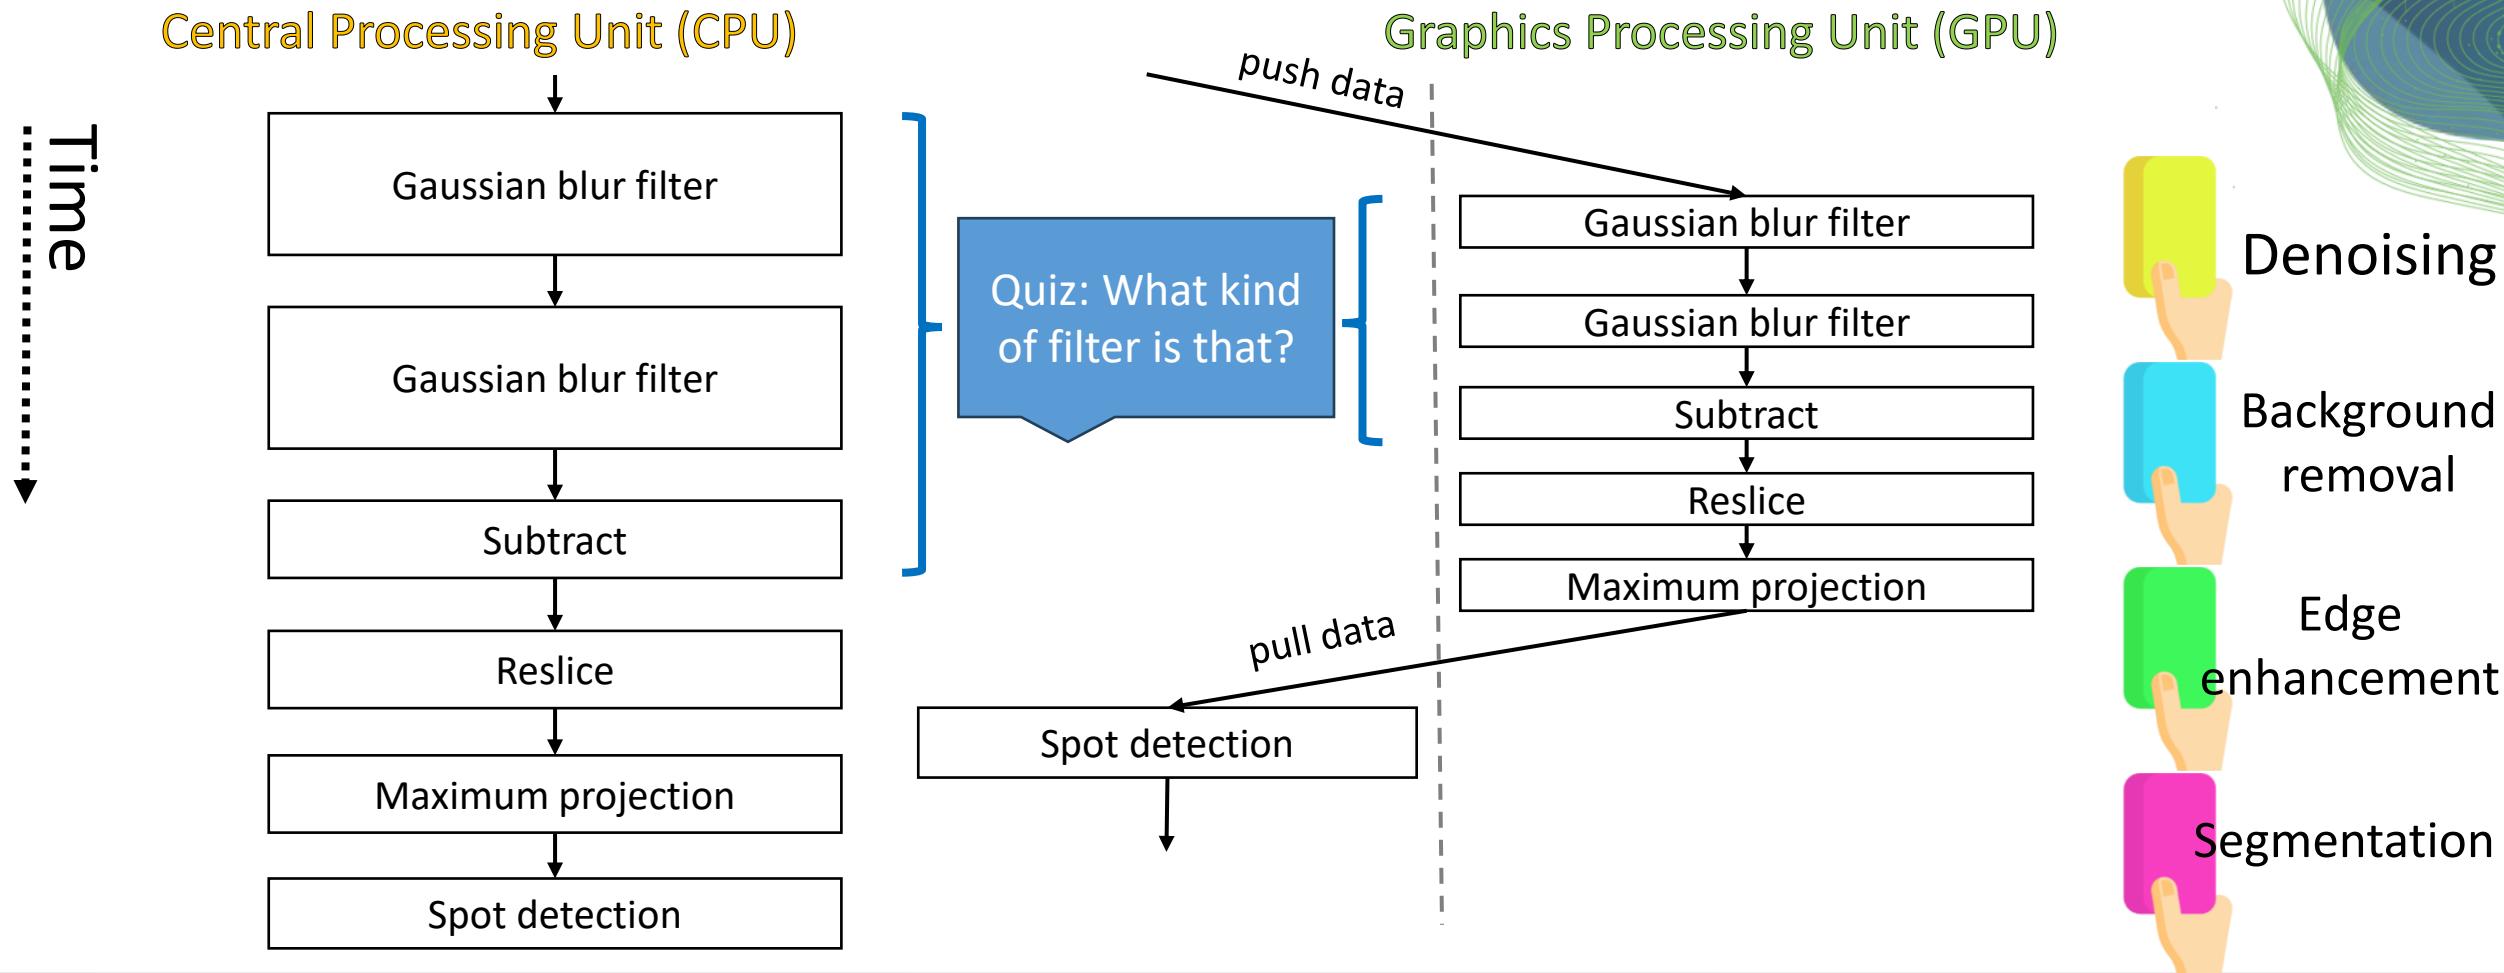
Time ↓

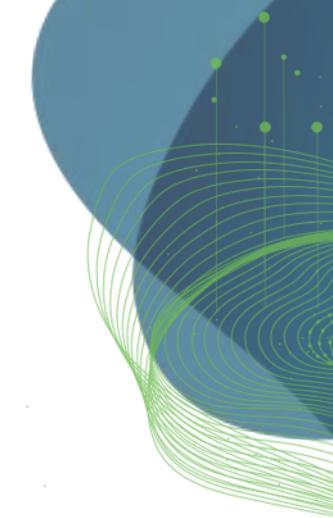
## Graphics Processing Unit (GPU)



# Build workflows consisting of many operations

GPU acceleration may suffer from data transfer between CPU and GPU





# GPU-accelerated Image Processing in Python: OpenCL / clesperanto

Robert Haase

Funded by



Bundesministerium  
für Bildung  
und Forschung



Diese Maßnahme wird gefördert durch die Bundesregierung  
aufgrund eines Beschlusses des Deutschen Bundestages.  
Diese Maßnahme wird mitfinanziert durch Steuermittel auf  
der Grundlage des von den Abgeordneten des Sächsischen  
Landtags beschlossenen Haushaltes.

# GPU-accelerated image processing

- Listing available GPUs

```
[2]: import pyclesperanto as cle  
  
# List available devices  
cle.available_device_names()
```

```
[2]: ['NVIDIA GeForce RTX 4070 Laptop GPU', 'gfx1103']
```

Integrated  
AMD GPU

# GPU-accelerated image processing

- Choosing a GPU

```
[2]: import pyclesperanto as cle  
  
# List available devices  
cle.available_device_names()
```

```
[2]: ['NVIDIA GeForce RTX 4070 Laptop GPU', 'gfx1103']
```

```
[3]: # select a specific device with only a part of its name  
cle.select_device("TX")
```

```
[3]: (OpenCL) NVIDIA GeForce RTX 4070 Laptop GPU (OpenCL 3.0 CUDA)  
      Vendor: NVIDIA Corporation  
      Driver Version: 560.94  
      Device Type: GPU  
      Compute Units: 36  
      Global Memory Size: 8187 MB  
      Local Memory Size: 0 MB  
      Maximum Buffer Size: 2046 MB  
      Max Clock Frequency: 1230 MHz  
      Image Support: Yes
```

Quiz: Why does it make sense to limit the maximum buffer size to about  $\frac{1}{4}$  of the global memory size?

# GPU-accelerated image processing

- Memory constraints

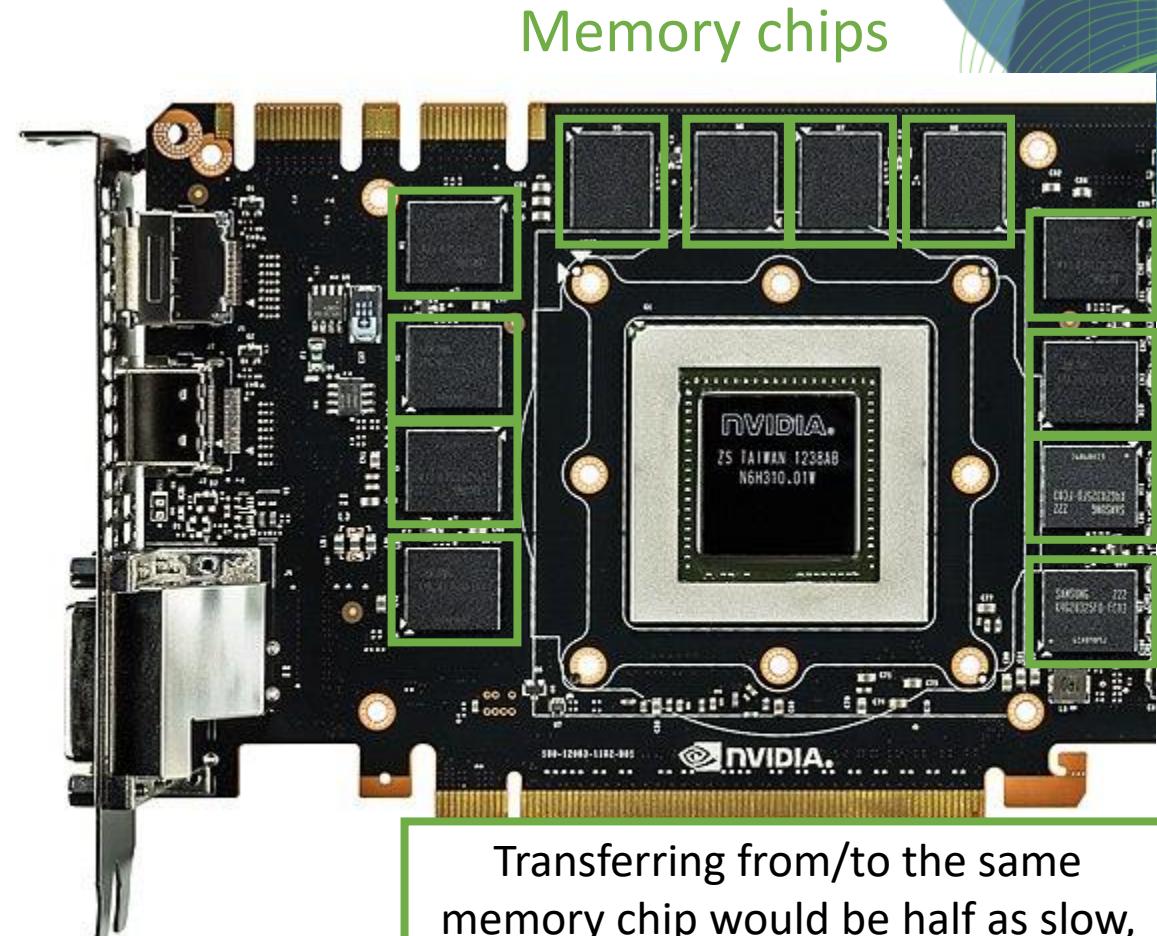
```
[2]: import pyclesperanto as cle

# List available devices
cle.available_device_names()

[2]: ['NVIDIA GeForce RTX 4070 Laptop GPU', 'gfx1103']

[3]: # select a specific device with only a part of its name
cle.select_device("TX")

[3]: (OpenCL) NVIDIA GeForce RTX 4070 Laptop GPU (OpenCL 3.0 CUDA)
      Vendor: NVIDIA Corporation
      Driver Version: 560.94
      Device Type: GPU
      Compute Units: 36
      Global Memory Size: 8187 MB
      Local Memory Size: 0 MB
      Maximum Buffer Size: 2046 MB
      Max Clock Frequency: 1230 MHz
      Image Support: Yes
```



Memory chips

# GPU-accelerated image processing

- Memory constraints

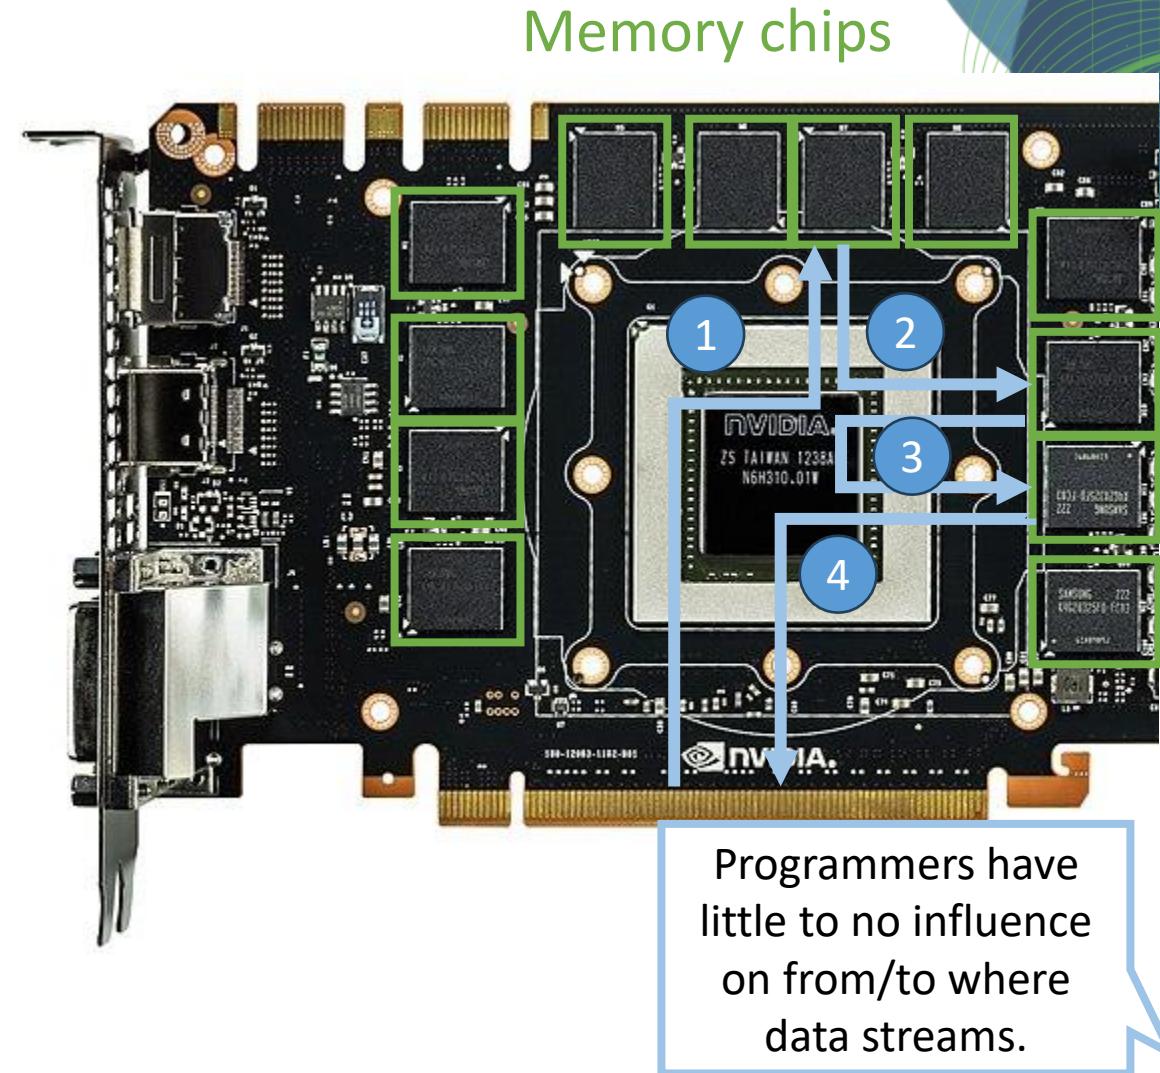
```
[2]: import pyclesperanto as cle

# List available devices
cle.available_device_names()

[2]: ['NVIDIA GeForce RTX 4070 Laptop GPU', 'gfx1103']

[3]: # select a specific device with only a part of its name
cle.select_device("TX")

[3]: (OpenCL) NVIDIA GeForce RTX 4070 Laptop GPU (OpenCL 3.0 CUDA)
      Vendor: NVIDIA Corporation
      Driver Version: 560.94
      Device Type: GPU
      Compute Units: 36
      Global Memory Size: 8187 MB
      Local Memory Size: 0 MB
      Maximum Buffer Size: 2046 MB
      Max Clock Frequency: 1230 MHz
      Image Support: Yes
```



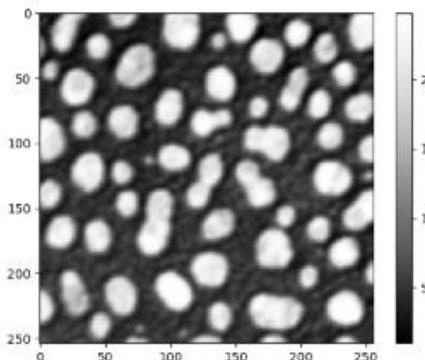
# GPU-accelerated image processing

- GPU-accelerated image processing also just uses functions.
- You used it in the recent weeks already.

```
[6]: # noise removal
```

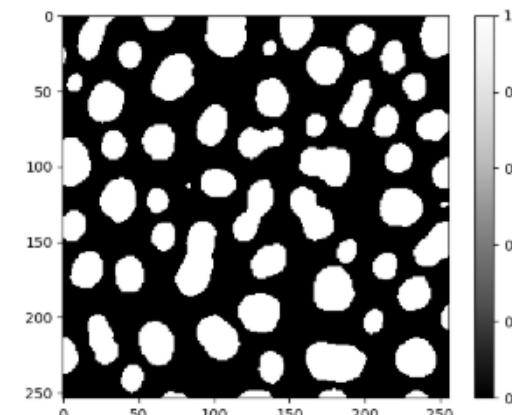
```
blurred = cle.gaussian_blur(image, sigma_x=1, sigma_y=1)  
blurred
```

```
[6]:
```



```
[7]:
```

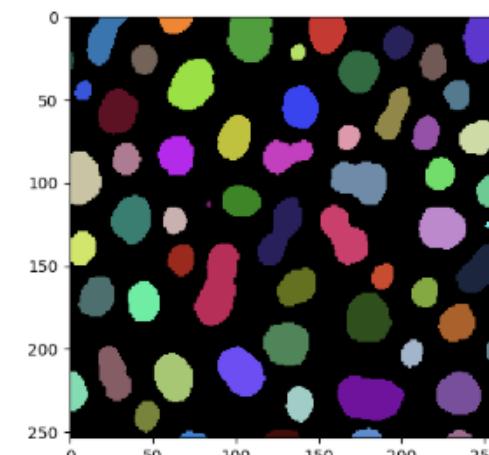
```
# binarization  
binary = cle.threshold_otsu(blurred)  
binary
```



```
[8]:
```

```
labels = cle.label(binary)  
labels
```

```
[8]:
```



**cle\_.image**

shape (254, 256)

dtype uint32

size 254.0 kB

min 0.0

max 62.0

# GPU-accelerated image processing

- Hint: Use SHIFT-Tab to discover functions in clesperanto. You can find them all in the main module. There are no sub-modules like in scikit-image.

```
[ ]: cle.gauss
f gauss_otsu_labeling          function
f gaussian_blur                function
f subtract_gaussian_background function
f divide_by_gaussian_background function
f difference_of_gaussian       function
```

```
[ ]: cle.otsu
f gauss_otsu_labeling          function
f eroded_otsu_labeling         function
f voronoi_otsu_labeling        function
f threshold_otsu               function
f remove_labels_with_map_values_out_of_range function
```

# Benchmarking

- The timeit library makes it straightforward to benchmark compute time in Python / Jupyter.

```
[7]: %timeit  
skimage.filters.gaussian(image, sigma=5, out=blurred_image, preserve_range=True)  
2.53 s ± 21.8 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)
```

```
[16]: %timeit  
cle.gaussian_blur(ocl_image, output_image=ocl_blurred, sigma_x=5, sigma_y=5, sigma_z=5)  
19.2 ms ± 465 µs per loop (mean ± std. dev. of 7 runs, 10 loops each)
```

Quiz: How much faster is cle compared to the scikit-image filter in this case?

# Benchmarking

- The timeit library makes it straightforward to benchmark compute time in Python / Jupyter.

```
[7]: %timeit
skimage.filters.gaussian(image, sigma=5, out=blurred_image, preserve_range=True)

2.53 s ± 21.8 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)
```

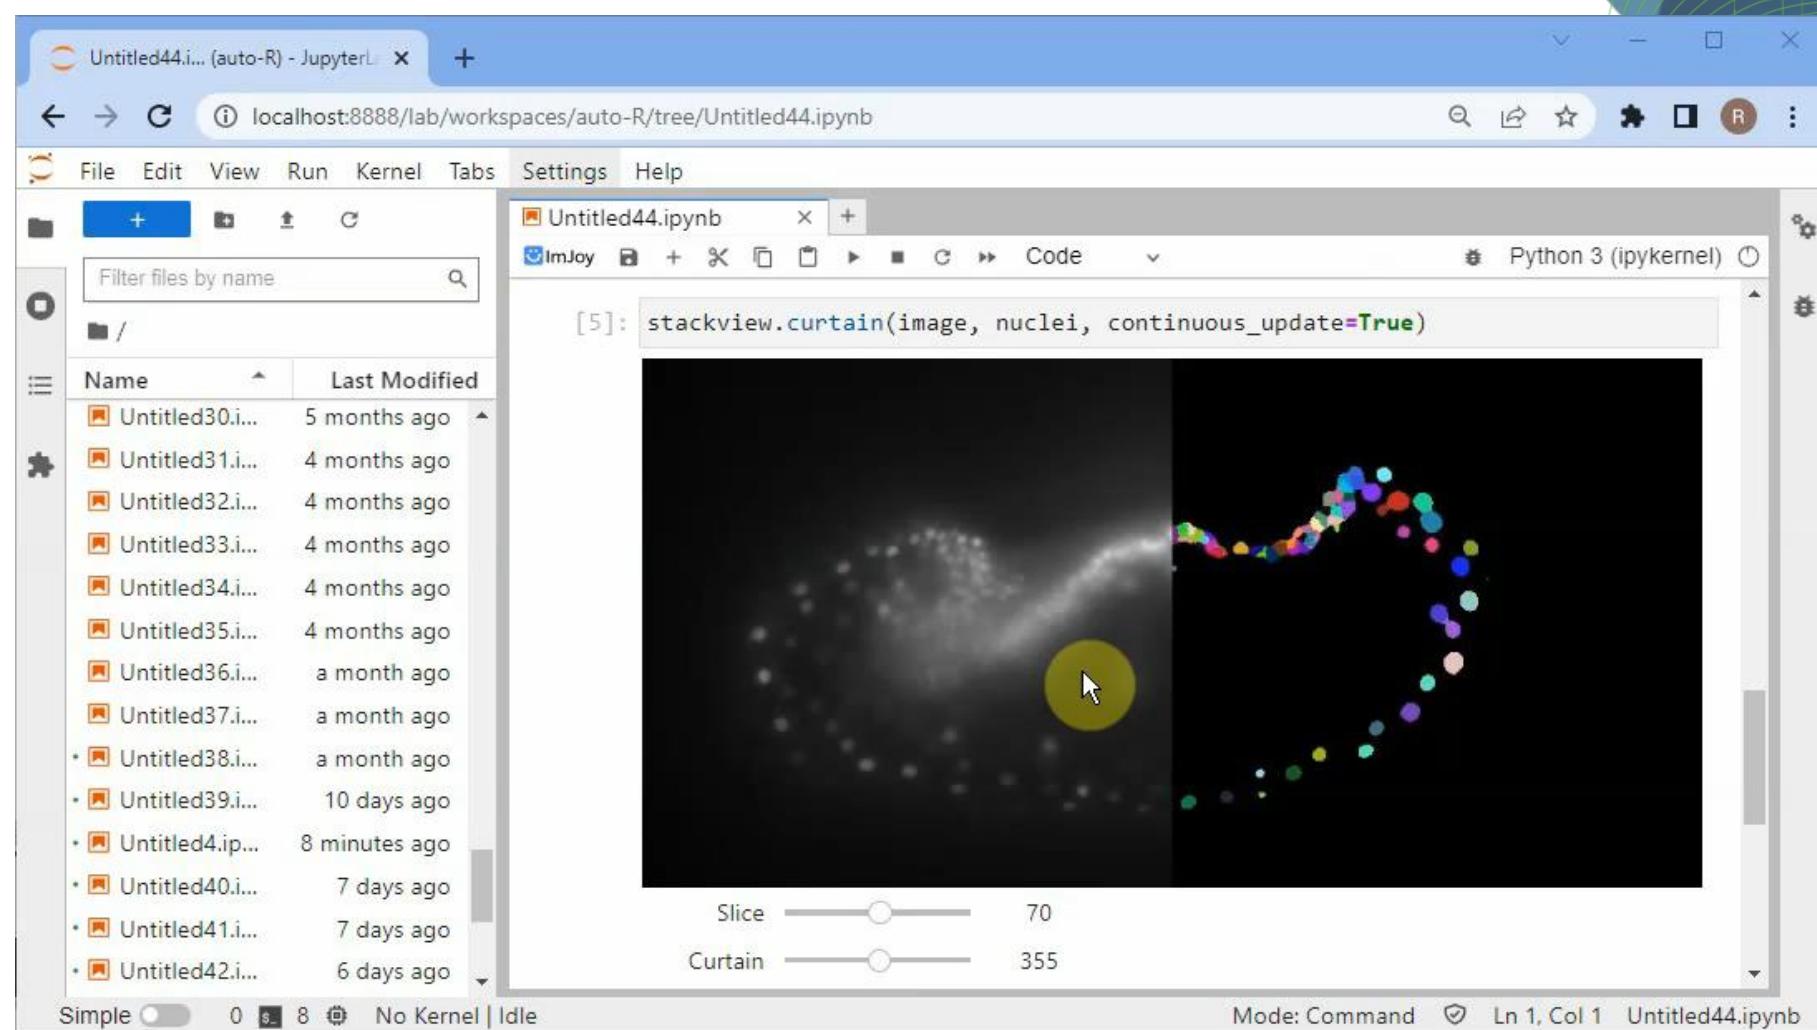
```
[16]: %timeit
cle.gaussian_blur(ocl_image, output_image=ocl_blurred, sigma_x=5, sigma_y=5, sigma_z=5)

19.2 ms ± 465 µs per loop (mean ± std. dev. of 7 runs, 10 loops each)
```

Quiz: Why do we pass the output-image here?

# 3D visualization on the cluster

- When working on the cluster / Jupyter Hub, consider using stackview instead of napari for inspecting images in 3D.





# GPU-accelerated Image Processing in Python: CUDA / cupy

Robert Haase

Funded by



Bundesministerium  
für Bildung  
und Forschung



Diese Maßnahme wird gefördert durch die Bundesregierung  
aufgrund eines Beschlusses des Deutschen Bundestages.  
Diese Maßnahme wird mitfinanziert durch Steuermittel auf  
der Grundlage des von den Abgeordneten des Sächsischen  
Landtags beschlossenen Haushaltes.

# cupy

- CUDA-based GPU-accelerated [image] data processing in Python

The image shows two side-by-side browser windows. The left window displays the official CuPy website at [cupy.dev](https://cupy.dev). It features a dark background with a green and white geometric cube logo and the word "CuPy". A green button labeled "NumPy/SciPy-compatible Array Library for GPU-accelerated Computing with Python" is highlighted with a red box. Below the button are three buttons: "GET STARTED" (red), "API REFERENCE" (teal), and "GITHUB" (teal). The right window shows the GitHub repository for CuPy at [github.com/cupy/cupy/](https://github.com/cupy/cupy/). The repository page includes a summary card with statistics: 451 issues, 63 pull requests, 128 watchers, 707 forks, and 7.1k stars. Below the card are sections for "main", "About", "Branches", and "Tags". A list of recent pull requests is shown, with the first one being a merge from user "takagi". The right sidebar contains a "cupy.dev" section with various tags like python, gpu, numpy, cuda, cublas, scipy, tensor, cudnn, etc.

# drop-in replacement for numpy and scipy

- The API of some cupy packages is close to the scipy/numpy API.
- This allows *easy* switching from scipy to cupy.

```
import scipy.ndimage as ndi
```

```
ndi.gaussian_filter(image, sigma=5)
```

```
import cupyx.scipy.ndimage as xdi
```

- However, image data still needs to be pushed to GPU memory.

```
xp_image = xp.asarray(image)
```

```
xdi.gaussian_filter(xp_image, sigma=5)
```

# Common patterns

- To make code independent from cupy availability, while minimizing if-else blocks, some common design patterns emerged:

```
try:  
    import cupy as xp  
except:  
    import numpy as xp  
  
import numpy as np
```

If this fails because cupy is not installed,  
This will execute and xp will be available.

We can still use np anyway.

The same pattern works with scipy.ndimage

```
try:  
    import cupyx.scipy.ndimage as xdi  
except:  
    import scipy.ndimage as xdi  
  
import scipy.ndimage as ndi
```

# Common patterns

- You can then call magic code like this, which will do different things depending on cupy-availability.

- If cupy is available:

```
[3]: image = imread("../data/blobs.tif")
```

```
xp_image = xp.asarray(image)
```

```
type(xp_image)
```

```
[3]: cupy.ndarray
```

- If cupy is not available:

```
[3]: image = imread("../data/blobs.tif")
```

```
xp_image = xp.asarray(image)
```

```
type(xp_image)
```

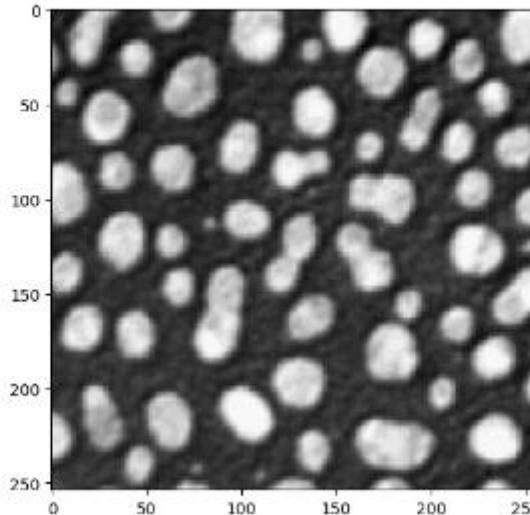
```
[3]: numpy.ndarray
```

# Common patterns

- Some if-else blocks are hard to avoid

```
[7]: if np == xp:  
    np_image = xp_image  
else:  
    np_image = xp.asarray(xp_image)  
  
imshow(np_image)
```

```
[7]: <matplotlib.image.AxesImage at  
0x24692ef85e0>
```

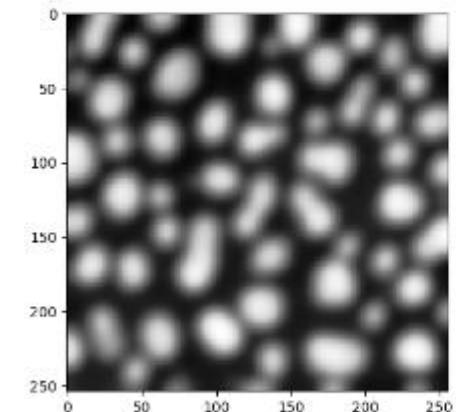


- Stackview aims to be cupy/numpy agnostic

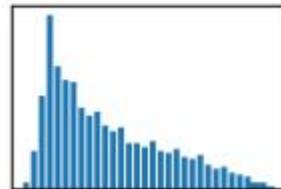
```
[4]: xp_blurred = xdi.gaussian_filter(xp_image, sigma=5)
```

```
[5]: stackview.insight(xp_blurred)
```

```
[5]:
```

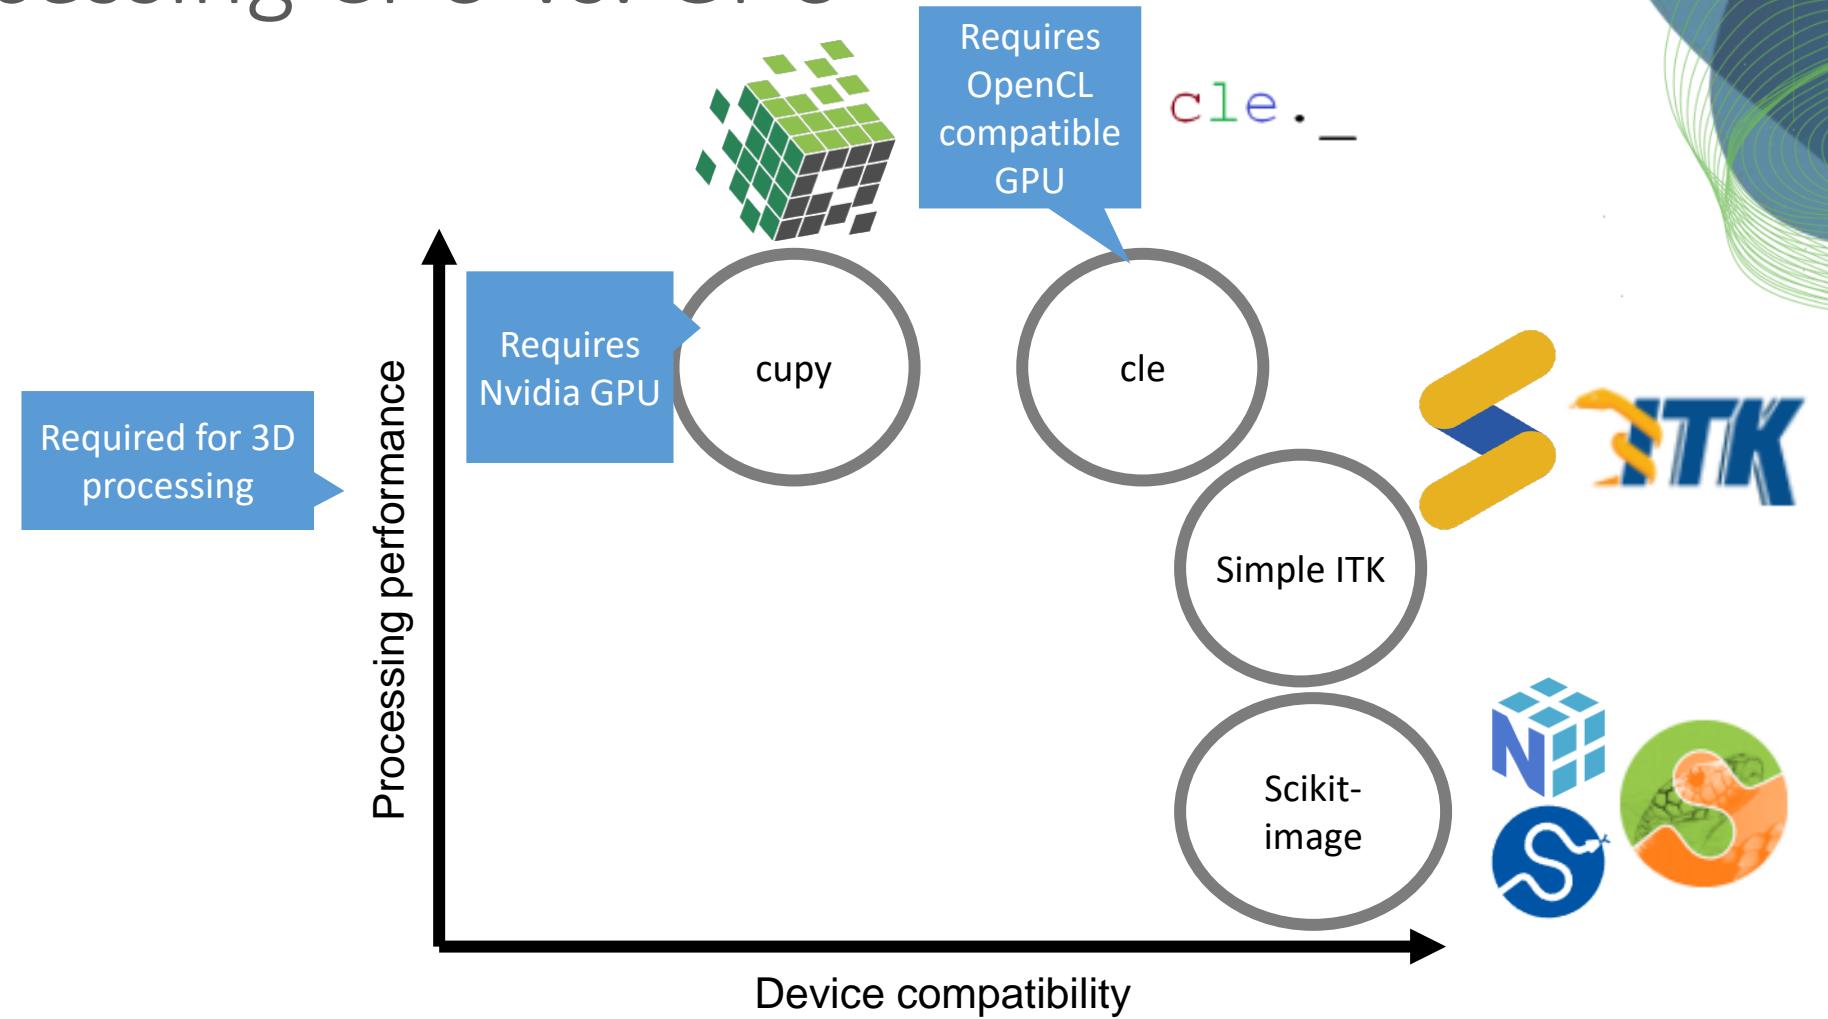


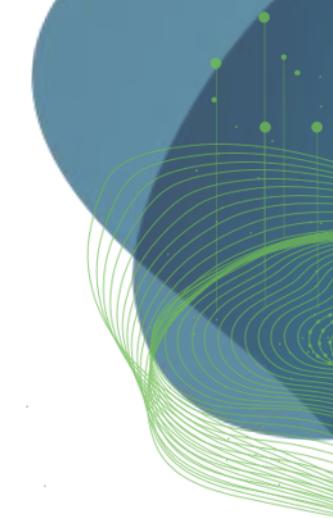
shape	(254, 256)
dtype	uint8
size	63.5 kB
min	35
max	237



# Image Processing CPU vs. GPU

- Performance versus compatibility





# Tiled image processing

## Robert Haase

Funded by



Bundesministerium  
für Bildung  
und Forschung

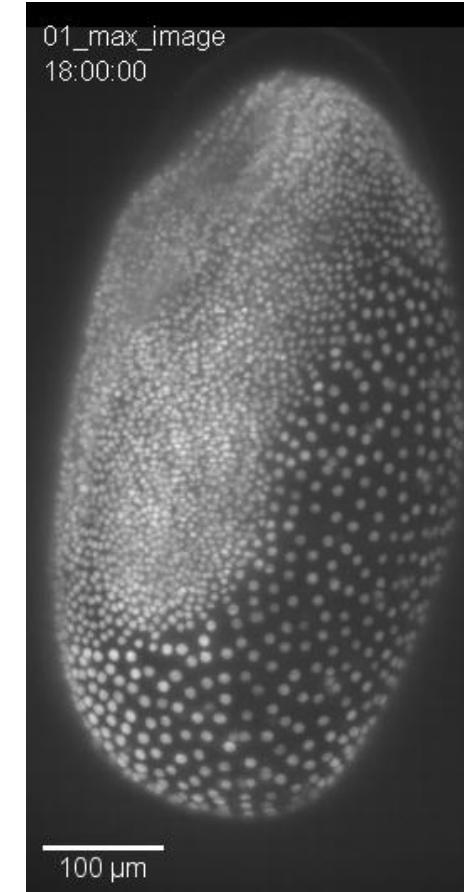


Diese Maßnahme wird gefördert durch die Bundesregierung  
aufgrund eines Beschlusses des Deutschen Bundestages.  
Diese Maßnahme wird mitfinanziert durch Steuermittel auf  
der Grundlage des von den Abgeordneten des Sächsischen  
Landtags beschlossenen Haushaltes.

# Quiz: Memory constraints

- Assume your computer runs out of memory while processing this image dataset.  
What can you do to avoid this?

- 1.
- 2.
- 3.
- 4.
5. Tiled image processing



The full data set is about  
1024 (width)  
2048 (height)  
100 (depth)  
3700 (frames)  
large

# Optimal performance through smart memory management

- The classical way of dealing with large image stacks...

Load data → Preprocessing

Load data

Load data

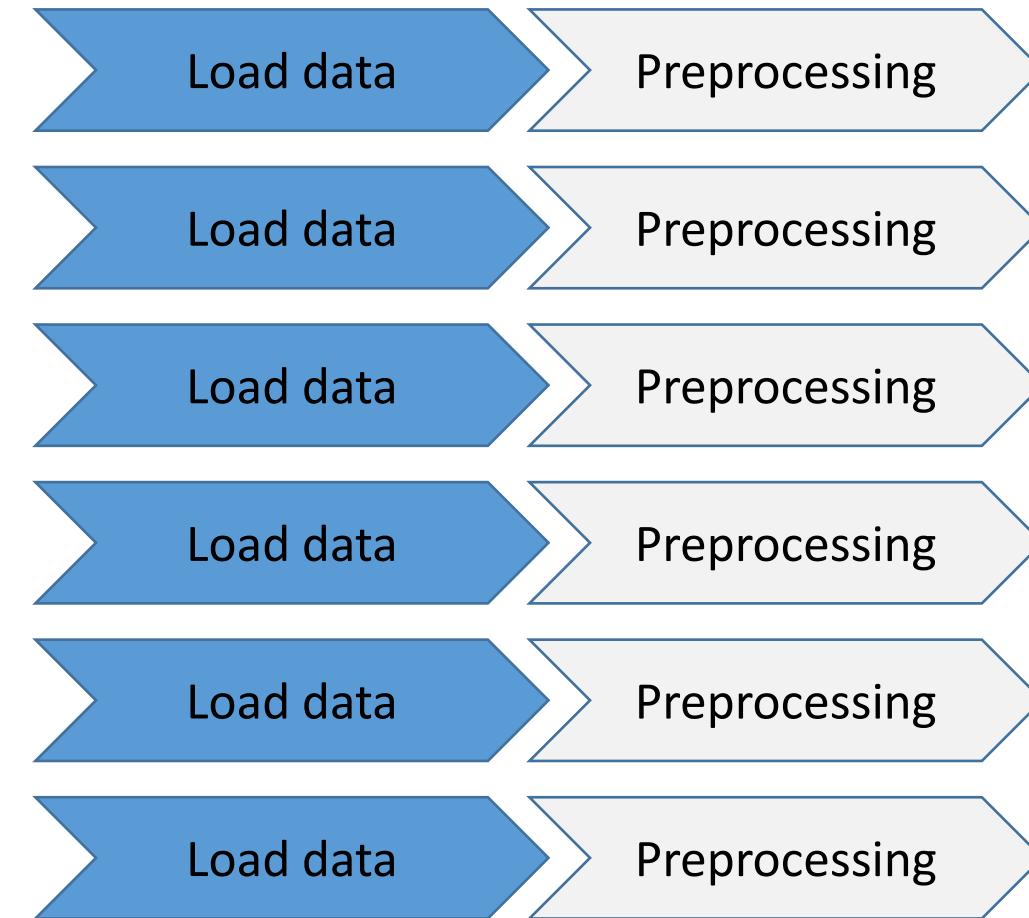
Load data

Load data

Load data

# Optimal performance through smart memory management

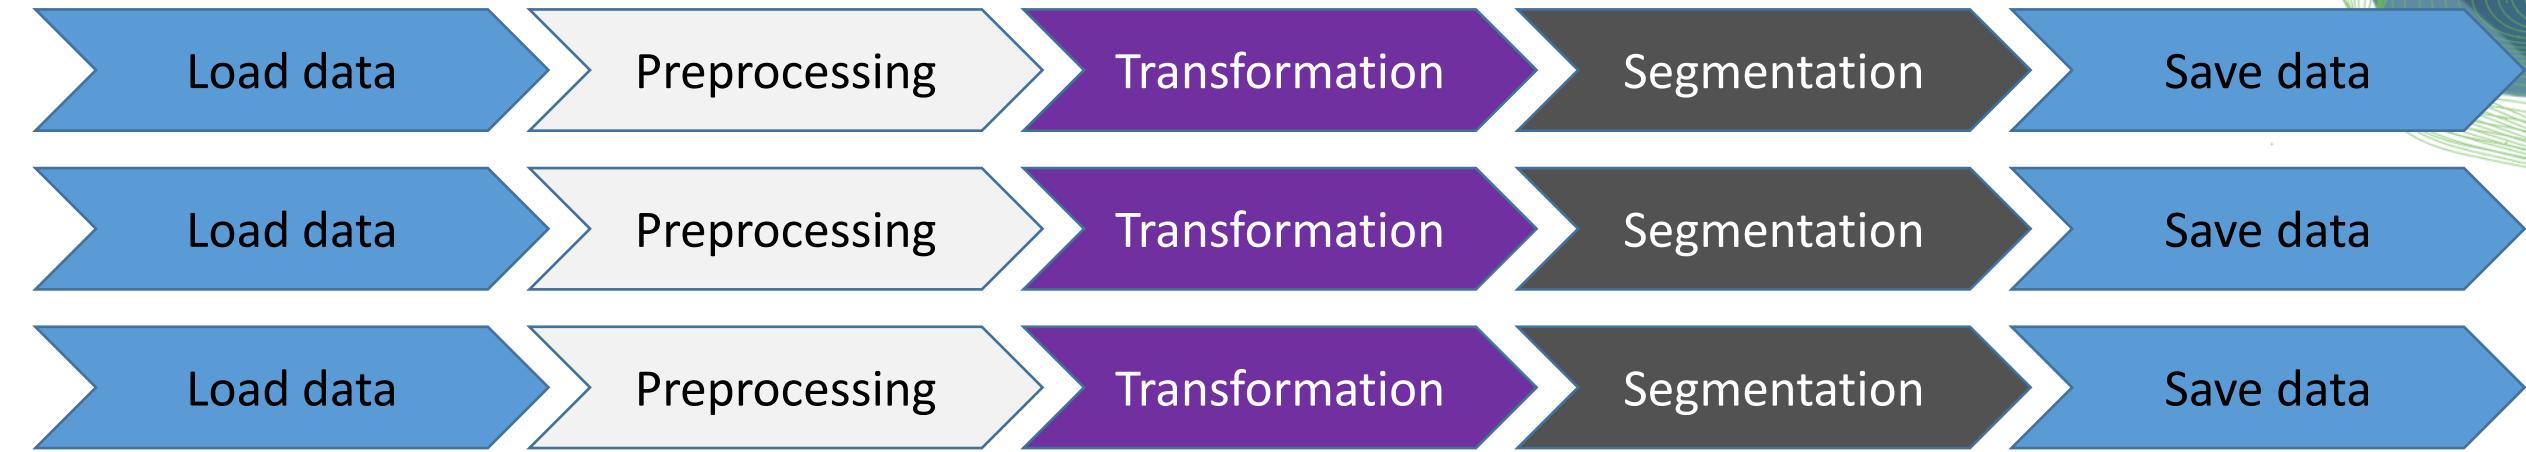
- The classical way of dealing with large image stacks... is suboptimal



This strategy does not just take long; it also costs a lot of memory!

# Optimal performance through smart memory management

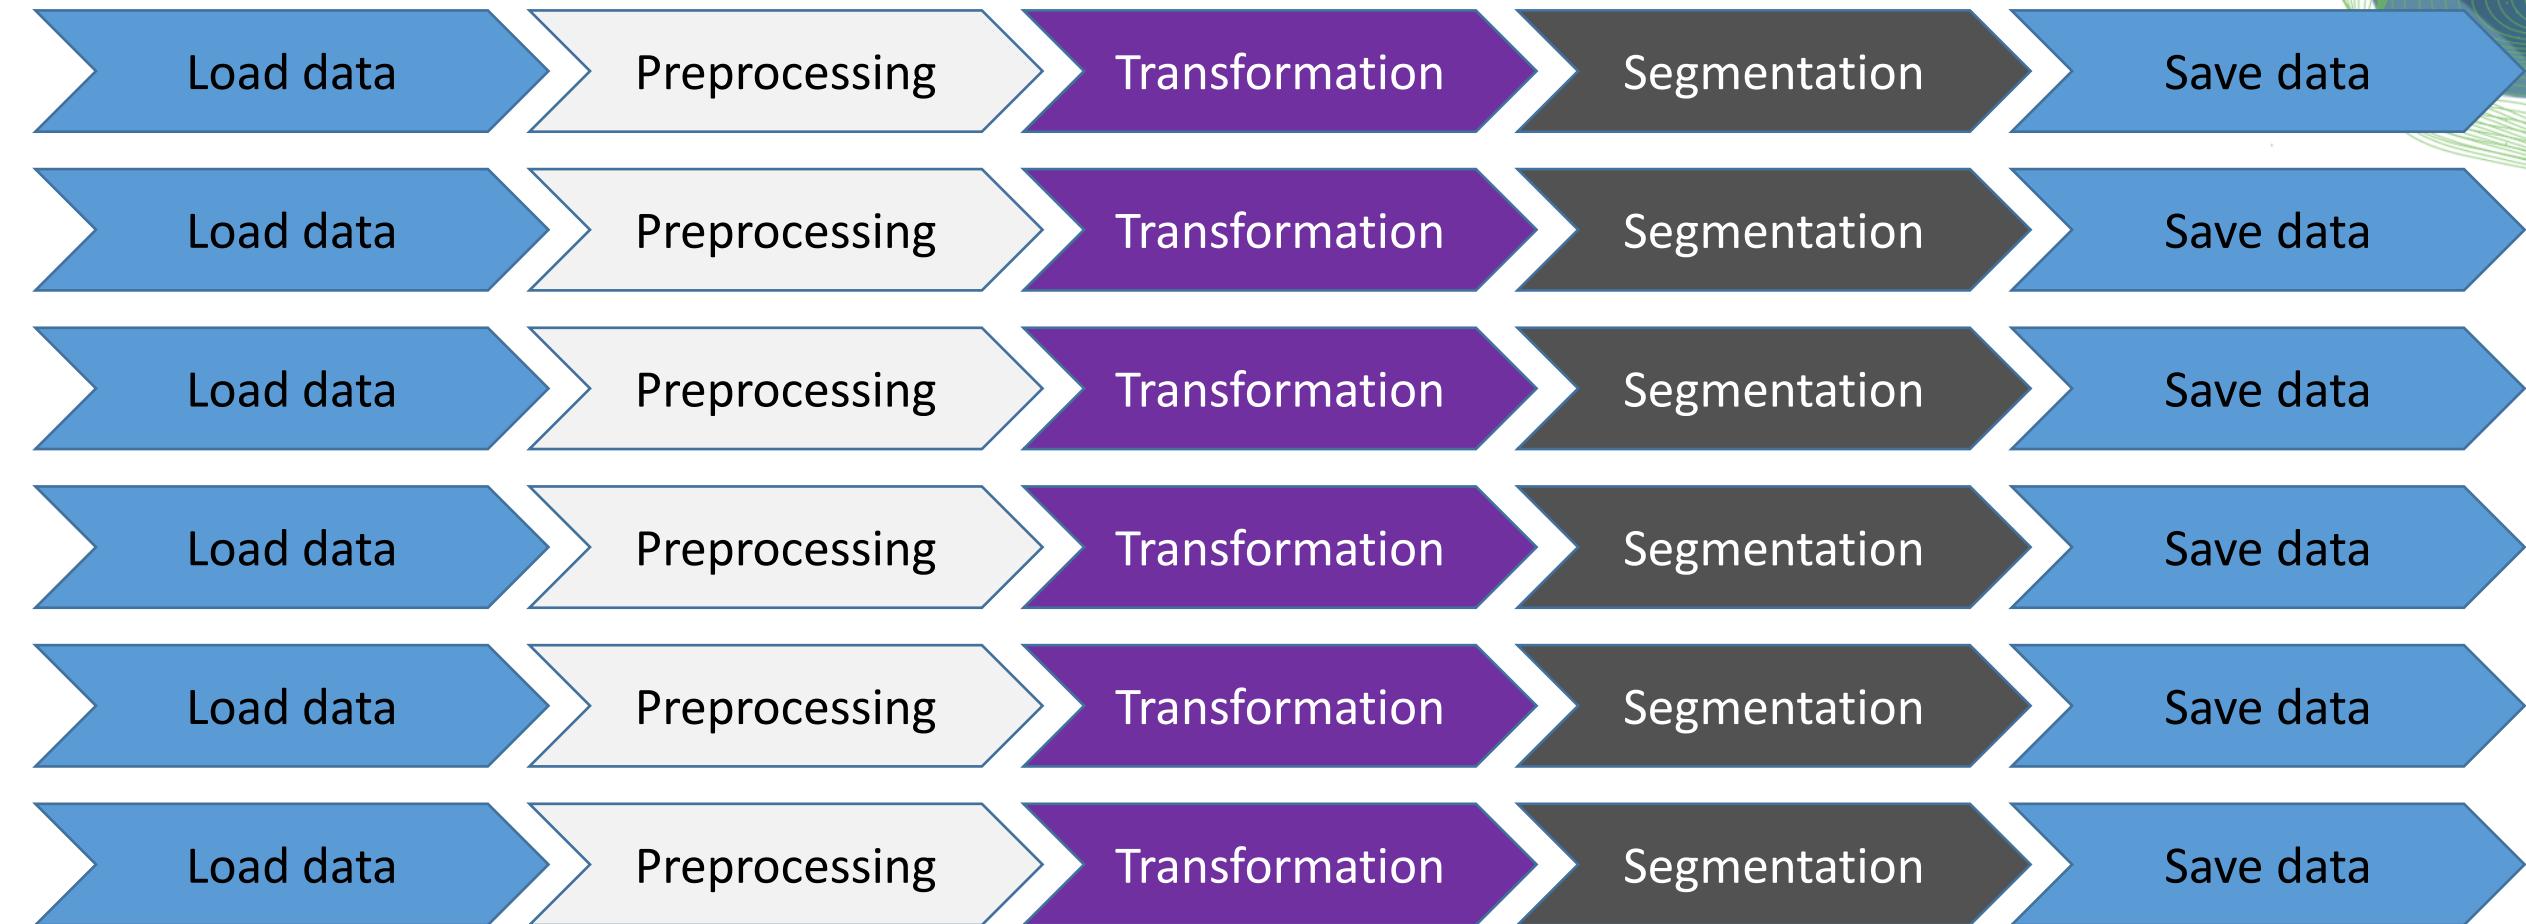
- Processing time-point by time-point is more efficient!



This strategy also works  
tile-by-tile on large 3D  
stacks and timelapse data!

# Optimal performance through smart memory management

- Even better: Distribute tasks between parallelized computation systems

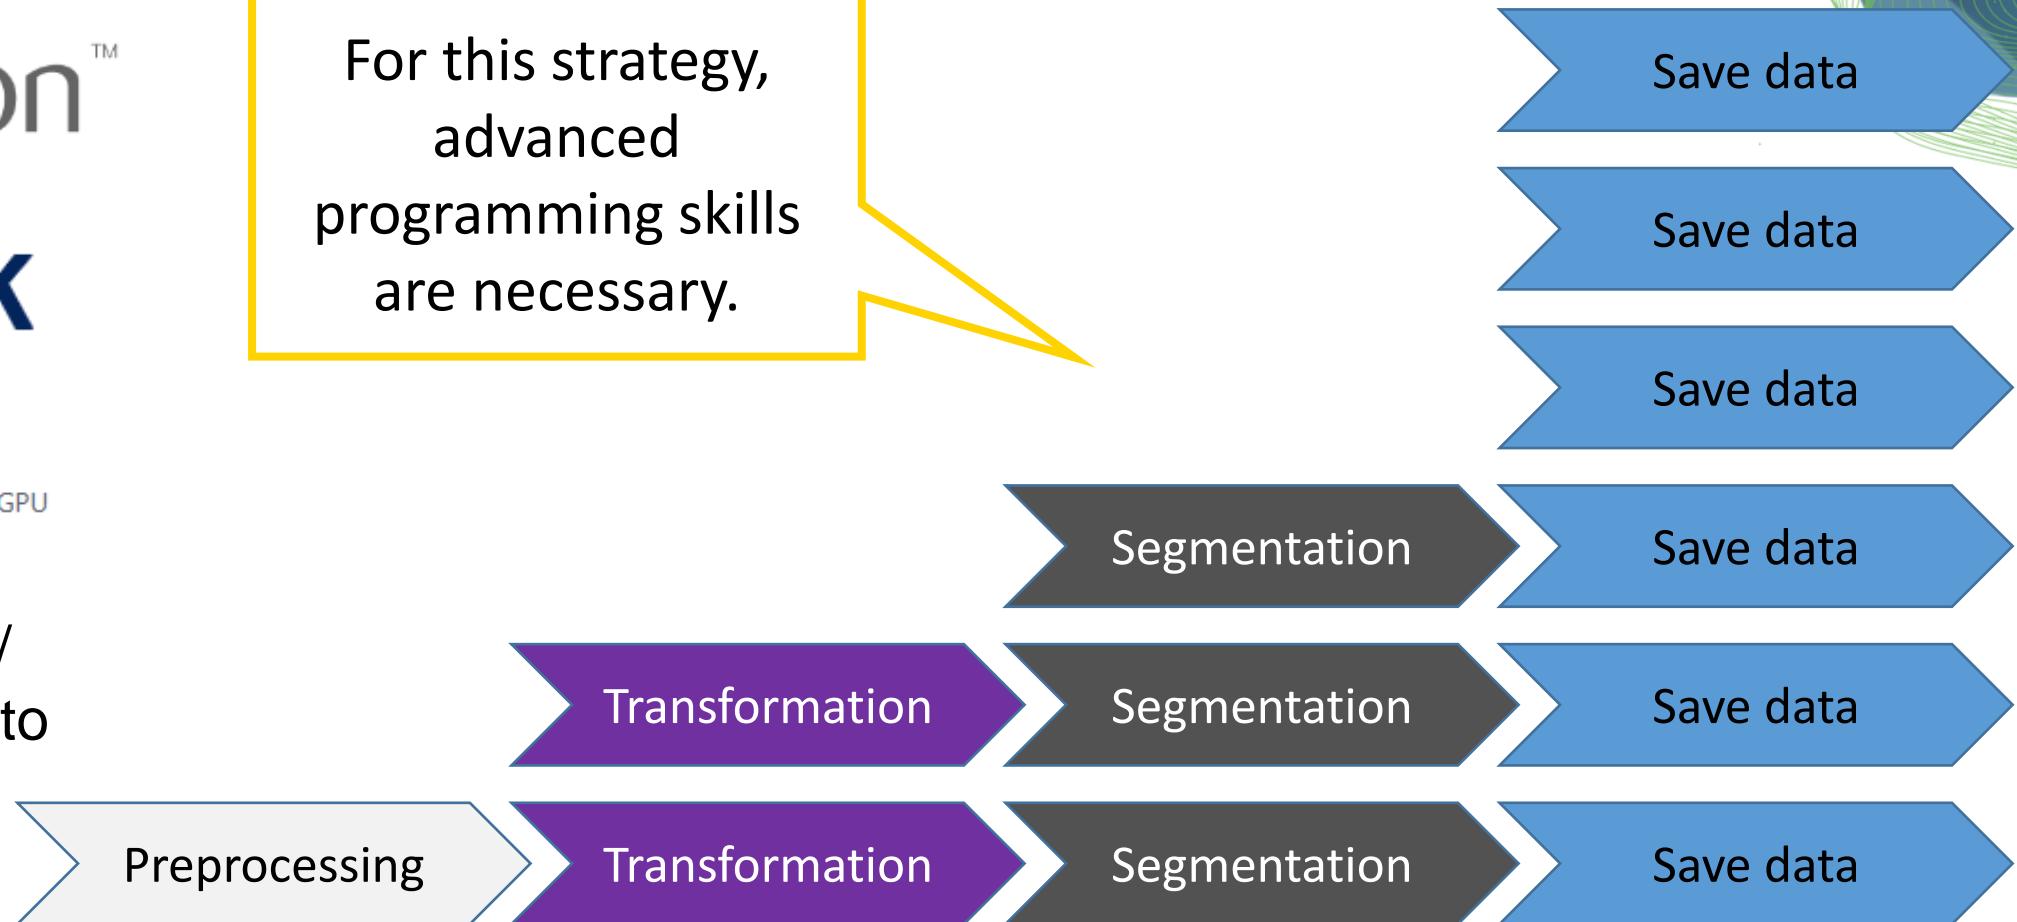


# Optimal performance through smart memory management

- Even better: Distribute tasks between parallelized computation systems

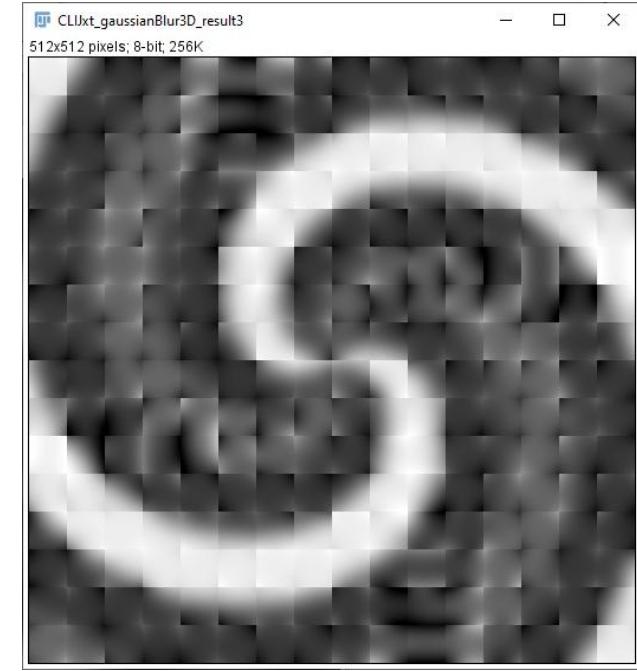
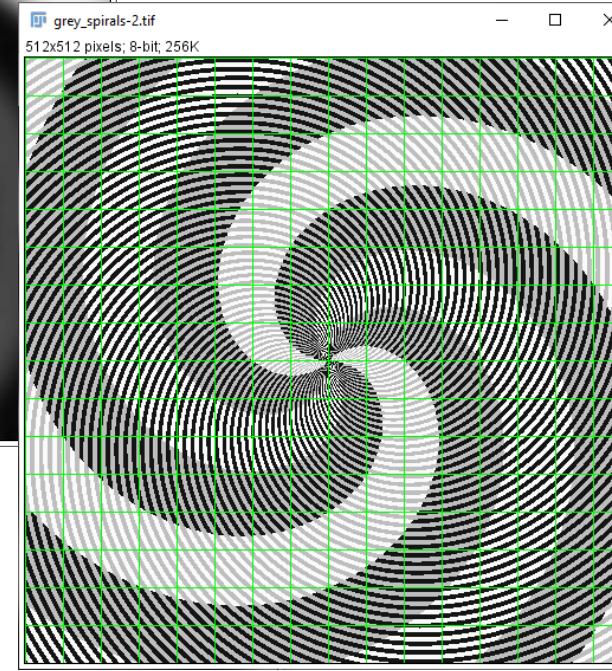
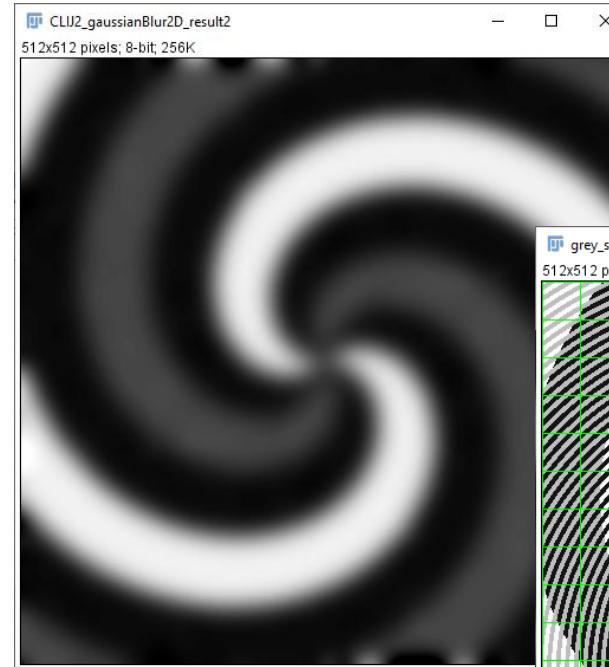
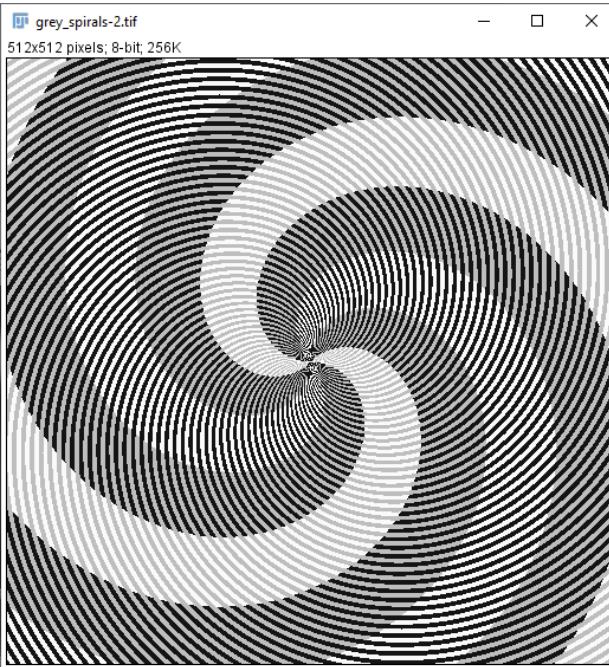


For this strategy,  
advanced  
programming skills  
are necessary.



# Tiling

- The **last** perimeter against big data



If the image is too large for the computer memory, image processing as a whole is *not possible*.

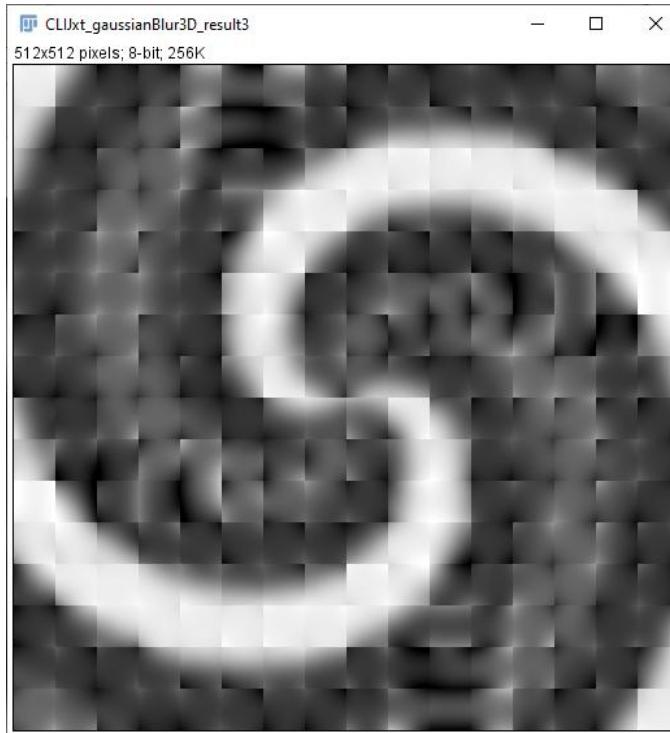
Processing tile-by-tile poses new challenges

# Tiling

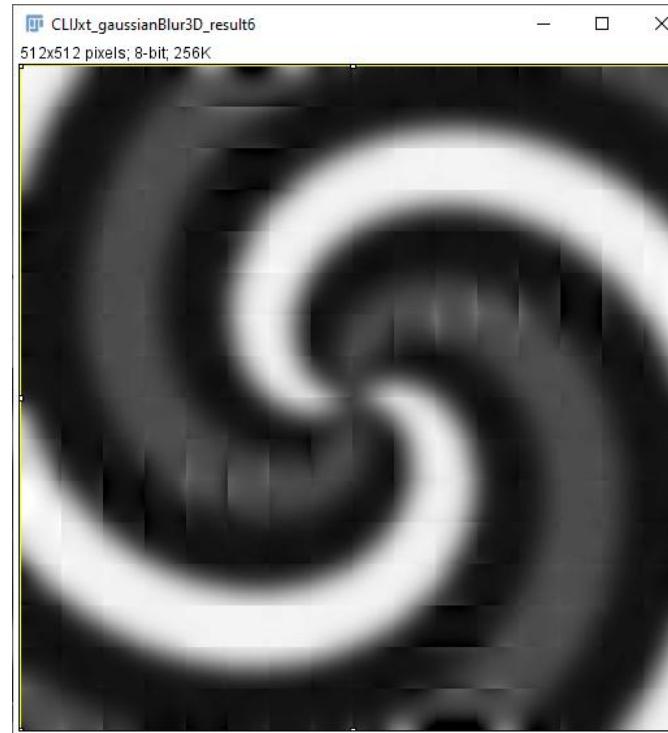
- Example: Gaussian blur (sigma = 20)
- Solution: Process with overlapping tiles (size + margin)

Optimal margin size depends  
on algorithm and its  
parameters

Margin: 0 pixels



Margin: 10 pixels



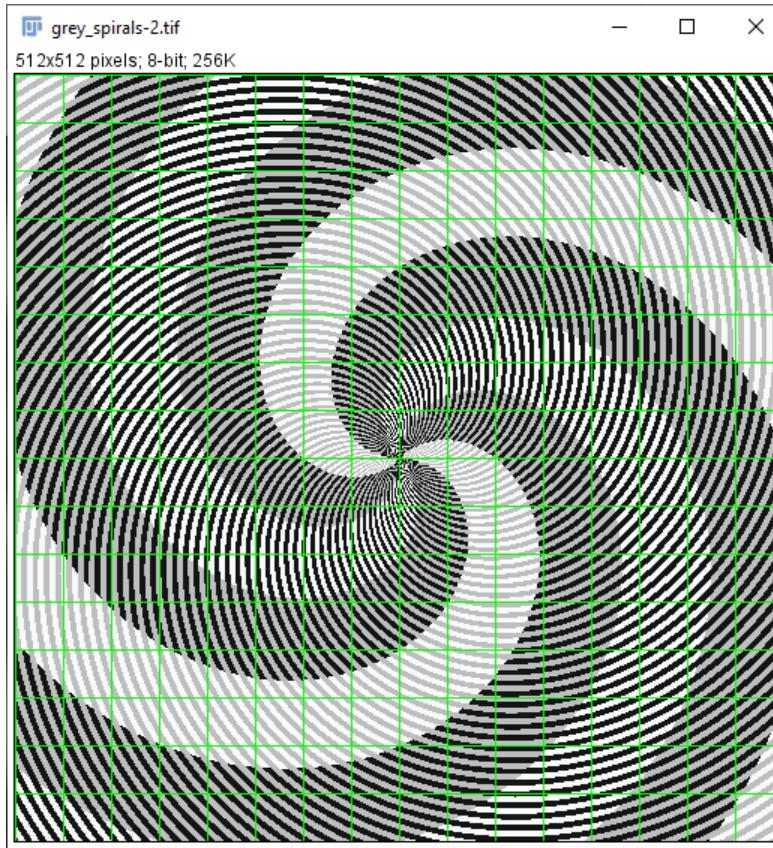
Margin: 20 pixels



# Tiling

- Example: Gaussian blur (sigma = 20 pixels)
- Solution: Process with overlapping tiles (size + margin)

Computation time depends  
on tile size and margin width



Margin: 20 pixels  
Size: 5x original

Margin: 10 pixels  
Size: 2.7x original

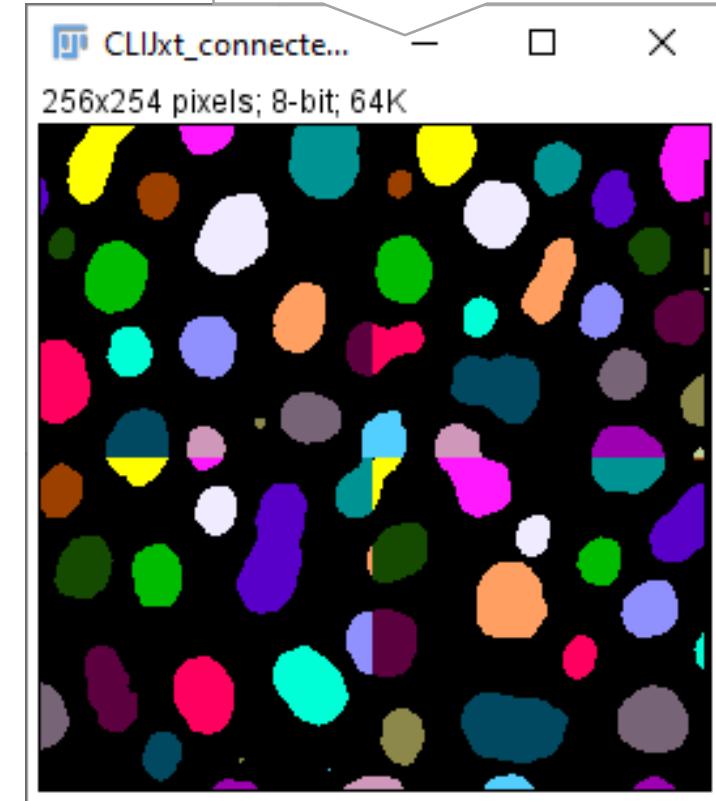
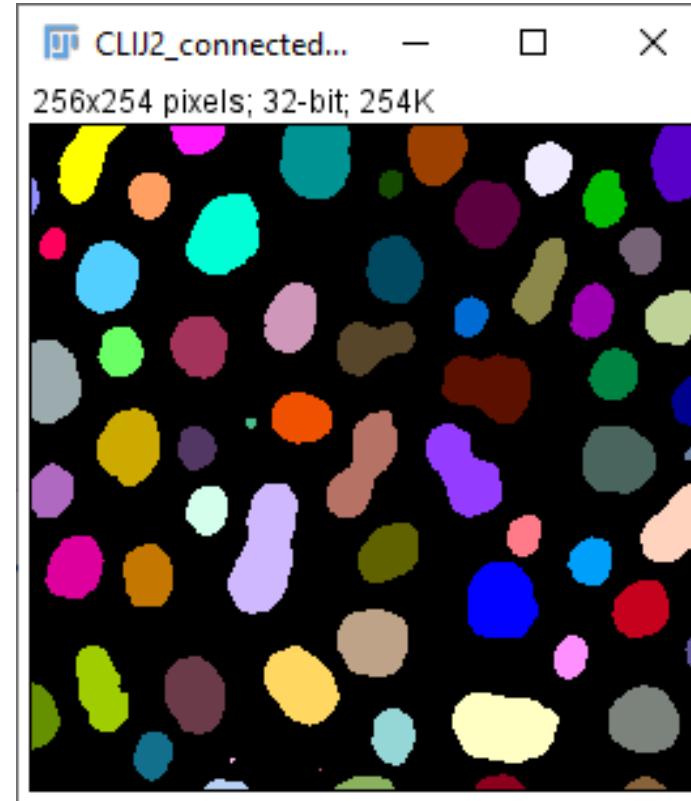
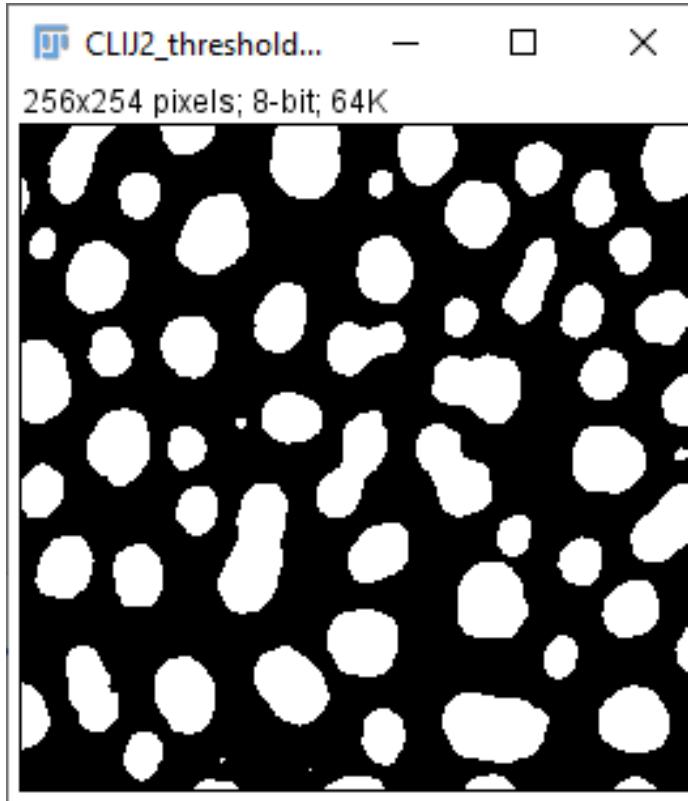
Tile  
32x32 pixels

52x52 pixels

72x72 pixels

# Tiling

- Some algorithms are hard to solve by processing tiles
- Example: Connected component analysis



Checking which labels touch and combine them is feasible.

# Tiling

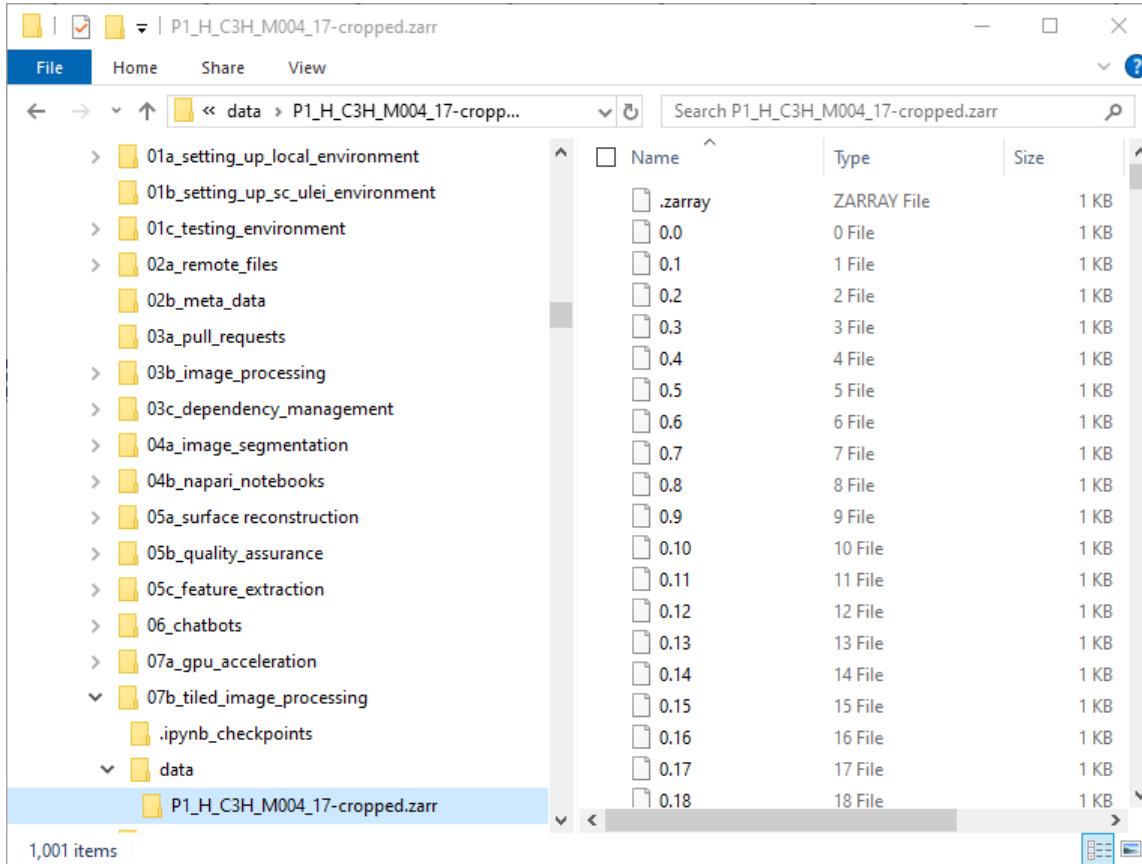
- Some algorithms are hard to solve by processing tiles
- Example: Connected component analysis

There are algorithms for that, but hardly available tools.



# Tiled image processing in Python

- Key: tiled file formats, for parallel, distributed, lazy loading



After executing this,  
no pixel has been  
read yet.

```
zarr_image = da.from_zarr(zarr_filename)  
zarr_image
```

	Array	Chunk
Bytes	9.54 MiB	9.77 kB
Shape	(2000, 5000)	(100, 100)
Dask graph	1000 chunks in 2 graph layers	
Data type	uint8 numpy.ndarray	

5000 2000

# Tiled image processing in Python

- Lazy processing

```
[5]: tile_map = da.map_blocks(count_nuclei, zarr_image)  
tile_map
```

```
Processing image of size (0, 0)  
(1, 1)  
Processing image of size (1, 1)  
(1, 1)
```

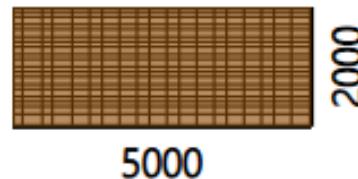
```
[5]:
```

	Array	Chunk
Bytes	76.29 MiB	78.12 kB
Shape	(2000, 5000)	(100, 100)

Dask graph 1000 chunks in 3 graph layers

Data type float64 numpy.ndarray

After executing this, no pixel has been read or processed yet.



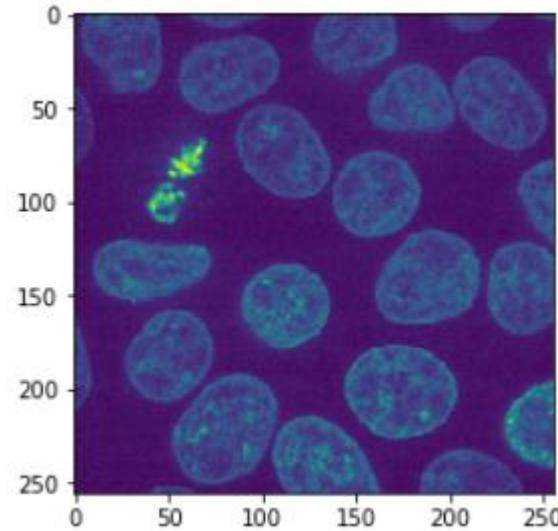
After that, results are available

```
result = tile_map.compute()
```

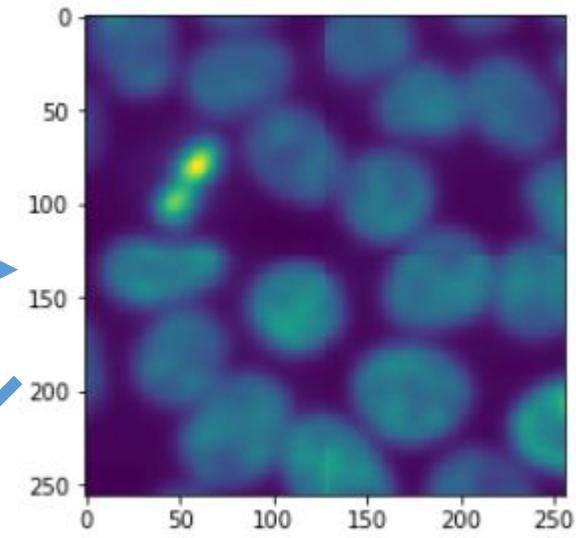
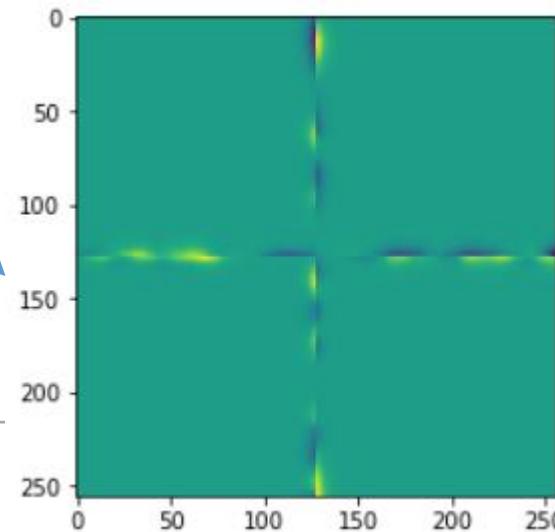
```
Processing image of size (100, 100)  
Processing image of size (100, 100)
```

# Tiling with/out overlap

- Processing of images in tiles: artifacts ad tile borders

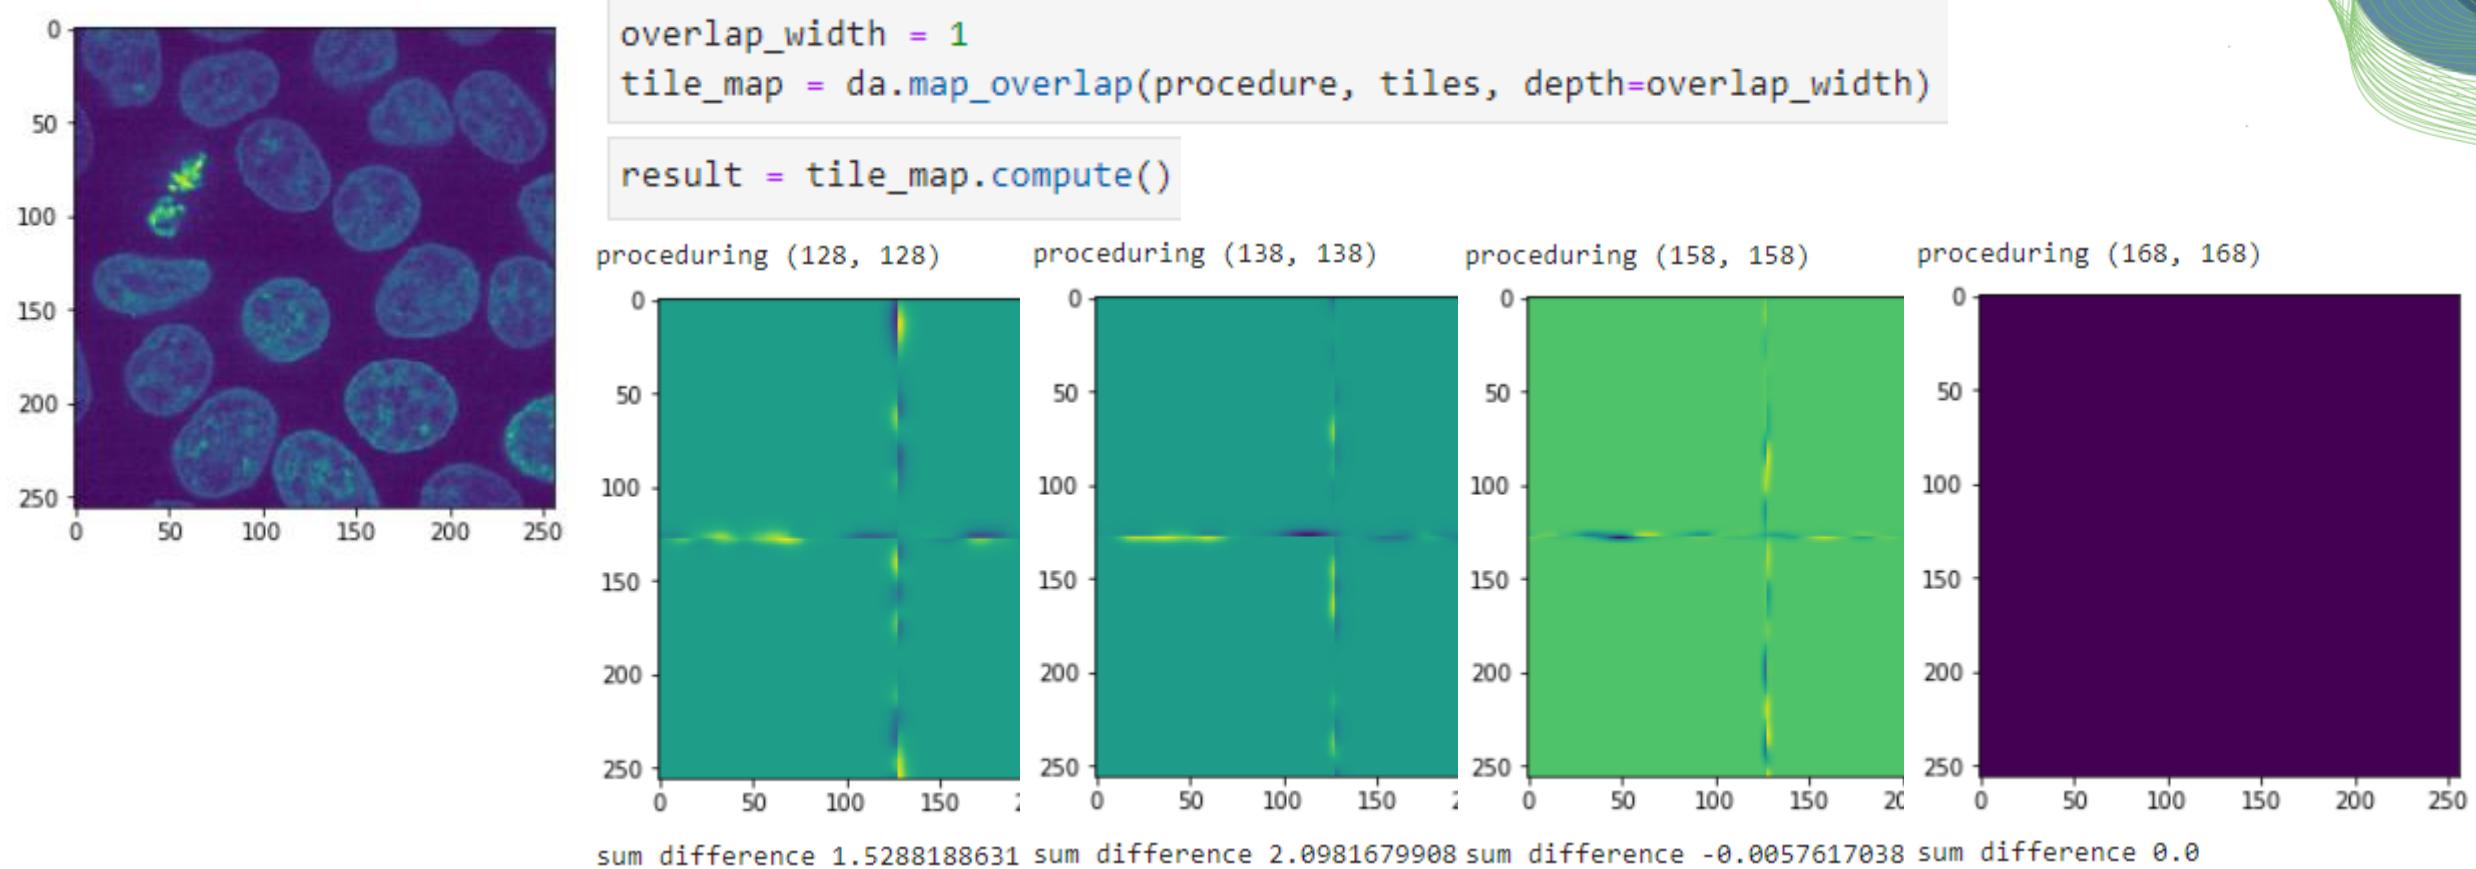


```
tile_map = da.map_blocks(procedure, tiles)  
  
result = tile_map.compute()
```



# Tiling with/out overlap

- Processing of images in tiles: artifacts at tile borders





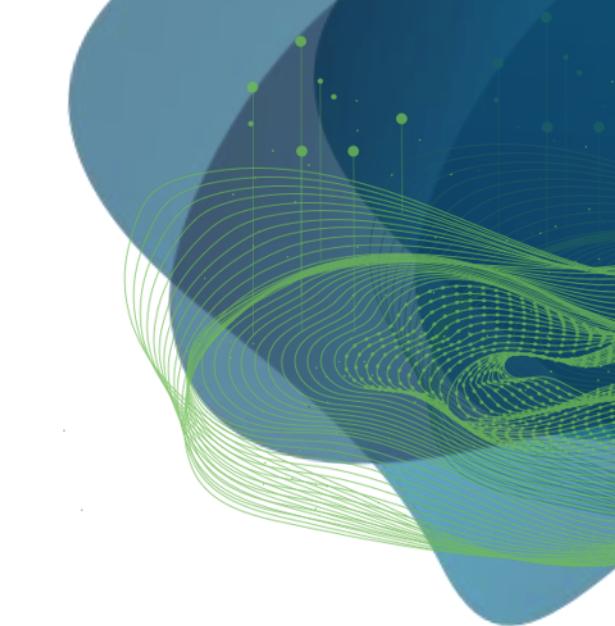
DRESDEN LEIPZIG

CENTER FOR SCALABLE DATA ANALYTICS  
AND ARTIFICIAL INTELLIGENCE

# Segmentation quality estimation

Robert Haase

Reusing materials from Lena Maier-Hein, Annika Reinke (DKFZ) et al.  
and Martin Schätz (Charles Uni Prague)



GEFÖRDERT VOM



Bundesministerium  
für Bildung  
und Forschung



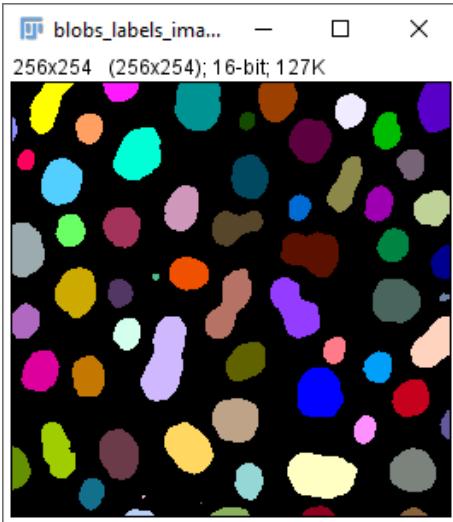
Diese Maßnahme wird gefördert durch die Bundesregierung  
aufgrund eines Beschlusses des Deutschen Bundestages.  
Diese Maßnahme wird mitfinanziert durch Steuermittel auf  
der Grundlage des von den Abgeordneten des Sächsischen  
Landtags beschlossenen Haushaltes.

# Goal

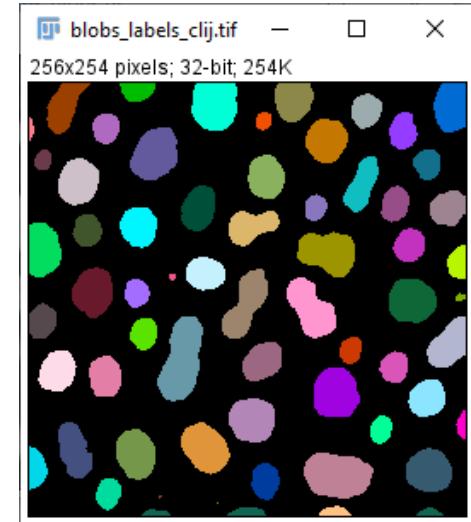
- Compare label images quantitatively, to know
  - how “good” a segmentation algorithms is and/or
  - how “variable” segmentations (from humans or computers) are

How can we know if  
these results are the  
same?

Human  
annotation  
("ground truth")



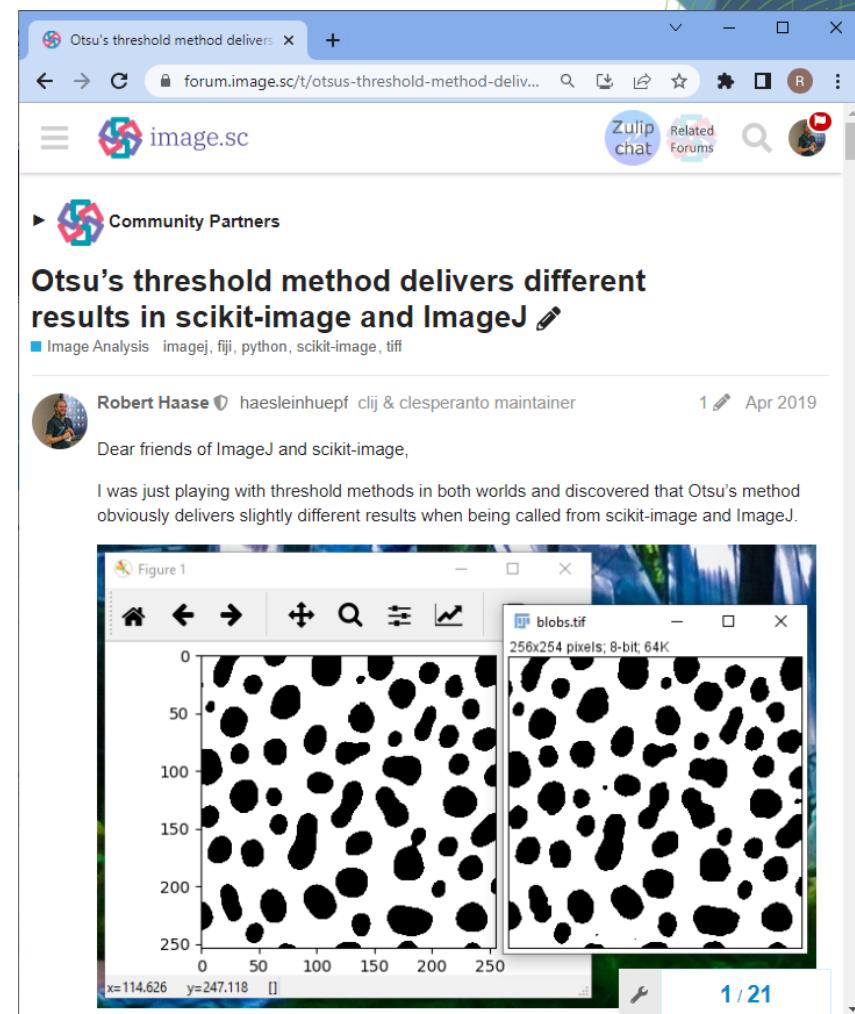
Algorithm result



# Why do results vary?

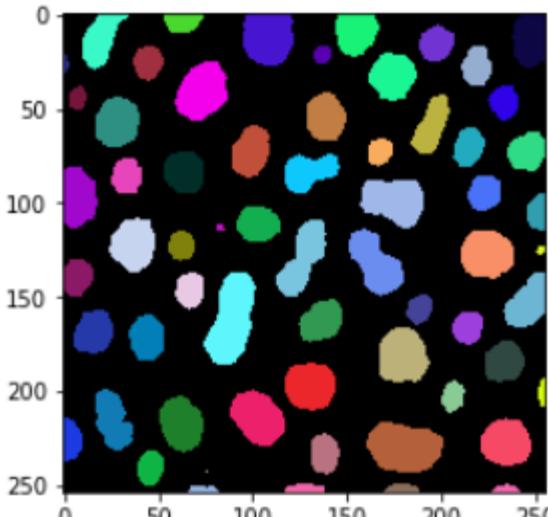
Potential reasons of same workflows delivering different results:

- Image data type (8/16/32-bit float/int)
- Workflow implementation
- How histograms are determined
- How the threshold is determined from the histogram
- Compute architecture (CPU, GPU, TPU, ...)
- Hardware vendor
- Software / driver versions

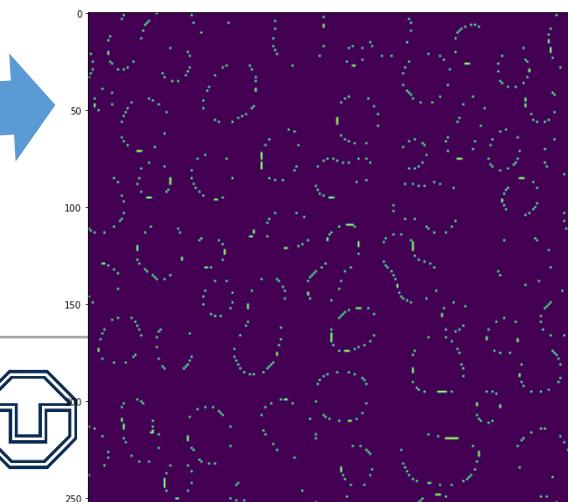
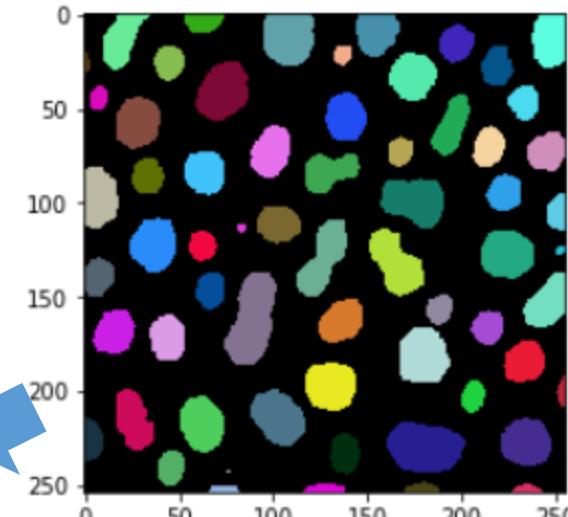
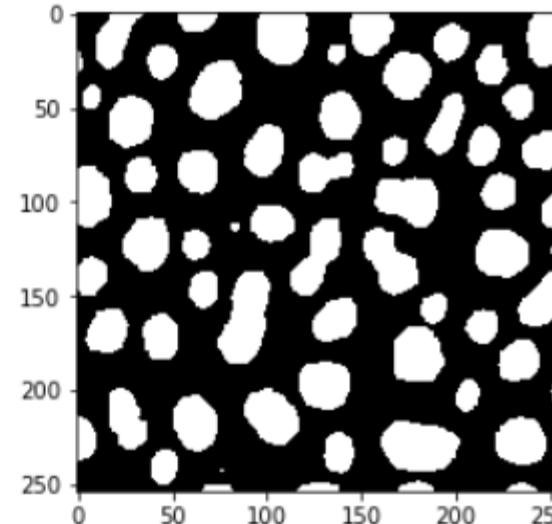
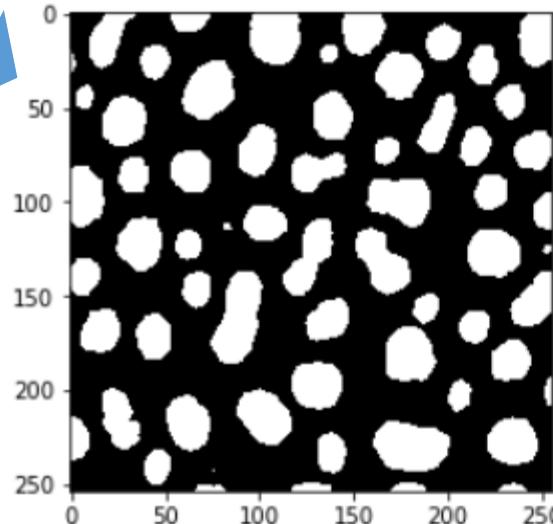


# Visual comparison

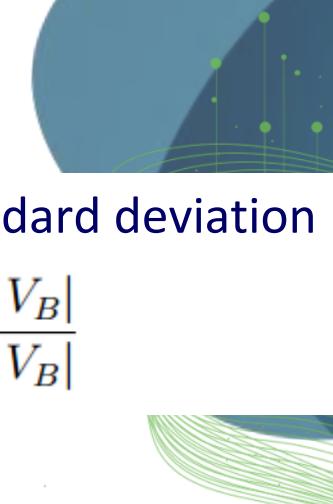
- The order in label images may be different. To compare them visually, we need to turn them into binary images first.



```
binary_blobs_imagej = imread(filenames[0]) > 0  
binary_blobs_skimage = imread(filenames[1]) > 0  
imshow(binary_blobs_imagej)  
imshow(binary_blobs_skimage)
```



# Quantitative comparison



- Voxel-wise Youden-Index

$$YI = p_{TP} + p_{TN} - 1$$

- Volume error

$$\Delta_V = V_A - V_B$$

$$\delta_V = \frac{\Delta_V}{V_B}$$

- Dice Index

$$DI(A, B) = \frac{2 |A \cap B|}{|A| + |B|}$$

- Jaccard Index

$$JI(A, B) = \frac{|A \cap B|}{|A \cup B|} = \frac{DI}{2 - DI}$$

- Contour distance

$$d_{e,min}(a, B) = \min(d_e(a, b) | b \in B)$$

$$\bar{d}_c(A, B) = \frac{\sum_{\forall a \in C(A)} d_{e,min}(a, C(B))}{|C(A)|}$$

$$\bar{d}_{bil,c}(A, B) = \frac{\bar{d}_c(A, B) + \bar{d}_c(B, A)}{2}$$

- Hausdorff distance

$$d_H(A, B) = \max(d_{e,min}(a, B) | a \in A)$$

$$d_{bil,H}(A, B) = \max(d_H(A, B), d_H(B, A))$$

- Simplified Hausdorff distance

$$d_H(A, B) = \max(d_{e,min}(a, C(B)) | a \in C(A))$$

- Volume standard deviation

$$\delta_{\bar{V}} = 2 \frac{|V_A - V_B|}{|V_A + V_B|}$$

- Classification error

$$e_{Class} = \frac{H}{|TP| + |FN|}$$

- Hamming distance

$$d_h = |A \cup B| - |A \cap B| \\ = |FP| + |FN|$$

# Choosing the right metric is key

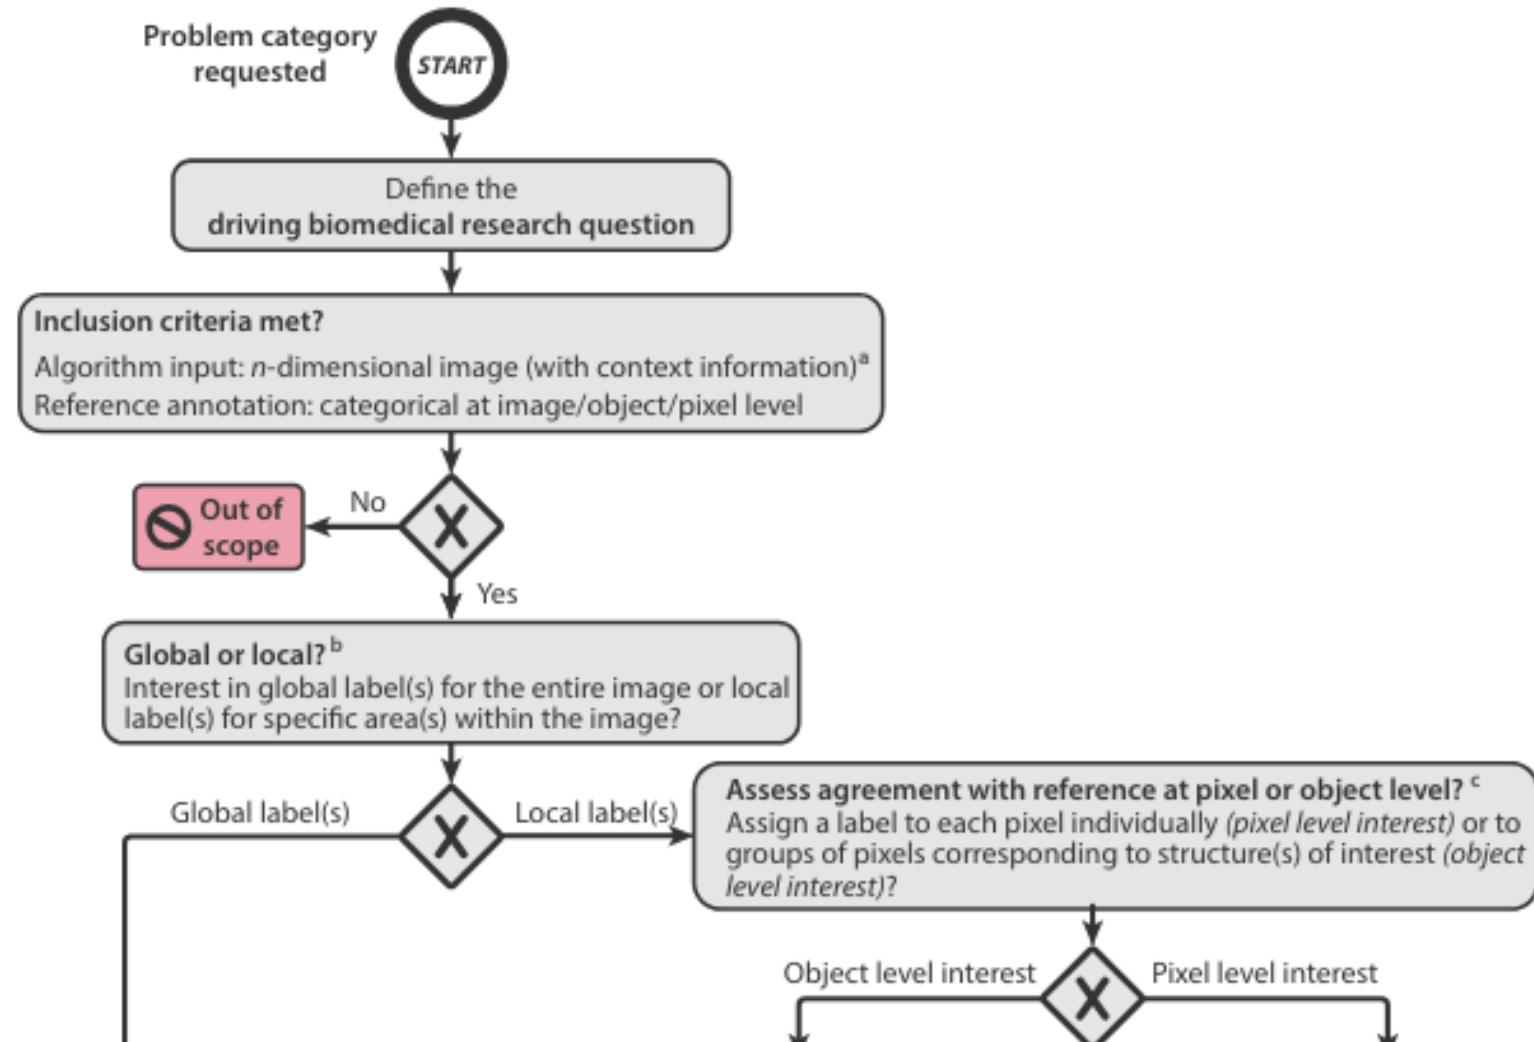
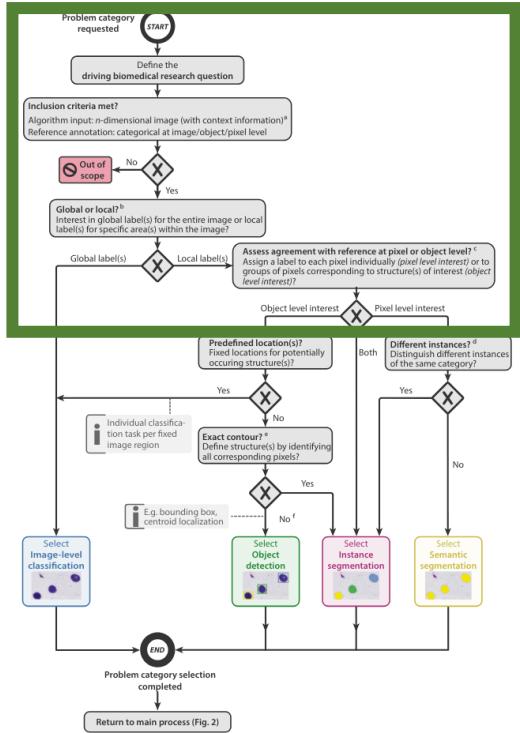
- Systematic overview by Maier-Hein, Reinke et al.

The screenshot shows a browser window for arXiv.org. The URL is arxiv.org/abs/2206.01653. The page title is "Metrics reloaded: Recommendations for image analysis validation". It's a Computer Science paper in Computer Vision and Pattern Recognition. The abstract discusses flaws in machine learning (ML) algorithm validation. The page includes a sidebar with download options (PDF, HTML, Tex Source, Other Formats), a CC-BY license, and citation links to NASAADS, Google Scholar, and Semantic Scholar.

The screenshot shows a browser window for nature.com. The URL is nature.com/articles/s41592-023-02151-z. The page title is "Metrics reloaded: recommendations for image analysis validation". It's published in Nature Methods, Perspectives section, on 12 February 2024. The page includes a sidebar with a download link ("Download PDF"), associated content ("Associated content" with links to "The future of bioimage analysis" and "Understanding metric-related pitfalls in image analysis validation"), and author information ("You have full access to this article via Universitätsbibliothek Leipzig").

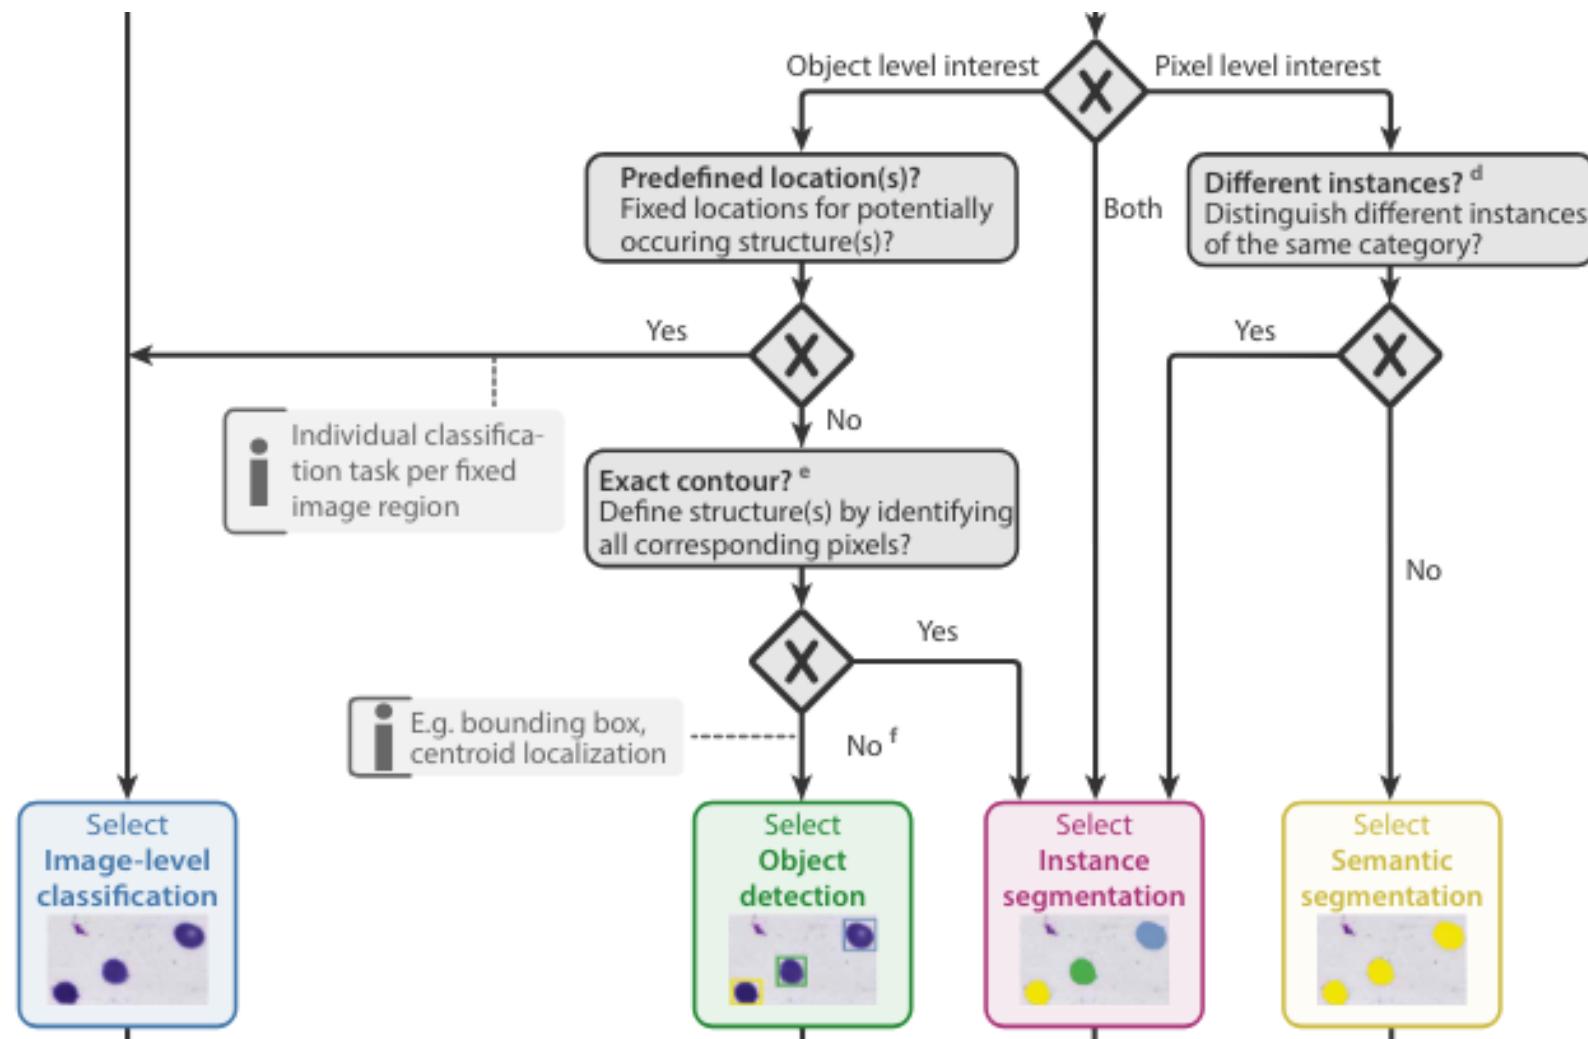
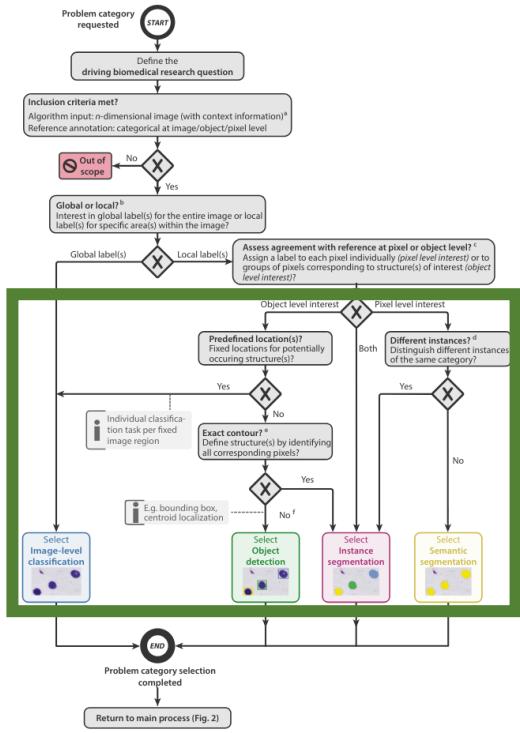
# Choosing the right metric is key

+ S1



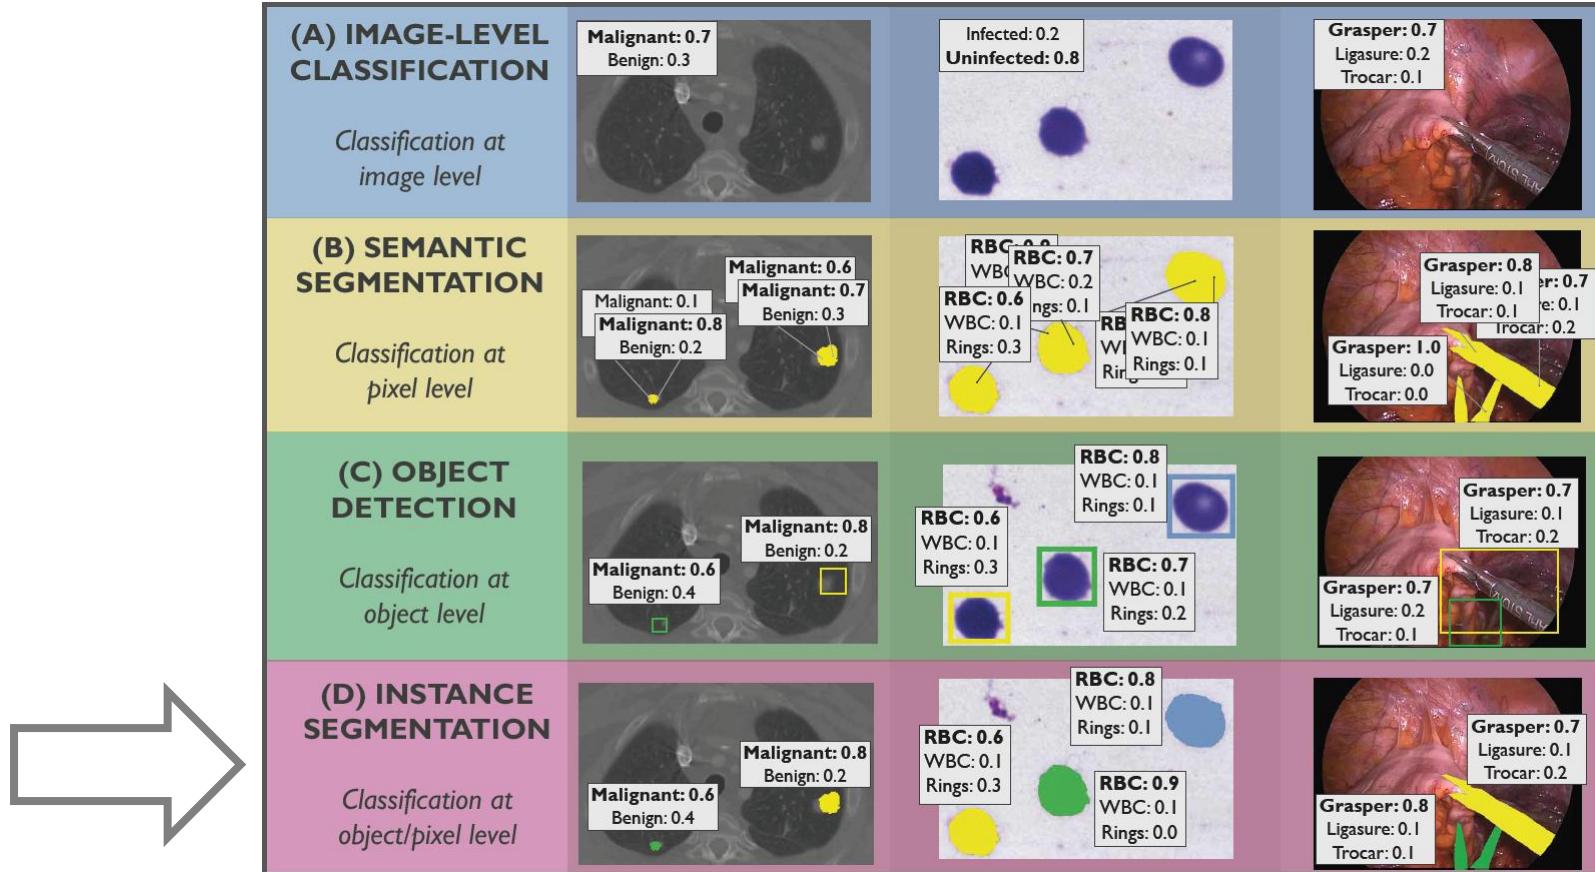
# Choosing the right metric is key

+ S1



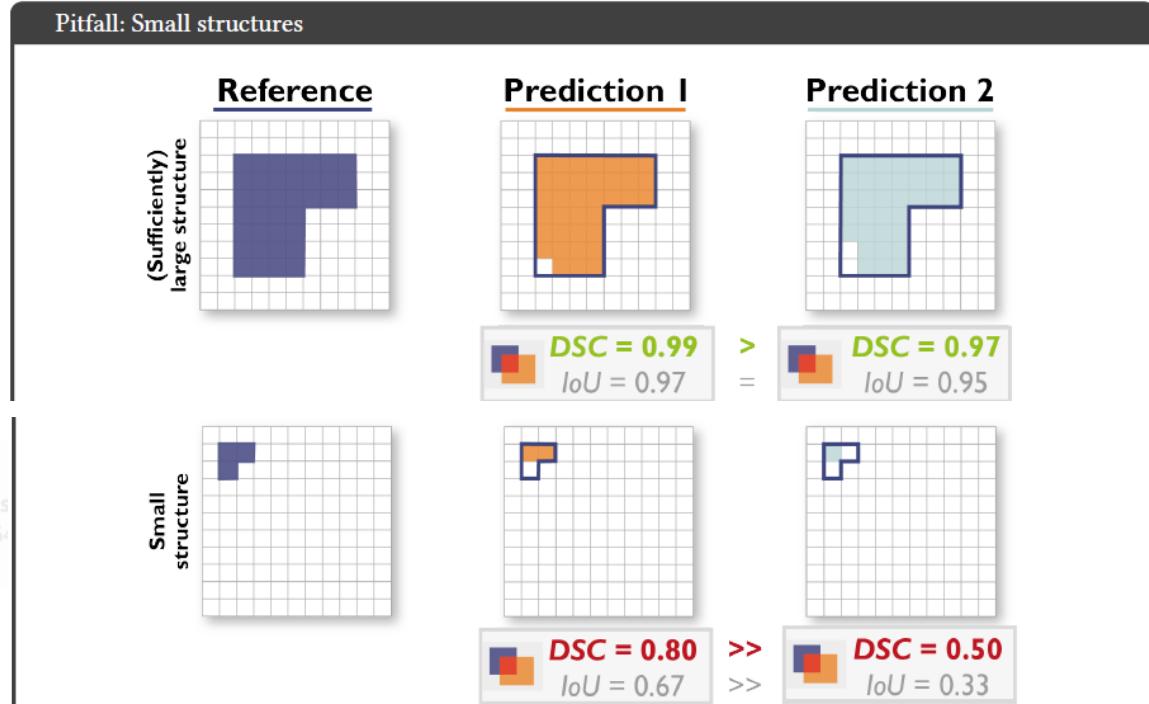
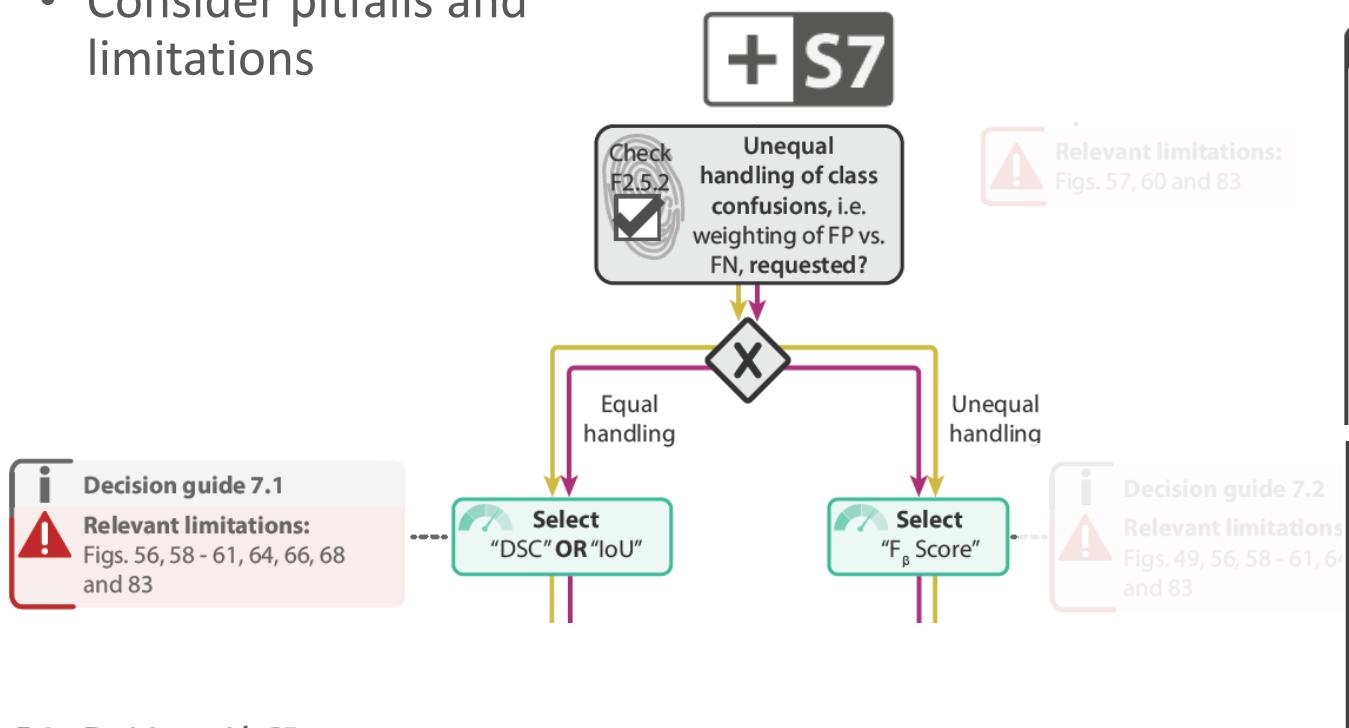
# Choosing the right metric is key

- Define your question



# Overlap metrics

- Consider pitfalls and limitations



(From Figure 56)

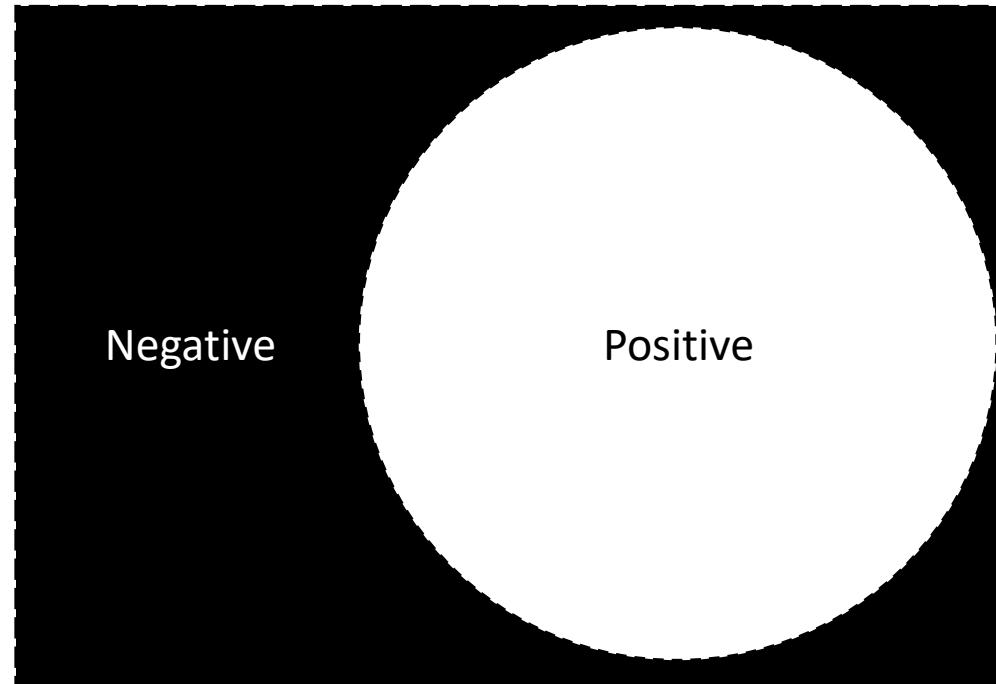
## E.6 Decision guide S7.

### D7.1: Dice vs IoU

The DSC is identical to the  $F_1$  Score on pixel level and closely related to the IoU, which, in turn, is identical to the Jaccard Index (see equations 5 and 6). The two metrics will yield the same ranking (of aggregated metric values) in most applications (theoretically, deviations are possible), such that there is no value in combining them. Commonly, the computer vision community prefers the IoU, while the medical image community favors the DSC.

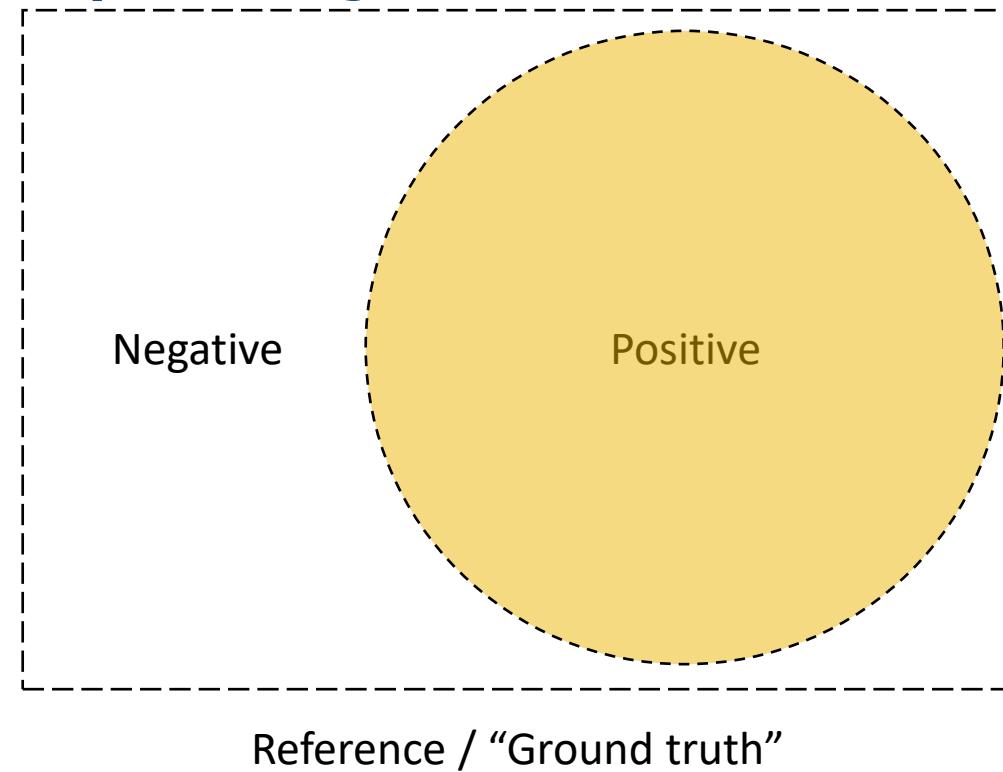
# Segmentation quality estimation

- In general
  - Define what's positive and what's negative.



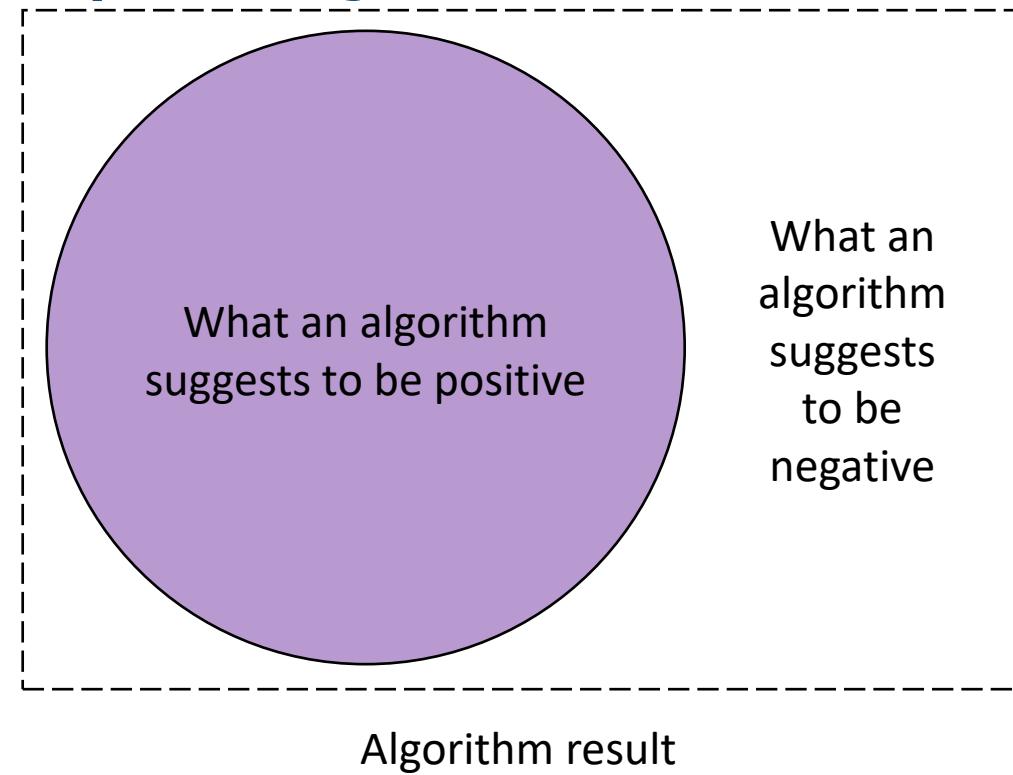
# Segmentation quality estimation

- In general
  - Define what's positive and what's negative.



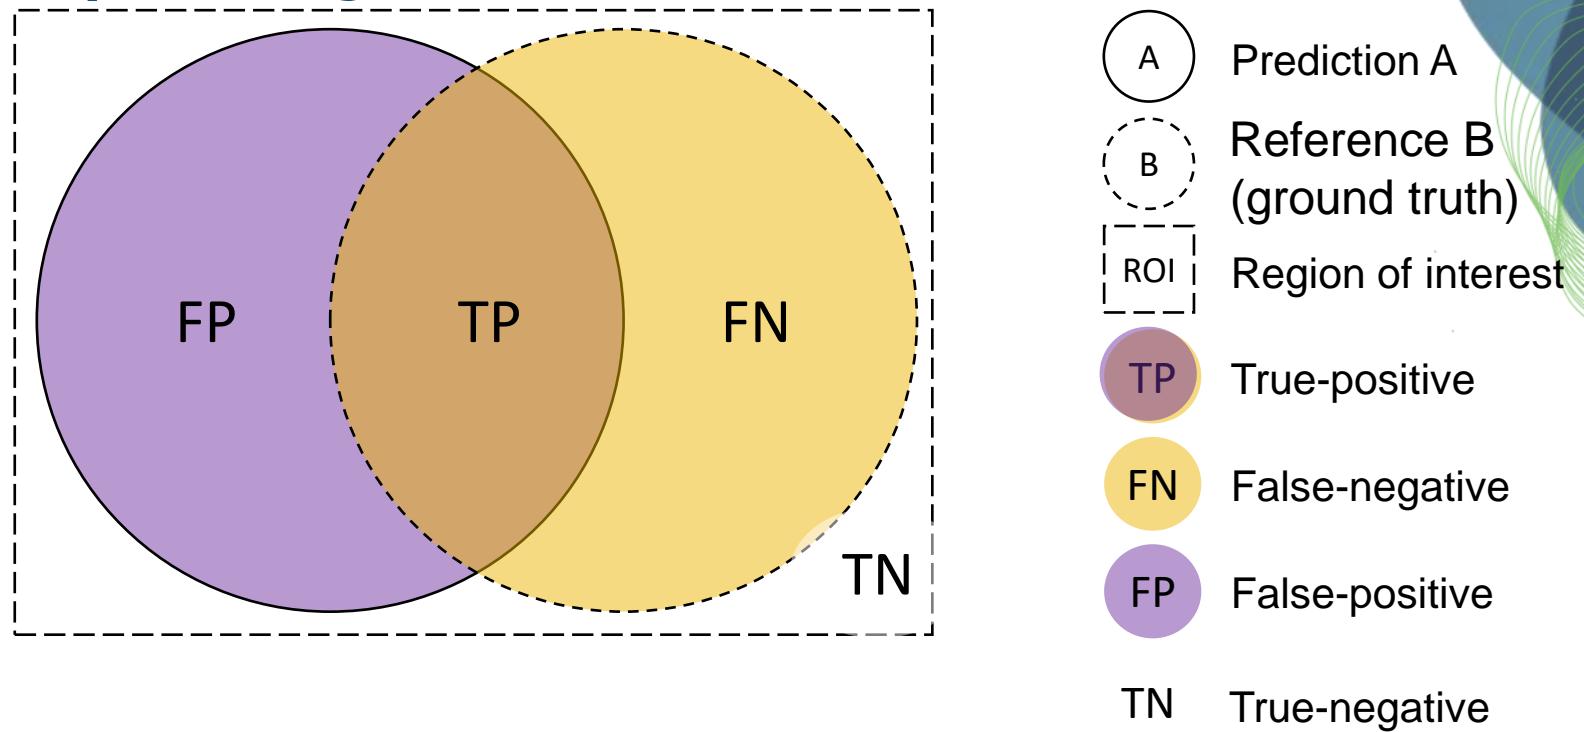
# Segmentation quality estimation

- In general
  - Define what's positive and what's negative.



# Segmentation quality estimation

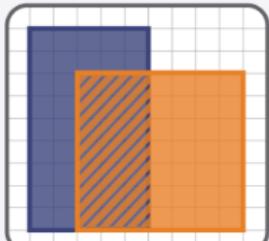
- In general
  - Define what's positive and what's negative.
  - Compare with a reference to figure out what was true and false
- Welcome to the Theory of Sets



# Dice versus Jaccard Index (IoU)

## DICE SIMILARITY COEFFICIENT (DSC)

*Synonyms:* Dice, Dice Coefficient, Sørensen–Dice Coefficient,  $F_1$  Score, Balanced F Score



$$\text{DSC}(A, B) = \frac{2 |A \cap B|}{|A| + |B|} = \frac{2 \text{PPV} \cdot \text{Sensitivity}}{\text{PPV} + \text{Sensitivity}}$$

VALUE RANGE: [0, 1] ↑

### DESCRIPTION

DSC measures the overlap between two structures.

### DEFINITION

[Dice, 1945]

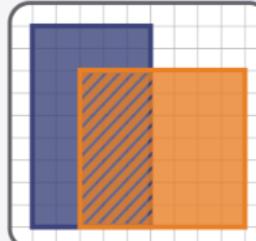
	RECOMMENDED FOR
ImLC	<input checked="" type="radio"/>
SemS	<input checked="" type="radio"/>
ObD	<input checked="" type="radio"/>
InS	<input checked="" type="radio"/>

### RECOMMENDATIONS

- An overlap-based metric (by default the DSC or IoU) should be used in most cases of segmentation assessment. An exception is the case of consistently tiny structures along with a noisy reference.
- DSC should generally be used in combination with a boundary-based metric if boundaries are of interest.
- DSC should generally not be considered if...
  - ... there is a high variability of structure sizes within an image or across images.
  - ... inter-rater variability is requested to be compensated.
  - ... over- and undersegmentation should be treated similarly.
- DSC should be considered as a metric in the medical community rather than in the computer vision and biology communities (where the almost identical IoU is preferred).

## INTERSECTION OVER UNION (IoU)

*Synonyms:* Jaccard Index, Tanimoto Coefficient



$$\begin{aligned}\text{IoU}(A, B) &= \frac{|A \cap B|}{|A| + |B| - |A \cap B|} \\ &= \frac{|A \cap B|}{|A \cup B|} = \frac{|A \cap B|}{|A| + |B| - |A \cap B|} \\ &= \frac{\text{PPV} \cdot \text{Sensitivity}}{\text{PPV} + \text{Sensitivity} - \text{PPV} \cdot \text{Sensitivity}}\end{aligned}$$

VALUE RANGE: [0, 1] ↑

### DESCRIPTION

IoU measures the overlap between two structures. It is often referred to as **Box IoU** when comparing bounding boxes, **Mask IoU** when comparing segmentation masks, or **Approx IoU** when comparing approximations of objects beyond bounding boxes.

### DEFINITION

[Jaccard, 1912]

### RECOMMENDED FOR

ImLC	SemS	ObD	InS
<input checked="" type="radio"/>	<input checked="" type="radio"/>	<input type="radio"/>	<input checked="" type="radio"/>

### IMPORTANT RELATIONS

$$\text{IoU} = \frac{\text{DSC}}{2 - \text{DSC}} \quad \text{IoU} = \frac{F_\beta}{2 - F_\beta} \quad \text{for } \beta = 1$$

### RECOMMENDATIONS

- An overlap-based metric (by default DSC or IoU) should be used in most cases for segmentation assessment. An exception is the case of consistently tiny structures along with a noisy reference.
- IoU should generally be used in combination with a boundary-based metric if boundaries are of interest.
- IoU should generally not be considered if...
  - ... there is a high variability of structure sizes within an image or across images.
  - ... inter-rater variability is requested to be compensated.
  - ... over- and undersegmentation should be treated similarly.
- IoU should be considered as a metric in the computer vision and biology communities rather than in the medical community (which prefers the almost identical DSC).



# Dice versus Jaccard Index (IoU)

## 2.7.5 Decision guide S6.

### DG6.1: Dice Similarity Coefficient (DSC) versus Intersection over Union (IoU)

Summary of DG6.1: DSC versus IoU

#### DSC

- ⇒ Identical to  $F_1$  Score
- ⇒ Close relation to IoU (see Eq. 5)
- ⇒ Preference in medical community

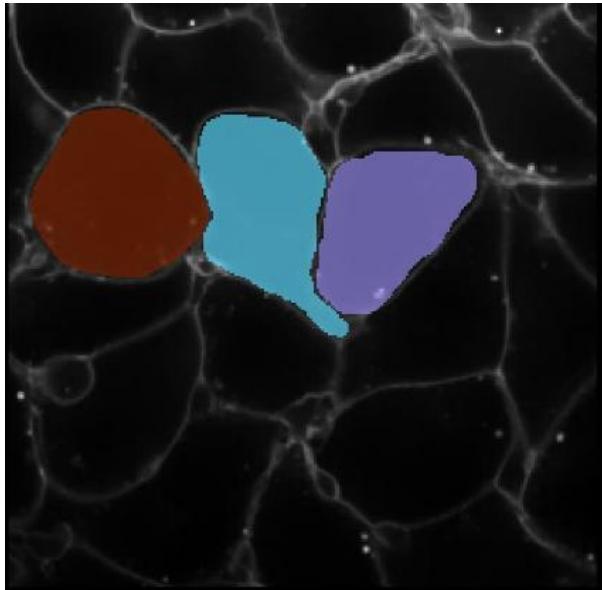
#### IoU

- ⇒ Identical to Jaccard Index
- ⇒ Close relation to DSC (see Eq. 4)
- ⇒ Preference in computer vision community

Extended Data Tab. SN 2.18. Comparison of Dice Similarity Coefficient (DSC) and Intersection over Union (IoU) in the context of the decision guide DG6.1 for Subprocess S6. Context: no exclusive interest in the center line of structures (FP2.3 = FALSE, FP3.3 = FALSE) and equal severity of class confusions (FP2.5.2 = FALSE).

# Sparse Jaccard Index

- For every annotated object, we compute the maximum IoU with any segmented object.
- We average this value over all annotated objects



Sparse  
instance annotation



IoU = 0.35



IoU = 0.66

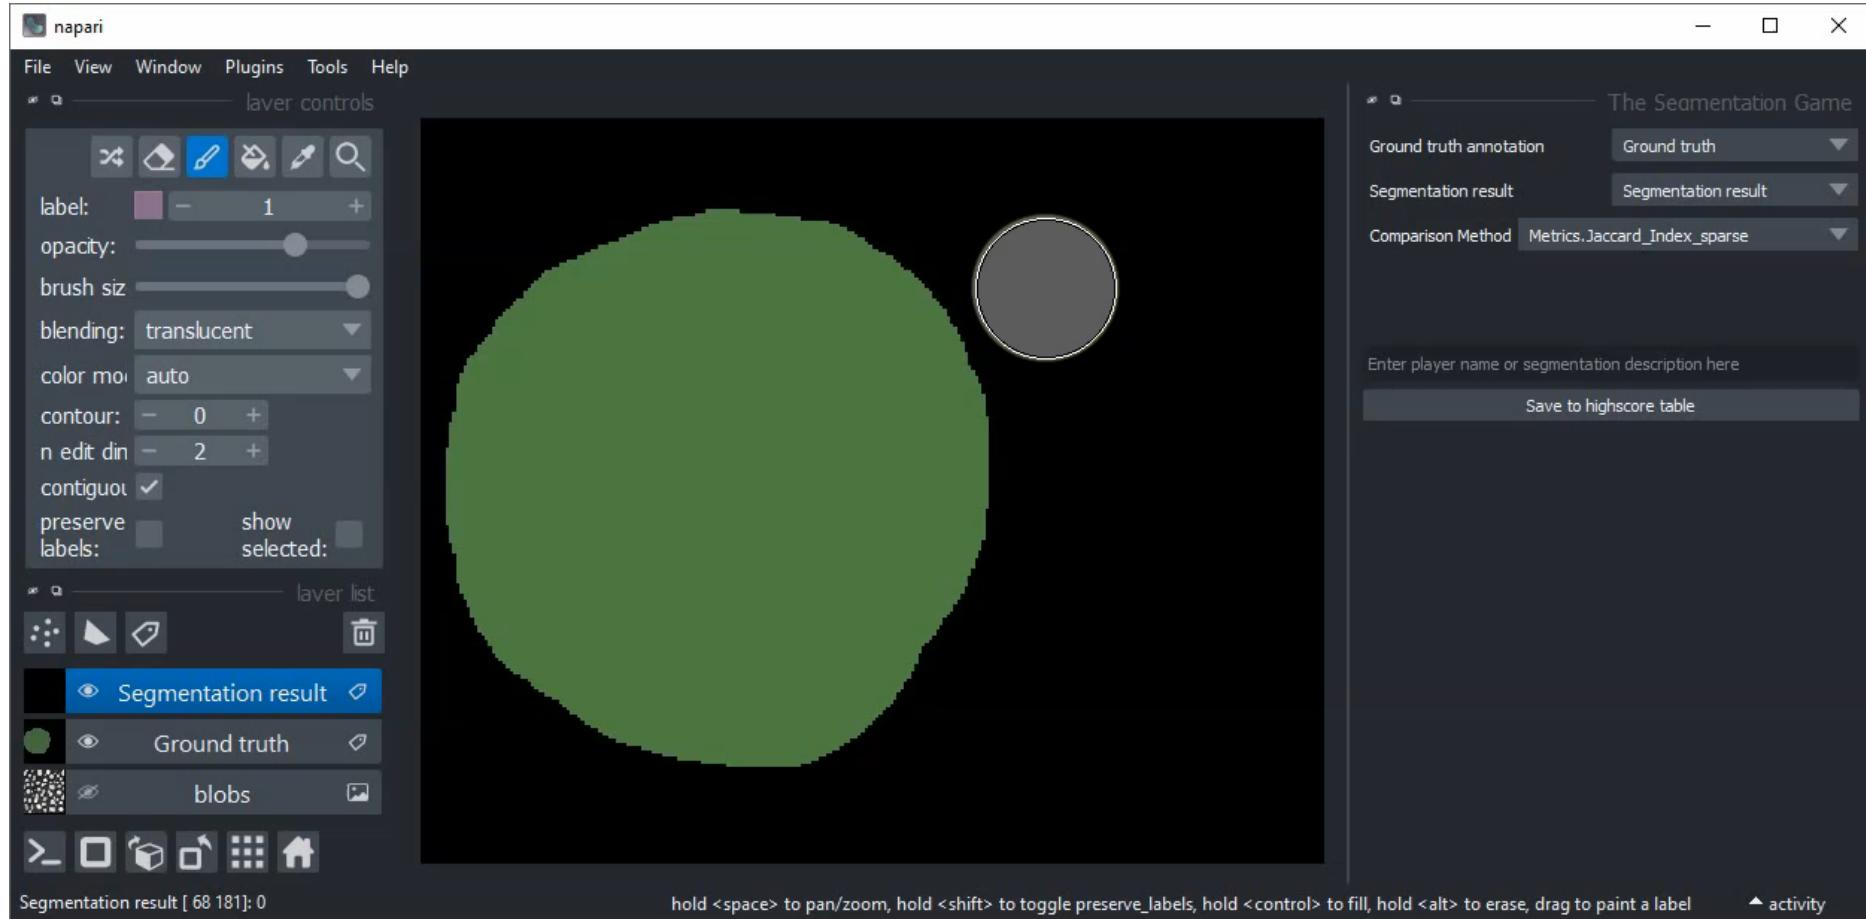


IoU = 0.69

Instance segmentation candidates

# Sparse Jaccard Index

- Play with metrics to gain understanding

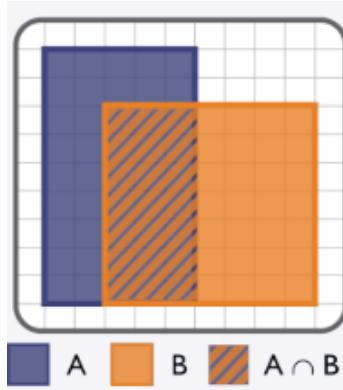
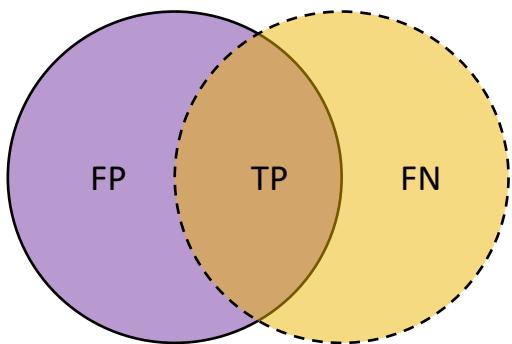


With Martin Schätz  
(Charles Uni, Prague)  
@schatzcz

# Pixel-wise versus Object-wise evaluation

- Pixel wise: Segmentation quality

Prediction      Ground truth



An object-level confusion matrix represented as a 3x3 grid. The columns are labeled 'A' (purple) and 'B' (yellow). The rows are labeled 'A' (purple), 'B' (yellow), and 'A ∩ B' (brown). The matrix values are: A-A: 1, A-B: 0, A ∩ B: 1; B-A: 0, B-B: 2, B ∩ A: 1; A ∩ B-A: 0, A ∩ B-B: 1, A ∩ B ∩ A ∩ B: 1.

	A	B	$A \cap B$
A	1	0	1
B	0	2	1
$A \cap B$	0	1	1

True-positive: 4

False-negative: 5

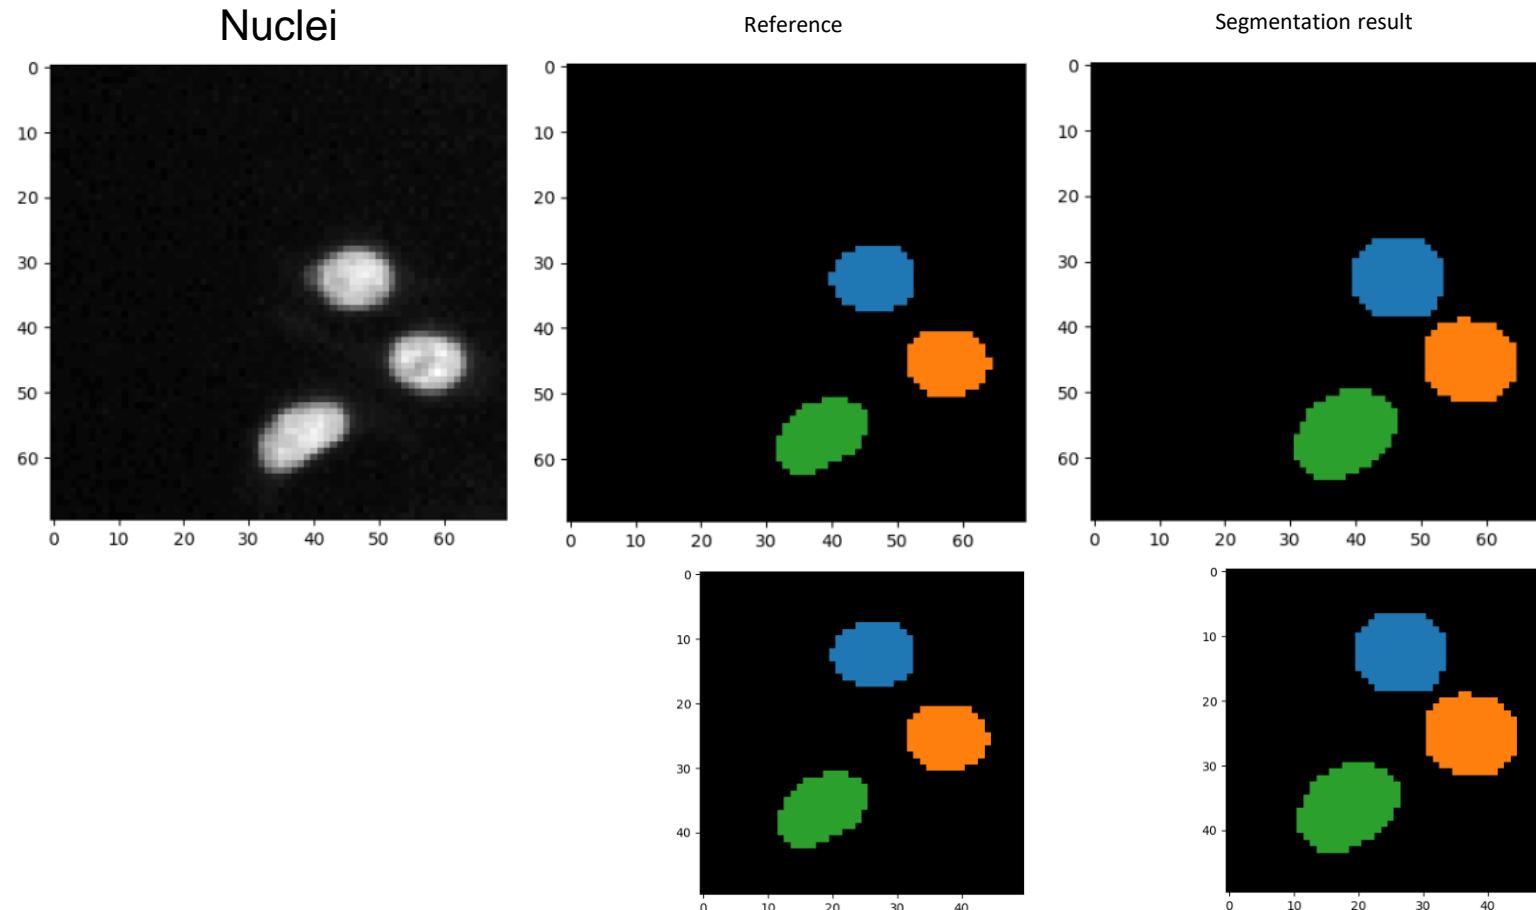
False-positive: 2

$$\begin{aligned} \text{IoU}(A, B) &= \frac{\text{A} \cap \text{B}}{\text{A} + \text{B} - \text{A} \cap \text{B}} \\ &= \frac{|\text{A} \cap \text{B}|}{|\text{A}| + |\text{B}| - |\text{A} \cap \text{B}|} = \frac{|\text{A} \cap \text{B}|}{|\text{A} \cup \text{B}|} \end{aligned}$$

$\text{IoU} = 4 / 11$

# Accuracy versus Jaccard Index (IoU)

- Side-effects of image size and number of nuclei



$$A = \frac{TP + TN}{FN + FP + TP + TN}$$

$$J = \frac{TP}{FN + FP + TP}$$

Accuracy: 0.97

Jaccard Index: 0.73

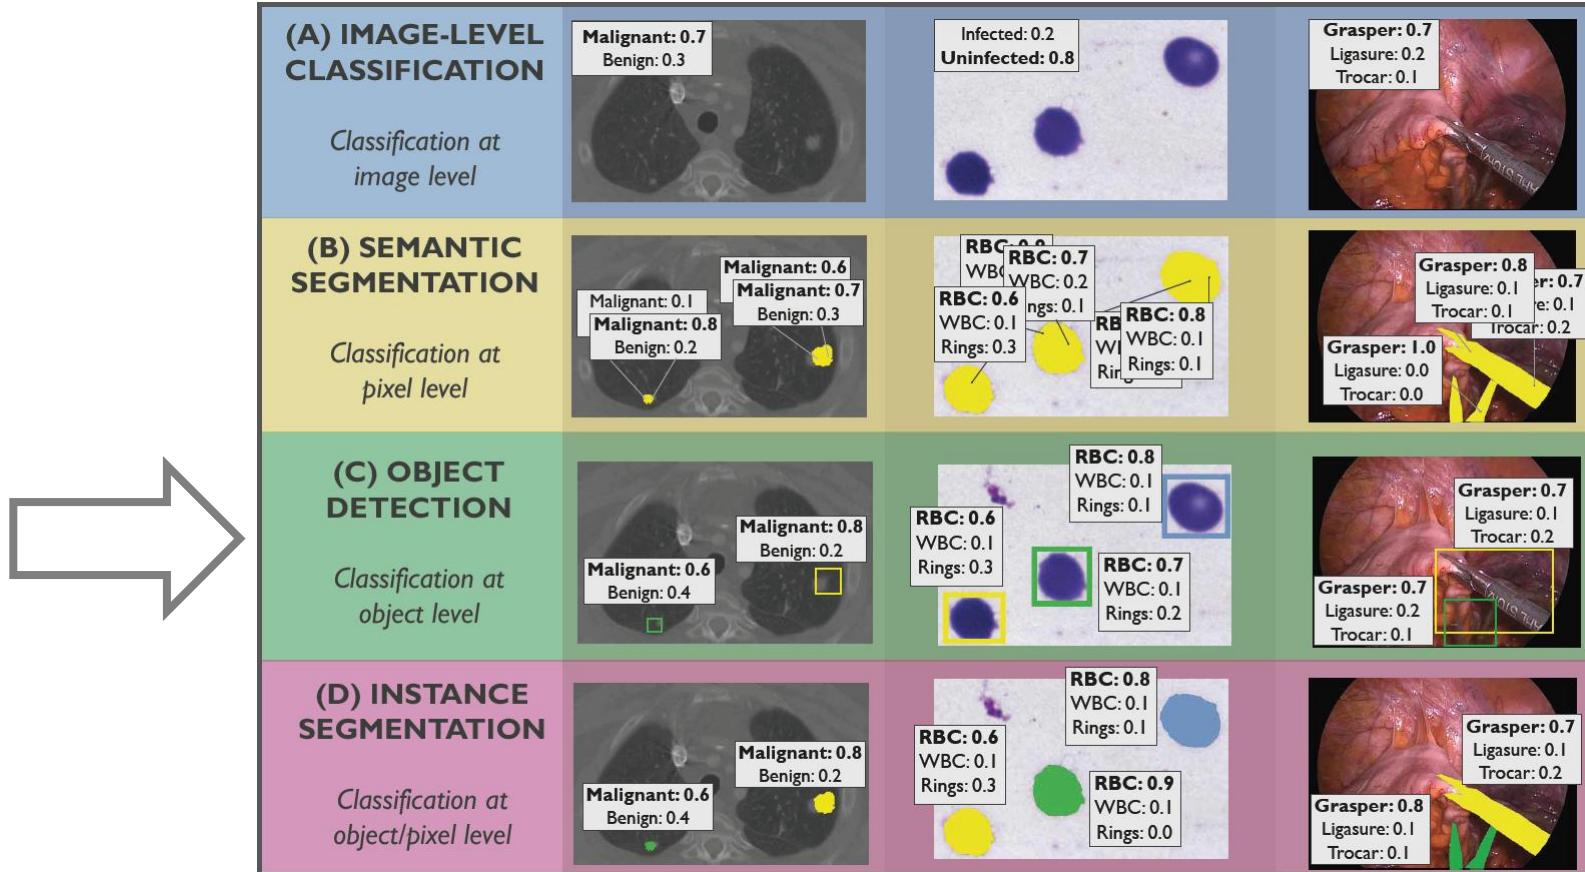
Accuracy decreases because  
there are less correct black  
pixels (TN)

Accuracy: 0.95

Jaccard Index: 0.73

# Choosing the right metric is key

- Define your question



# Pixel-wise versus Object-wise evaluation

- Are these objects overlapping?

		FN	FN			
FP	TP	FN				
FP	FP					

$$\text{IoU} = 1 / 7$$

		FP	TP	FN		
		FP	TP	FN		

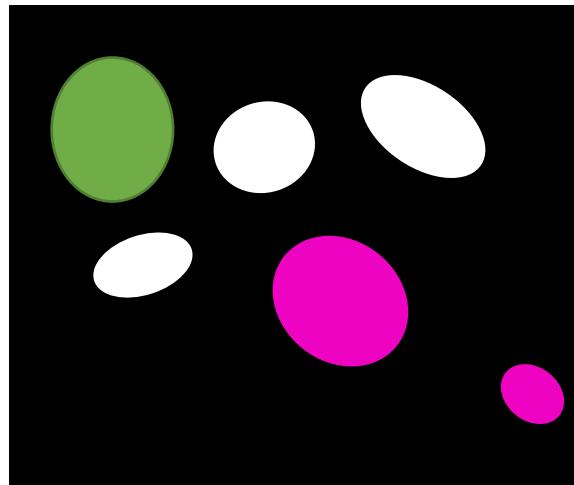
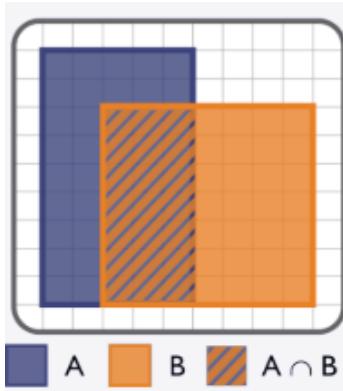
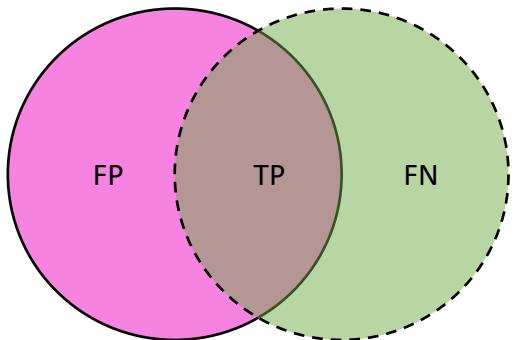
$$\text{IoU} = 1 / 3$$

		FP	TP	TP	FN	
		FP	TP	TP	FN	

$$\text{IoU} = 1 / 2$$

# Pixel-wise versus Object-wise evaluation

- Object wise: Detection quality



True-positive: 3

False-negative: 1

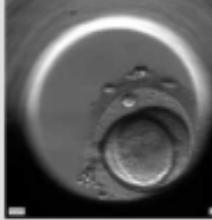
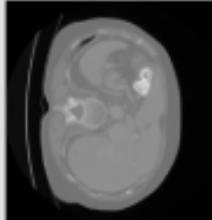
False-positive: 2

$$\begin{aligned} \text{IoU}(A, B) &= \frac{\text{A} \cap \text{B}}{\text{A} + \text{B} - \text{A} \cap \text{B}} \\ &= \frac{|\text{A} \cap \text{B}|}{|\text{A}| + |\text{B}| - |\text{A} \cap \text{B}|} = \frac{|\text{A} \cap \text{B}|}{|\text{A} \cup \text{B}|} \end{aligned}$$

$$\text{IoU} = 1 / 2$$

# Special case in medical imaging

- We elaborate segmentation quality of one / large object:

PROBLEM DESCRIPTION	ID	SCENARIO	SAMPLE INPUT IMAGE	RECOMMENDED OUPUT	RECOMMENDATION
Segmentation of large objects	SemS-1	Embryo segmentation from microscopy images			<p><b>Problem category:</b> <b>Semantic segmentation</b></p> <p><b>Overlap-based metric (S6):</b> Dice Similarity Coefficient (DSC)</p> <p><b>Boundary-based metric (S7):</b> Normalized Surface Distance (NSD)</p> <p><b>Specific property-related metric:</b> Liver segmentation: Absolute Volume Difference</p>
	SemS-2	Liver segmentation in computed tomography (CT) images			

# Further reading

Quality assurance of segmentation results

**FocalPlane**  
Where biology meets microscopy

Home About us Topics Gallery Jobs Events  
Network Contact us Log in/register

Home / Default / Quality assurance of segmentation result ...

## Quality assurance of segmentation results

Posted by Mara Lampert, on 13 April 2023

This blog post revolves around determining and improving the quality of segmentation results. A common problem is that this step is often omitted and done rather than by actually quantifying it. This blogpost aims to achieve this quantification as this leads to reproducibility.

The Jaccard index

The **Jaccard Index** is a measure to investigate the similarity or difference of sample sets. Hereby, the so called **sparse Jaccard Index** measures the overlap lying between 0 (no overlap) and 1 (perfect overlap). Therefore, the ground truth label is compared to the segmented label by determination of the maximum overlap. The metric result is the mean overlap of all investigated labels (Jaccard, 1902) (see this jupyter notebook).

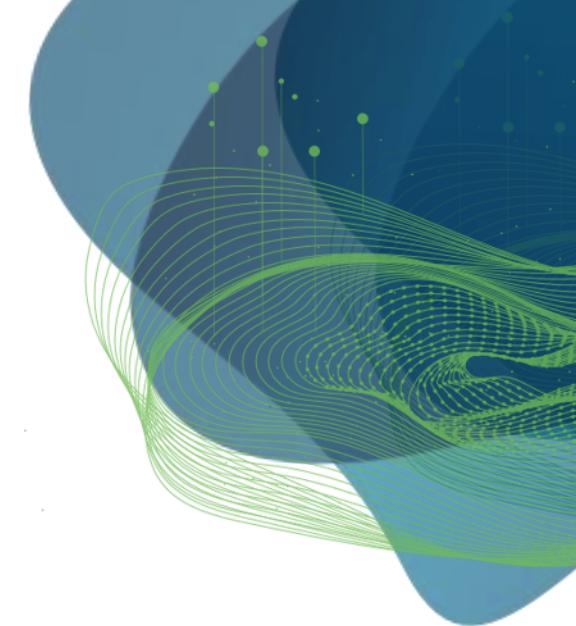


Mara Lampert  
@marabuuuu.bsky.social



DRESDEN LEIPZIG

CENTER FOR SCALABLE DATA ANALYTICS  
AND ARTIFICIAL INTELLIGENCE



# Exercises

Robert Haase

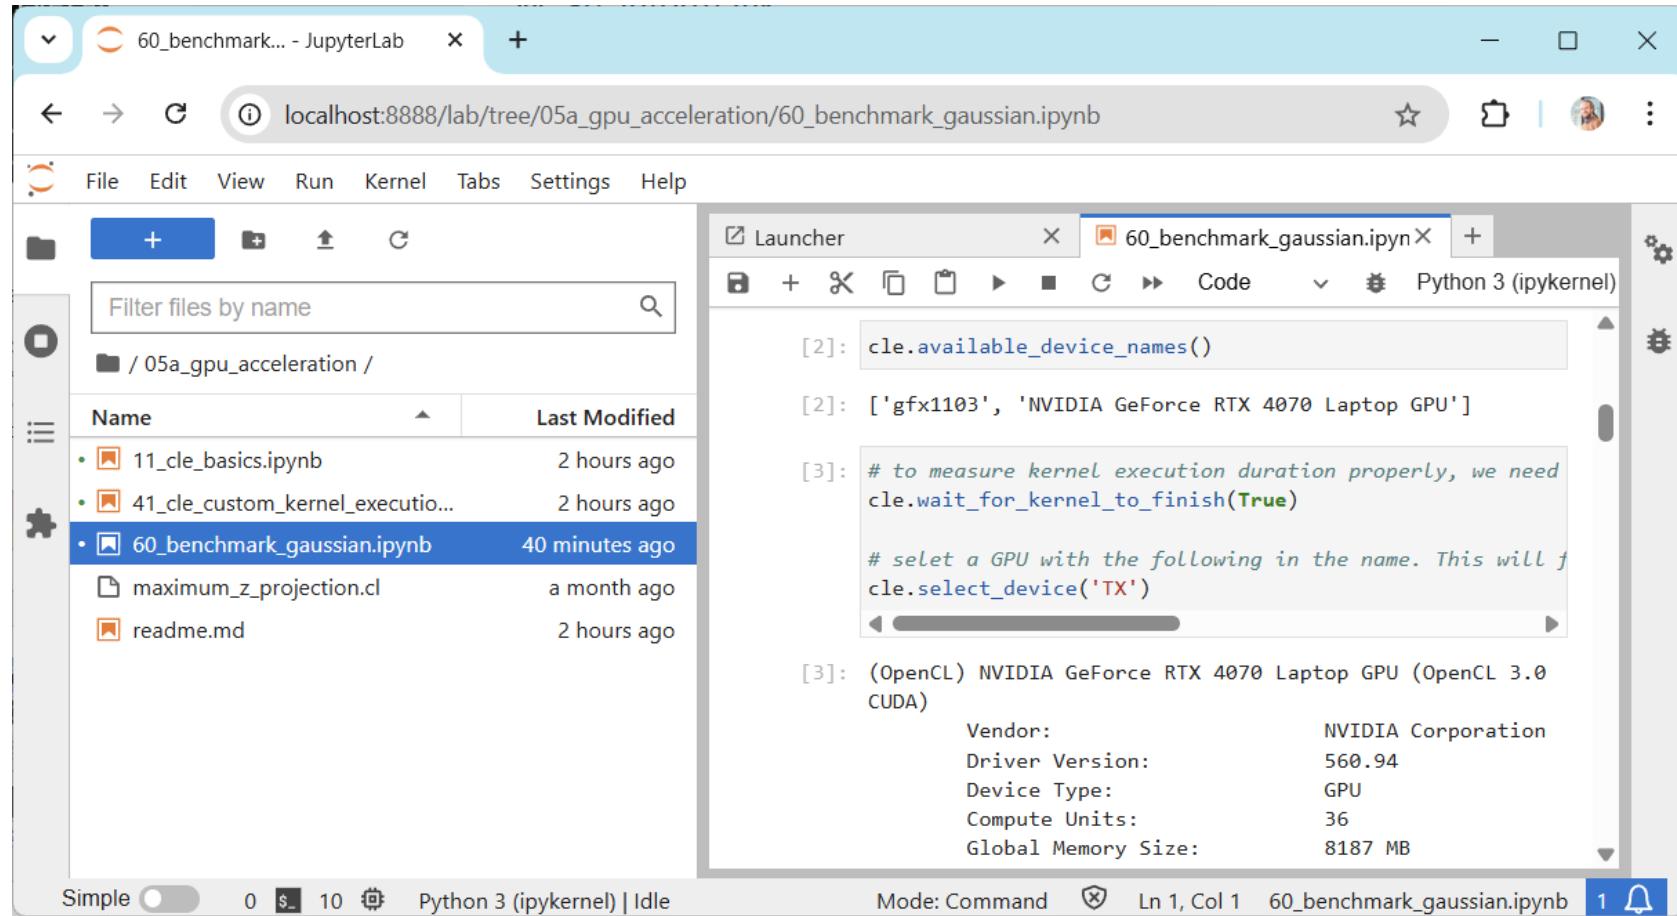
GEFÖRDERT VOM



Bundesministerium  
für Bildung  
und Forschung

# Exercise: GPU-accelerated image processing

- Compare CPU processing speed with a GPU, consider using the cluster for this



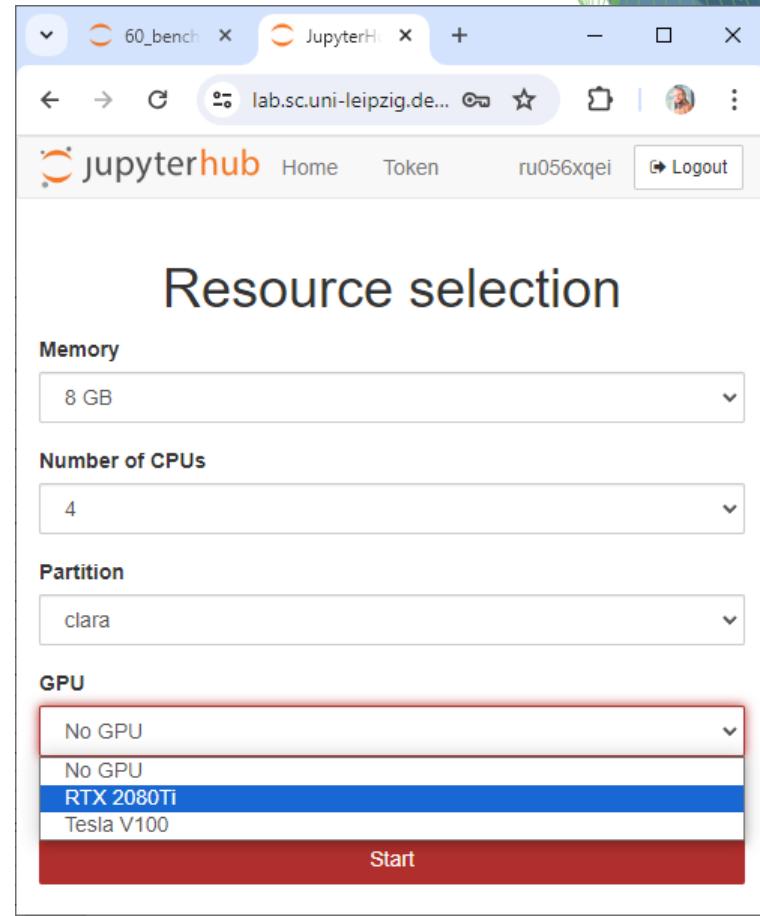
The screenshot shows a JupyterLab interface. On the left, a file tree displays several notebooks and files under the directory `/05a_gpu_acceleration/`. The notebook `60_benchmark_gaussian.ipynb` is selected and open in the main area. The code editor contains the following Python code:

```
[2]: cle.available_device_names()
[2]: ['gfx1103', 'NVIDIA GeForce RTX 4070 Laptop GPU']
[3]: # to measure kernel execution duration properly, we need
       cle.wait_for_kernel_to_finish(True)

# select a GPU with the following in the name. This will f
       cle.select_device('TX')

[3]: (OpenCL) NVIDIA GeForce RTX 4070 Laptop GPU (OpenCL 3.0
       CUDA)
       Vendor: NVIDIA Corporation
       Driver Version: 560.94
       Device Type: GPU
       Compute Units: 36
       Global Memory Size: 8187 MB
```

The status bar at the bottom indicates "Python 3 (ipykernel) | Idle".



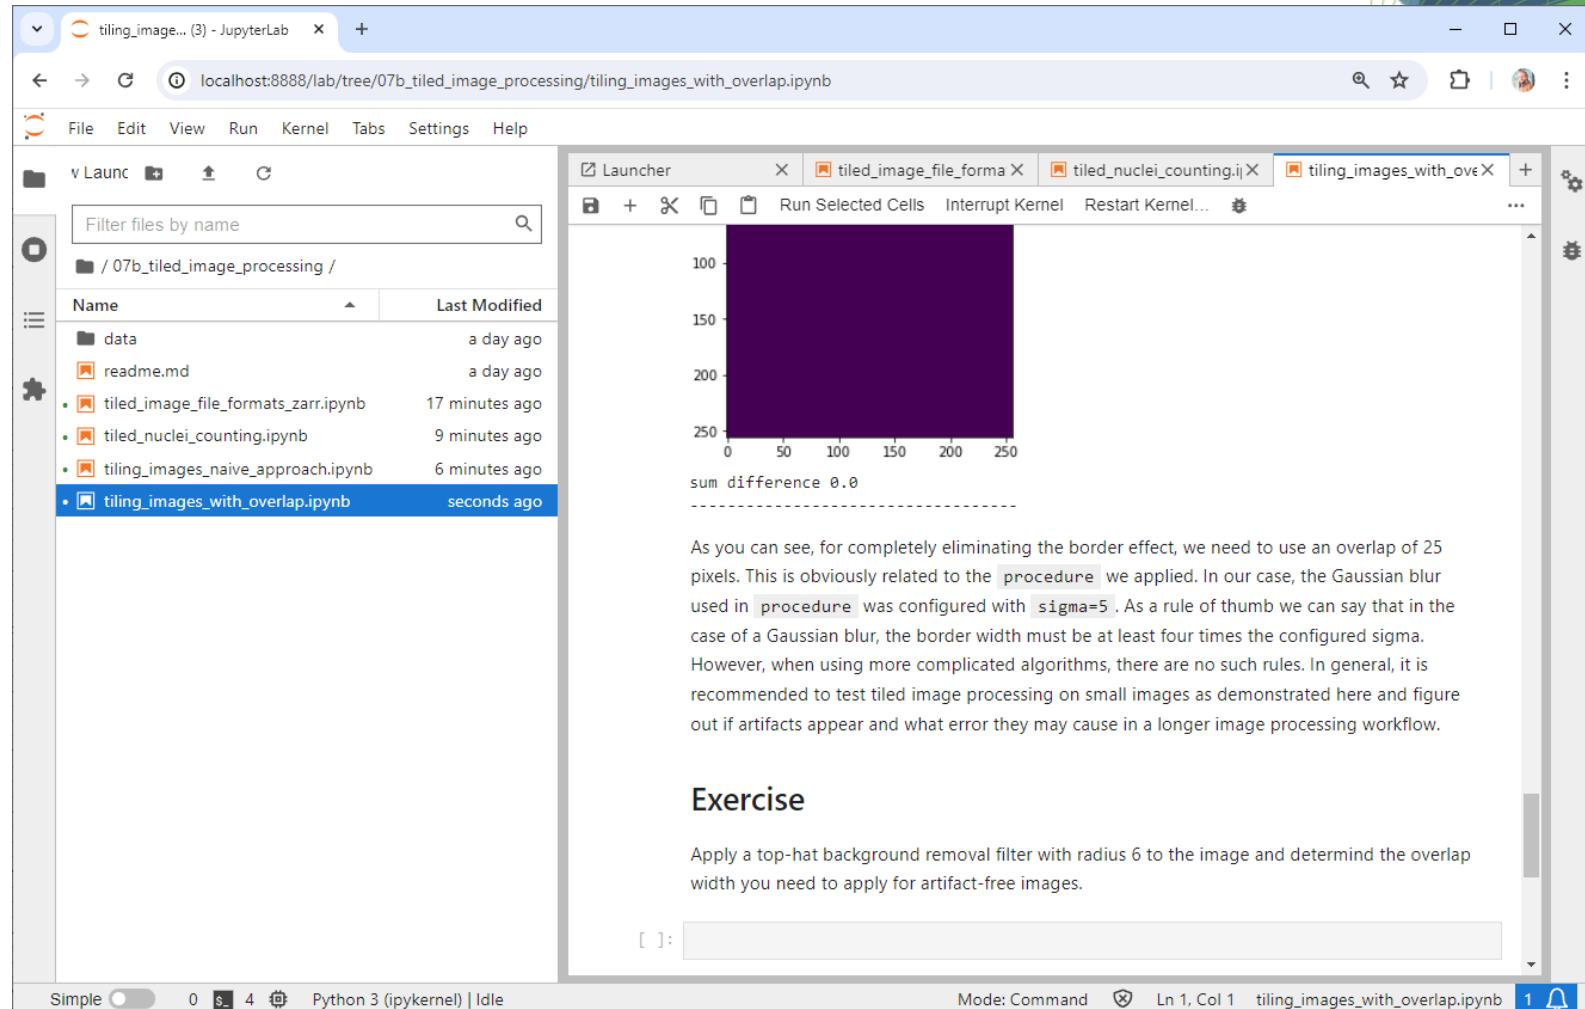
The screenshot shows a JupyterHub interface. The title bar includes the URL `lab.sc.uni-leipzig.de...`, a user icon, and a "Logout" button. The main content area is titled "Resource selection" and contains the following configuration fields:

Memory	8 GB
Number of CPUs	4
Partition	clara
GPU	No GPU No GPU <b>RTX 2080Ti</b> Tesla V100

A red "Start" button is located at the bottom right of the GPU dropdown menu.

# Exercise: Tiled image processing

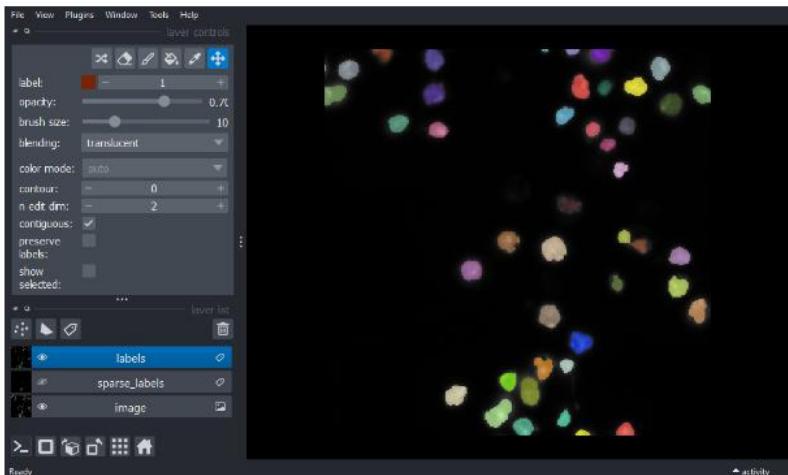
- Apply background-removal to an image in tiles. Determine the overlap width that's necessary to have artifact-free results.



# Segmentation quality

- Measure segmentation quality of a given algorithm applied to a folder of images.

```
[7]: def my_segmentation_algorithm(input_image):  
  
    # background subtraction  
    background_subtracted = nsbatwm.white_tophat(input_image, radius = 10)  
  
    # instance segmentation / labeling  
    labels_result = nsbatwm.voronoi_otsu_labeling(background_subtracted,  
  
    return labels_result
```



## Quality estimation: Sparse Jaccard Index

From the two label images loaded and produced above we can compute the sparse Jaccard Index.

```
[9]: metrics.jaccard_index_sparse(sparse_labels, labels)  
[9]: 0.8357392602053431
```

## Exercise

Use the following for-loop and code snippets from above to compute the segmentation quality of all images in the folder. Provide the average quality over all images.

```
[10]: for image_filename in os.listdir(image_folder):  
    print(image_folder + image_filename)  
  
data/BBBC007_batch/17P1_POS0013_D_1UL.tif  
data/BBBC007_batch/20P1_POS0005_D_1UL.tif  
data/BBBC007_batch/20P1_POS0007_D_1UL.tif  
data/BBBC007_batch/20P1_POS0010_D_1UL.tif  
data/BBBC007_batch/A9_p7d.tif  
data/BBBC007_batch/AS_09125_040701150004_A02f00d0.tif
```

# Optional exercise

- Play Bio-image Analysis Bubbles and test your skills: Can you build a workflow that reproduces the results?



Major parts of this game  
were written by the AI in  
cursor.  
<https://www.cursor.com/>